

Introduction to Generalised Low-Rank Model and Missing Values



Strata+Hadoop
WORLD

Jo-fai (Joe) Chow

Data Scientist

joe@h2o.ai

@matlabulus

Based on work by Anqi Fu, Madeleine Udell, Corinne Horn,
Reza Zadeh & Stephen Boyd.

About H2O.ai

- H2O is an open-source, distributed machine learning library written in Java with APIs in R, Python, Scala and REST/JSON.
- Produced by H2O.ai in Mountain View, CA.
- H2O.ai advisers are Trevor Hastie, Rob Tibshirani and Stephen Boyd from Stanford.



About Me

- 2005 - 2015
- Water Engineer
 - Consultant for Utilities
 - EngD Research
- 2015 - Present
- Data Scientist
 - Virgin Media
 - Domino Data Lab
 - H2O.ai

About This Talk

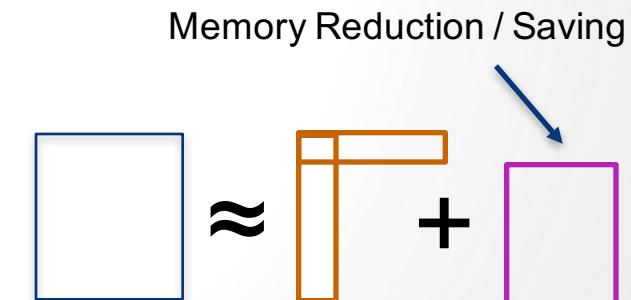
- Overview of generalised low-rank model (GLRM).
- Four application examples:
 - Basics.
 - How to accelerate machine learning.
 - How to visualise clusters.
 - How to impute missing values.
- Q & A.

GLRM Overview

- GLRM is an extension of well-known matrix factorisation methods such as Principal Component Analysis (PCA).
- Unlike PCA which is limited to numerical data, GLRM can also handle categorical, ordinal and Boolean data.
- **Given:** Data table A with m rows and n columns
- **Find:** Compressed representation as numeric tables X and Y where k is a small user-specified number

$$m \left\{ \overbrace{\begin{bmatrix} A \end{bmatrix}}^n \right\} \approx m \left\{ \begin{bmatrix} X \end{bmatrix} \left[\overbrace{\begin{bmatrix} Y \end{bmatrix}}^n \right] \right\} k$$

- Y = archetypal features created from columns of A
- X = row of A in reduced feature space
- GLRM can approximately reconstruct A from product XY



GLRM Key Features

- **Memory**
 - Compressing large data set with minimal loss in accuracy
- **Speed**
 - Reduced dimensionality = short model training time
- **Feature Engineering**
 - Condensed features can be analysed visually
- **Missing Data Imputation**
 - Reconstructing data set will automatically impute missing values

GLRM Technical References

- Paper
 - arxiv.org/abs/1410.0342
- Other Resources
 - [H2O World Video](#)
 - [Tutorials](#)

Generalized Low Rank Models

Madeleine Udell, Corinne Horn, Reza Zadeh, and Stephen Boyd

May 6, 2015. (Original version posted September 2014.)

Abstract

Principal components analysis (PCA) is a well-known technique for approximating a tabular data set by a low rank matrix. Here, we extend the idea of PCA to handle arbitrary data sets consisting of numerical, Boolean, categorical, ordinal, and other data types. This framework encompasses many well known techniques in data analysis, such as nonnegative matrix factorization, matrix completion, sparse and robust PCA, k -means, k -SVD, and maximum margin matrix factorization. The method handles heterogeneous data sets, and leads to coherent schemes for compressing, denoising, and imputing missing entries across all data types simultaneously. It also admits a number of interesting interpretations of the low rank factors, which allow clustering of examples or of features. We propose several parallel algorithms for fitting generalized low rank models, and describe implementations and numerical results.

Example 1: Motor Trend Car Road Tests

“mtcars” dataset in R

$m = 32$

$n = 11$

> mtcars

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.97	2.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.93	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.93	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.07	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.335	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.730	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	4.93	5.250	18.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	4.00	5.424	18.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.90	4.423	17.82	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.280	18.30	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.90	1	1	4	2
Toyota Corolla	33.9	4	71.1	55	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	127.0	97	3.70	2.465	20.00	0	0	3	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

Original Data Table

$$m \left\{ \begin{bmatrix} n \\ A \end{bmatrix} \right\} \approx m \left\{ \begin{bmatrix} k \\ X \end{bmatrix} \left[\begin{bmatrix} n \\ Y \end{bmatrix} \right] k \right\}$$

Example 1: Training a GLRM

```
# Initialise H2O Cluster
library(h2o)
h2o.init(nthreads = -1)

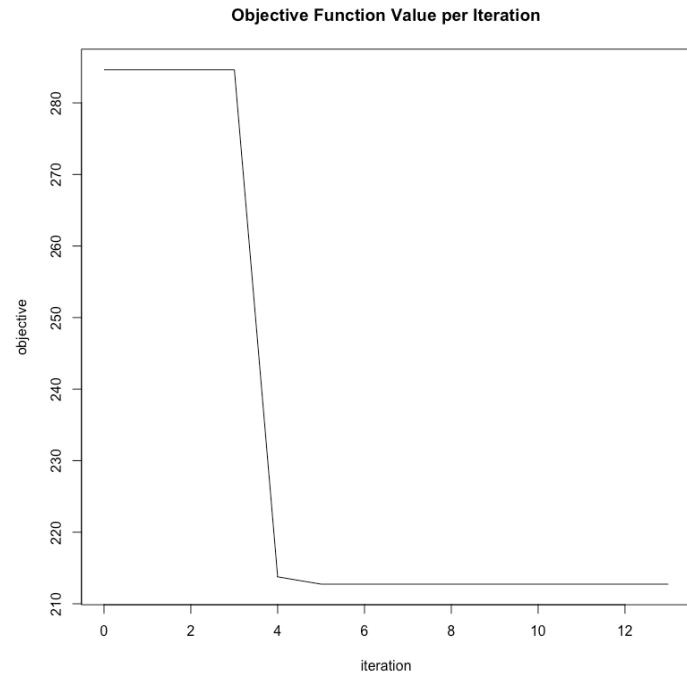
# Load and import mtcars data
data(mtcars)
hex_mtcars ← as.h2o(mtcars)

# Build a GLRM and compress data into a 3-column matrix
glrm_mtcars ← h2o.glrm(training_frame = hex_mtcars,
                         k = 3, seed = 1234)

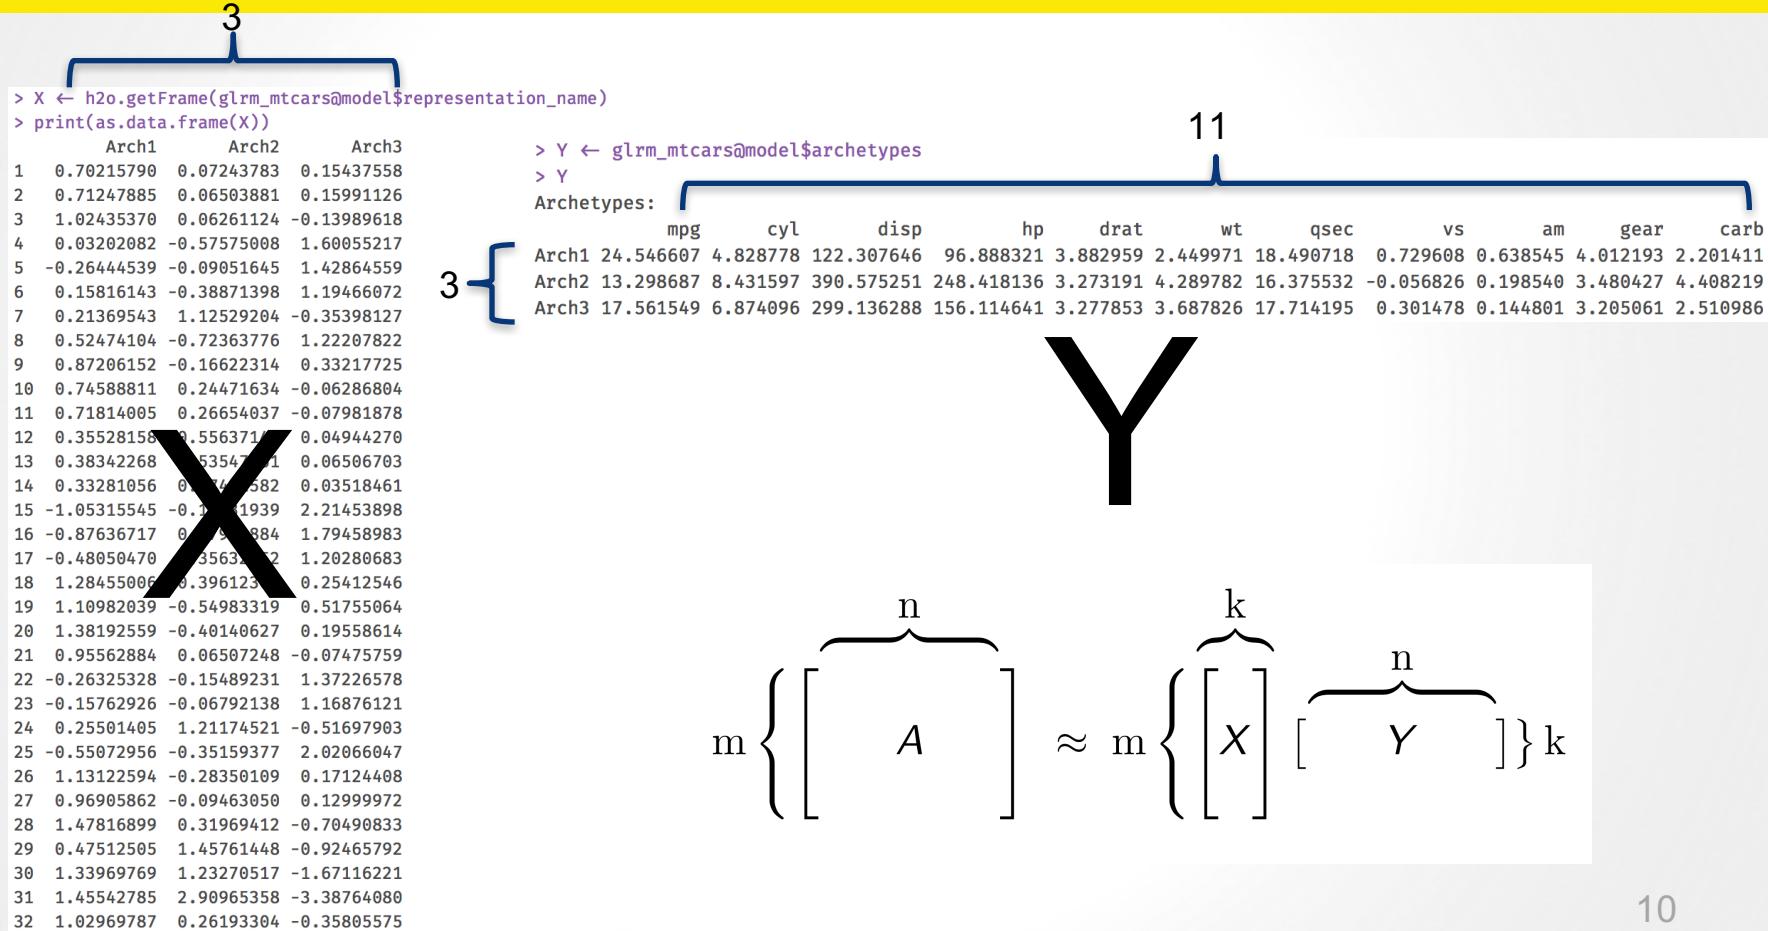
# Check convergence
plot(glrm_mtcars)
```

$$m \left\{ \begin{bmatrix} n \\ A \end{bmatrix} \right\} \approx m \left\{ \begin{bmatrix} k \\ X \end{bmatrix} \left[\begin{bmatrix} n \\ Y \end{bmatrix} \right] \right\} k$$

Check convergence



Example 1: X and Y from GLRM



Example 1: Summary

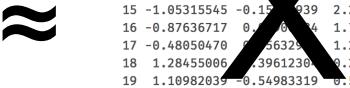
> mtcars

```

mpg cyl disp hp drat wt qsec vs am gear carb
Mazda RX4    21.0   6 160.0 110 3.98 2.620 16.46  0  1   4   4
Mazda RX4 Wag 21.0   6 160.0 110 3.98 2.875 17.02  0  1   4   4
Datsun 710   22.8   4 108.0  93 3.85 2.320 18.61  1  1   4   1
Hornet 4 Drive 21.4   6 258.0 110 3.08 3.215 19.44  1  0   3   1
Hornet Sportabout 18.7   8 360.0 175 3.15 3.440 17.02  0  0   3   2
Valiant   18.1   6 225.0 105 2.76 3.460 20.22  1  0   3   1
Duster 360   14.3   8 360.0 245 3.21 3.570 15.84  0  0   3   4
Merc 240D   24.4   4 146.7  62 3.69 3.190 20.00  1  0   4   2
Merc 230    22.8   4 140.8  95 3.92 3.150 22.90  1  0   4   2
Merc 280    19.2   6 167.6 123 3.92 3.440 18.30  1  0   4   4
Merc 280C   17.8   6 167.6 123 3.02 3.440 18.90  1  0   4   4
Merc 450SE   16.4   8 275.8 180 3.02 4.070 17.40  0  0   3   3
Merc 450SL   17.3   8 275.8 180 3.07 3.730 17.60  0  0   3   3
Merc 450SLC  15.2   8 275.8 180 3.07 3.780 18.00  0  0   3   3
Cadillac Fleetwood 10.4   8 472.0 205 4.93 5.420 17.98  0  0   3   4
Lincoln Continental 10.4   8 460.0 225 3.93 5.760 17.82  0  0   3   4
Chrysler Imperial 14.7   8 440.0 235 3.23 5.330 17.42  0  0   3   4
Fiat 128    32.4   4  78.7  65 4.08 2.200 19.47  1  1   4   1
Honda Civic  30.4   4  75.7  52 4.93 1.615 18.52  1  1   4   2
Toyota Corolla 33.9   4  71.1  65 4.22 1.835 19.90  1  1   4   1
Toyota Corona 21.5   4 120.1  97 3.70 2.465 20.01  1  0   3   1
Dodge Challenger 15.5   8 318.0 150 2.76 3.520 16.87  0  0   3   2
AMC Javelin  15.2   8 304.0 150 3.15 3.435 17.30  0  0   3   2
Camaro Z28   13.3   8 350.0 245 3.73 3.840 15.41  0  0   3   4
Pontiac Firebird 19.2   8 400.0 175 3.08 3.845 17.05  0  0   3   2
Fiat X1-9    27.3   4  79.0  66 4.08 1.930 18.90  1  1   4   1
Porsche 914-2 26.0   4 120.3  91 4.43 2.140 16.70  0  1   5   2
Lotus Europa  30.4   4  95.1 113 3.77 1.513 16.90  1  1   5   2
Ford Pantera L 15.8   8 351.0 264 4.22 3.170 14.50  0  1   5   4
Ferrari Dino  19.7   6 145.0 175 3.62 2.770 15.50  0  1   5   6
Maserati Bora 15.0   8 301.0 335 3.54 3.570 14.60  0  1   5   8
Volvo 14E    21.4   4 121.0 109 4.11 2.780 18.60  1  1   4   2

```

$$m \left\{ \begin{bmatrix} A \end{bmatrix} \right\} \approx m \left\{ \begin{bmatrix} X \end{bmatrix} \left[\begin{bmatrix} Y \end{bmatrix} \right] \right\} k$$



```

> X <- h2o.getFrame(glrm_mtcars@model$representation_name)
> print(as.data.frame(X))

```

```

Arch1      Arch2      Arch3
1  0.70215790  0.07243783  0.15437558
2  0.71247885  0.06503881  0.15991126
3  1.02435370  0.06261124 -0.13989618
4  0.03202082  -0.57575088  1.60055217
5  -0.26444539  -0.09051645  1.42864559
6  0.15816143 -0.38871398  1.19466072
7  0.21369543  1.12529284 -0.35398127
8  0.52474104 -0.72363776  1.22207822
9  0.87206152 -0.16622314  0.33217225
10 0.74588811  0.24471634 -0.06286804
11 0.71814005  0.26654037 -0.07981878
12 0.35528158  0.55637148  0.04944270
13 0.38342268  0.65547801  0.06506703
14 0.33281056  0.55547862  0.03518461
15 -1.05315545 -0.15204939  2.21453898
16 -0.87636717  0.15204934  1.79459893
17 -0.48050470  0.55632915  1.20280683
18 1.28455006 -0.39612300  0.25412546
19 1.10982039 -0.54983319  0.51755064
20 1.38192559 -0.40140627  0.19558614
21 0.955562884  0.06507248 -0.07475759
22 -0.26325328 -0.15489231  1.37226578
23 -0.15762926 -0.06792138  1.16876121
24 0.25501405  1.21174521 -0.51697903
25 -0.55072956 -0.35159377  2.02066047
26 1.13122594 -0.28350109  0.17124468
27 0.96905862 -0.09463050  0.12999772
28 1.47816899  0.31969412 -0.70499833
29 0.47512505  1.45761448 -0.92465792
30 1.33969769  1.23270517 -1.67116221
31 1.45542785  2.09065358 -3.38764080
32 1.02969787  0.26193304 -0.35805575

```

Memory Reduction / Saving



Example 2: ML Acceleration

- About the dataset
 - R package “mlbench”
 - Multi-spectral scanner image data
 - 6k samples
 - x1 to x36: predictors
 - Classes:
 - 6 levels
 - Different type of soil
 - Use GLRM to compress predictors

```
> head(Satellite[sample(1:nrow(Satellite), 100),], 10)
   x.1 x.2 x.3 x.4 x.5 x.6 x.7 x.8 x.9 x.10 x.11 x.12 x.13 x.14 x.15 x.16 x.17 x.18 x.19 x.20 x.21
2474 64 112 128 103 64 116 122 99 64 121 122 96 67 111 123 100 67 116 123 100 71
4800 68 71 75 59 60 57 60 45 53 54 53 38 63 61 63 42 55 51 50 29 55
4165 63 96 108 89 66 96 112 89 66 100 112 92 67 99 110 94 67 103 114 94 71
5924 67 92 101 76 63 102 114 90 67 102 114 94 64 85 98 76 64 89 106 83 64
1327 59 56 89 70 59 60 80 63 66 63 76 66 63 57 86 72 59 57 82 68 59
6393 50 73 94 79 53 81 102 83 53 77 98 79 46 79 96 78 50 79 96 81 53
1838 72 81 82 65 76 81 82 65 72 85 86 68 71 83 83 63 76 79 79 67 71
2826 71 107 118 96 67 103 113 96 67 107 118 96 63 109 117 100 63 104 117 96 63
3397 55 79 93 75 51 75 93 75 51 79 96 79 56 83 96 85 56 83 108 85 56
1175 84 95 100 78 80 91 96 81 71 87 91 74 79 95 96 79 79 91 93 75 79
   x.22 x.23 x.24 x.25 x.26 x.27 x.28 x.29 x.30 x.31 x.32 x.33 x.34 x.35 x.36 classes
2474 111 128 100 66 113 122 100 66 113 127 100 66 113 122 100 red soil
4800 54 57 37 63 67 69 52 59 56 62 48 56 53 66 48 vegetation stubble
4165 103 114 98 67 106 114 94 70 106 119 94 70 106 119 94 red soil
5924 102 115 91 68 87 96 78 68 87 100 78 64 95 104 81 red soil
1327 60 82 68 57 56 82 73 57 53 85 76 57 56 82 65 vegetation stubble
6393 79 96 81 48 68 89 75 48 75 89 79 51 75 96 79 red soil
1838 83 87 63 71 79 85 62 71 79 85 62 71 79 85 67 very damp grey soil
2826 109 112 92 63 103 119 90 63 103 119 94 67 103 119 94 red soil
3397 83 100 81 56 81 93 83 56 88 101 83 56 88 105 83 red soil
1175 91 96 75 82 96 100 81 78 91 96 78 78 91 96 78 damp grey soil
```

Example 2: Use GLRM to Speed Up ML

```
# Initialise H2O Cluster
library(h2o)
h2o.init(nthreads = -1)

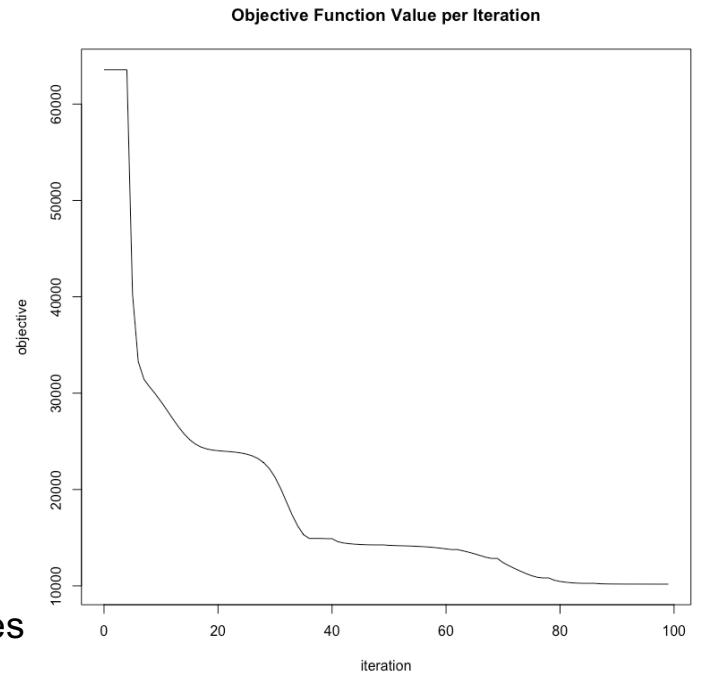
# Use Satellite data from "mlbench"
library(mlbench)
library(data.table)
data("Satellite")

hex_sat <- as.h2o(Satellite)

# Build a GLRM and compress data into a 10-column matrix
glrm_sat <- h2o.glrm(training_frame = hex_sat[, -37], # exclude response
                      k = 6, seed = 1234,
                      transform = "STANDARDIZE",
                      regularization_x = "Quadratic",
                      regularization_y = "L1",
                      max_iterations = 100)

# Check convergence
plot(glrm_sat)
```

k = 6
Reduce to 6 features



Example 2: Random Forest

- Train a vanilla H2O Random Forest model with ...
 - Full data set (36 predictors)
 - Compressed data set (6 predictors)

```
# Train DRF with Full Dataset
time_start <- proc.time()
model_drf_full <- h2o.randomForest(training_frame = hex_sat,
                                      x = features_full,
                                      y = "classes",
                                      ntrees = 500,
                                      nfolds = 10)
time_end_full <- timetaken(time_start)
model_drf_full

# Train DRF with Compressed Dataset
time_start <- proc.time()
model_drf_glrm <- h2o.randomForest(training_frame = hex_sat_glrm,
                                      x = features_glrm,
                                      y = "classes",
                                      ntrees = 500,
                                      nfolds = 10)
time_end_glrm <- timetaken(time_start)
model_drf_glrm
```

Example 2: Results Comparison

Data	Time	10-fold Cross Validation	
		Log Loss	Accuracy
Raw data (36 Predictors)	4 mins 26 sec	0.24553	91.80%
Data compressed with GLRM (6 Predictors)	1 min 24 sec	0.25792	90.59%

- Benefits of GLRM
 - Shorter training time
 - Quick insight before running models on full data set

Example 3: Clusters Visualisation

- About the dataset
 - Multi-spectral scanner image data
 - Same as example 2
 - x_1 to x_{36} : predictors
 - Use GLRM to compress predictors to 2D representation
 - Use 6 classes to colour clusters

```
> head(Satellite[sample(1:nrow(Satellite), 100),], 10)
   x.1 x.2 x.3 x.4 x.5 x.6 x.7 x.8 x.9 x.10 x.11 x.12 x.13 x.14 x.15 x.16 x.17 x.18 x.19 x.20 x.21
2474 64 112 128 103 64 116 122 99 64 121 122 96 67 111 123 100 67 116 123 100 71
4800 68 71 75 59 60 57 60 45 53 54 53 38 63 61 63 42 55 51 50 29 55
4165 63 96 108 89 66 96 112 89 66 100 112 92 67 99 110 94 67 103 114 94 71
5924 67 92 101 76 63 102 114 90 67 102 114 94 64 85 98 76 64 89 106 83 64
1327 59 56 80 70 59 60 80 63 66 63 76 66 63 57 86 72 59 57 82 68 59
6393 50 73 94 79 53 81 102 83 53 77 98 79 46 79 96 78 50 79 96 81 53
1838 72 81 82 65 76 81 82 65 72 85 86 68 71 83 83 63 76 79 79 67 71
2826 71 107 118 96 67 103 113 96 67 107 118 96 63 109 117 100 63 104 117 96 63
3397 55 79 93 75 51 75 93 75 51 79 96 79 56 83 96 85 56 83 108 85 56
1175 84 95 100 78 80 91 96 81 71 87 91 74 79 95 96 79 79 91 93 75 79
   x.22 x.23 x.24 x.25 x.26 x.27 x.28 x.29 x.30 x.31 x.32 x.33 x.34 x.35 x.36 classes
2474 111 128 100 66 113 122 100 66 113 127 100 66 113 122 100 red soil
4800 54 57 37 63 67 69 52 59 56 62 48 56 53 66 48 vegetation stubble
4165 103 114 98 67 106 114 94 70 106 119 94 70 106 119 94 red soil
5924 102 115 91 68 87 96 78 68 87 100 78 64 95 104 81 red soil
1327 60 82 68 57 56 82 73 57 53 85 76 57 56 82 65 vegetation stubble
6393 79 96 81 48 68 89 75 48 75 89 79 51 75 96 79 red soil
1838 83 87 63 71 79 85 62 71 79 85 62 71 79 85 67 very damp grey soil
2826 109 112 92 63 103 119 90 63 103 119 94 67 103 119 94 red soil
3397 83 100 81 56 81 93 83 56 88 101 83 56 88 105 83 red soil
1175 91 96 75 82 96 100 81 78 91 96 78 78 91 96 78 damp grey soil
```

Example 3: Clusters Visualisation

```
# Initialise H2O Cluster
library(h2o)
h2o.init(nthreads = -1)

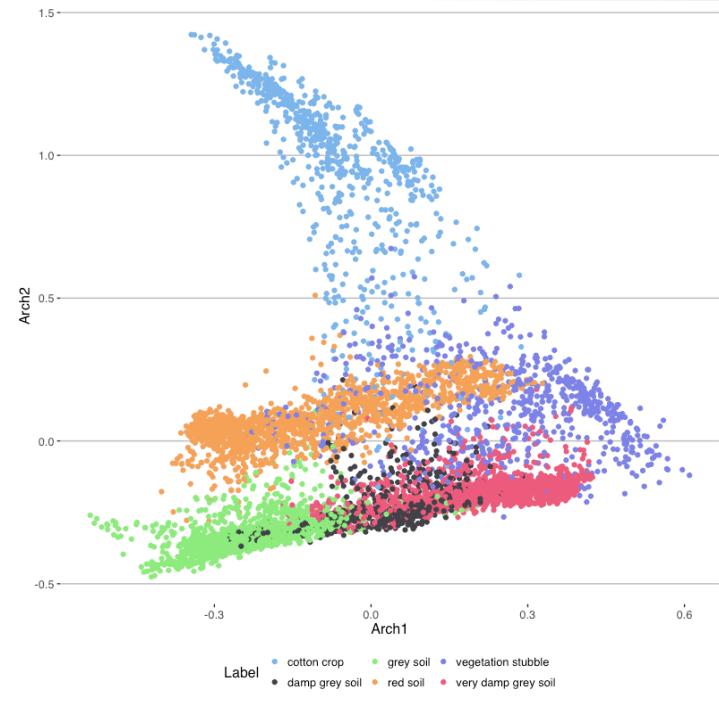
# Use Satellite data from "mlbench"
library(mlbench)
data("Satellite")
hex_sat <- as.h2o(Satellite)

# Build a GLRM and compress data into a 10-column matrix
glrm_sat <- h2o.glrm(training_frame = hex_sat[, -37], # exclude response
                      k = 2, seed = 1234,
                      transform = "STANDARDIZE",
                      regularization_x = "Quadratic",
                      regularization_y = "L1",
                      max_iterations = 100)

# Extract X from GLRM
X <- as.data.frame(h2o.getFrame(glrm_sat@model$representation_name))

# Create data frame for data visualisation
d_sat <- data.frame(X, as.data.frame(hex_sat[, 37]))
colnames(d_sat) <- c("Arch1", "Arch2", "Label")
d_sat$Label <- as.factor(d_sat$Label)

# Visualise
library(ggplot2)
library(ggthemes)
ggplot(d_sat, aes(Arch1, Arch2, colour = Label)) +
  geom_point() + theme_hc() + scale_colour_hc()
  ggtitle("Clusters in Satellite Dataset")
```



Example 4: Imputation

"mtcars" – same dataset for example 1

> mtcars											
	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
Fiat X-1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2



Randomly introduce 50% missing values

> mtcars_with_na											
	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	NA	110	3.90	NA	16.46	0	NA	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	NA	NA	NA	4	4
Datsun 710	22.8	4	108.0	NA	3.85	NA	NA	NA	1	NA	1
Hornet 4 Drive	21.4	NA	258.0	110	3.08	NA	NA	1	NA	NA	1
Hornet Sportabout	18.7	NA	NA	NA	NA	NA	17.02	0	0	NA	2
Valiant	NA	6	NA	105	2.76	3.460	20.22	1	NA	NA	1
Duster 360	14.3	8	NA	245	NA	3.570	15.84	NA	NA	NA	NA
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	NA	NA	4	NA
Merc 230	NA	4	NA	NA	3.92	NA	22.90	1	0	4	2
Merc 280	NA	6	NA	123	NA	NA	NA	1	NA	NA	NA
Merc 280C	NA	6	167.6	NA	NA	NA	18.90	NA	NA	4	4
Merc 450SE	NA	8	275.8	NA	NA	4.070	NA	0	NA	NA	NA
Merc 450SL	NA	NA	275.8	3.07	3.730	NA	0	NA	3	3	3
Merc 450SLC	NA	8	NA	180	3.780	NA	0	0	3	3	3
Cadillac Fleetwood	10.4	8	NA	NA	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	NA	NA	5.424	17.82	0	NA	NA	NA
Chrysler Imperial	14.7	8	440.0	230	3.23	NA	17.42	0	NA	3	NA
Fiat 128	32.4	4	78.7	66	NA	2.200	NA	NA	NA	4	1
Honda Civic	30.4	4	NA	NA	4.93	NA	18.52	1	1	4	NA
Toyota Corolla	NA	4	NA	65	4.22	1.835	NA	1	NA	4	NA
Toyota Corona	21.5	4	120.1	NA	NA	2.465	20.01	NA	0	NA	NA
Dodge Challenger	15.5	8	318.0	150	2.76	NA	16.87	0	0	3	NA
AMC Javelin	15.2	NA	304.0	150	NA	NA	17.30	0	0	NA	2
Camaro Z28	NA	8	350.0	NA	3.73	3.840	NA	0	0	NA	NA
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	NA	0	0	3	2
Fiat X-1-9	27.3	4	79.0	NA	4.08	1.935	NA	1	1	NA	1
Porsche 914-2	NA	4	120.3	91	4.43	NA	NA	0	NA	5	2
Lotus Europa	NA	4	95.1	113	3.77	1.513	16.90	NA	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
Ferrari Dino	NA	6	145.0	175	NA	NA	NA	NA	5	6	6
Maserati Bora	NA	NA	NA	NA	NA	3.570	NA	0	NA	5	8
Volvo 142E	21.4	4	121.0	109	4.11	NA	NA	NA	NA	4	NA

Example 4: GLRM with NAs

When we reconstruct the table using GLRM,
missing values are automatically imputed.

```
# Convert to H2O Data Frame
hex_mtcars_na ← as.h2o(mtcars_with_na)

# Build a GLRM using data table with missing values
glrm_mtcars ← h2o.glrm(training_frame = hex_mtcars_na,
                        k = 10, seed = 1234,
                        init = "SVD",
                        regularization_x = "Quadratic",
                        regularization_y = "Quadratic",
                        max_iterations = 100)

# Use GLRM to impute missing values
mtcars_recon ← h2o.predict(glrm_mtcars, hex_mtcars_na)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	222.2	110.0	3.9	2.7	16.4	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110.0	3.9	2.9	17.7	0	1	4	4
Datsun 710	22.8	4	108.0	149.3	3.8	2.4	18.6	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110.0	3.1	3.2	17.7	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175.0	3.2	3.2	17.0	0	0	3	2
Valiant	19.7	6	222.2	105.0	2.7	3.5	20.2	1	0	4	1
Duster 360	14.3	8	222.2	245.0	3.2	3.6	15.8	0	0	3	4
Merc 240D	24.4	4	146.7	62.0	3.7	3.1	20.0	0	1	4	2
Merc 230	22.8	4	140.8	95.0	3.9	3.1	22.9	1	0	4	2
Merc 280	19.2	6	222.2	123.0	3.9	3.1	17.7	1	0	4	3
Merc 280C	19.7	6	167.6	149.3	3.9	3.1	18.9	1	0	4	4
Merc 450SE	16.4	8	275.8	180.0	3.1	4.1	17.4	0	0	3	3
Merc 450SL	17.3	8	275.8	180.0	3.1	3.7	17.6	0	0	3	3
Merc 450SLC	15.2	8	275.8	180.0	3.1	3.8	18.0	0	0	3	3
Cadillac Fleetwood	10.4	8	222.2	205.0	2.9	5.2	18.0	0	0	3	4
Lincoln Continental	10.4	8	460.0	215.0	3.0	5.5	17.8	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230.0	3.2	5.3	17.4	0	0	3	3
Fiat 128	32.4	4	78.7	66.0	4.1	2.2	19.5	0	1	4	1
Honda Civic	30.4	4	75.7	52.0	4.9	1.7	18.5	1	1	4	2
Toyota Corolla	33.9	4	71.1	65.0	4.2	1.8	17.7	1	1	4	1
Toyota Corona	21.5	4	120.1	97.0	3.7	2.5	20.0	1	0	3	1
Dodge Challenger	15.5	8	318.0	150.0	2.8	3.2	16.9	0	0	3	2
AMC Javelin	15.2	8	304.0	150.0	3.1	3.2	17.3	0	0	3	2
Camaro Z28	13.3	8	350.0	245.0	3.7	3.8	15.4	0	0	3	4
Pontiac Firebird	19.2	8	400.0	175.0	3.1	3.8	17.1	0	0	3	2
Fiat X1-9	27.3	4	79.0	66.0	4.1	2.0	18.9	1	1	4	1
Porsche 914-2	19.7	4	120.3	91.0	4.5	2.1	16.7	0	1	5	2
Lotus Europa	19.6	4	95.1	113.0	3.8	1.5	16.9	0	1	5	2
Ford Pantera L	15.8	8	351.0	264.0	4.2	3.2	14.5	0	1	5	4
Ferrari Dino	19.7	6	145.0	175.0	3.6	2.7	15.5	0	1	5	6
Maserati Bora	15.0	8	301.0	335.0	3.5	3.6	14.6	0	1	5	8
Volvo 142E	21.4	4	121.0	109.0	4.1	2.8	18.6	1	1	4	2

Example 4: Results Comparison

- We are asking GLRM to do a difficult job
 - 50% missing values
 - Imputation results look reasonable

Absolute difference between original and imputed values.

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	0.0	0	62.2	0	0	0.0	0.0	0	0	0	0
Mazda RX4 Wag	0.0	0	0.0	0	0	0.0	0.7	0	0	0	0
Datsun 710	0.0	0	0.0	56	0	0.1	0.0	0	0	0	0
Hornet 4 Drive	0.0	0	0.0	0	0	0.0	-1.7	0	0	0	0
Hornet Sportabout	0.0	0	0.0	0	0	-0.3	0.0	0	0	0	0
Valiant	1.6	0	-2.8	0	0	0.0	0.0	0	0	1	0
Duster 360	0.0	0	-137.8	0	0	0.0	0.0	0	0	0	0
Merc 240D	0.0	0	0.0	0	0	0.0	0.0	-1	1	0	0
Merc 230	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Merc 280	0.0	0	54.6	0	0	-0.3	-0.6	0	0	0	-1
Merc 280C	1.9	0	0.0	26	0	-0.3	0.0	0	0	0	0
Merc 450SE	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Merc 450SL	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Merc 450SLC	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Cadillac Fleetwood	0.0	0	-249.8	0	0	0.0	0.0	0	0	0	0
Lincoln Continental	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Chrysler Imperial	0.0	0	0.0	0	0	0.0	0.0	0	0	0	-1
Fiat 128	0.0	0	0.0	0	0	0.0	0.0	-1	0	0	0
Honda Civic	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Toyota Corolla	0.0	0	0.0	0	0	0.0	-2.2	0	0	0	0
Toyota Corona	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Dodge Challenger	0.0	0	0.0	0	0	-0.4	0.0	0	0	0	0
AMC Javelin	0.0	0	0.0	0	0	-0.3	0.0	0	0	0	0
Camaro Z28	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Pontiac Firebird	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Fiat X1-9	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Porsche 914-2	-6.3	0	0.0	0	0	0.0	0.0	0	0	0	0
Lotus Europa	-10.8	0	0.0	0	0	0.0	0.0	-1	0	0	0
Ford Pantera L	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Ferrari Dino	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Maserati Bora	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0
Volvo 142E	0.0	0	0.0	0	0	0.0	0.0	0	0	0	0

Conclusions

- Use GLRM to
 - Save memory
 - Speed up machine learning
 - Visualise clusters
 - Impute missing values
- A great tool for data pre-processing
 - Include it in your data pipeline

Any Questions?

- Contact
 - joe@h2o.ai
 - [@matlabulous](https://twitter.com/matlabulous)
 - github.com/woobe
- Slides & Code
 - github.com/h2oai/h2o-meetups
- H2O in London
 - Meetups / Office (soon)
 - www.h2o.ai/careers
- H2O Help Docs & Tutorials
 - www.h2o.ai/docs
 - university.h2o.ai