

Jakub Háva  
[jakub@h2o.ai](mailto:jakub@h2o.ai)

# Introduction to Spark

Firemní semináře @ MFF UK, Praha  
October 24, 2018

**Spark**<sup>★</sup> + **H<sub>2</sub>O**

---

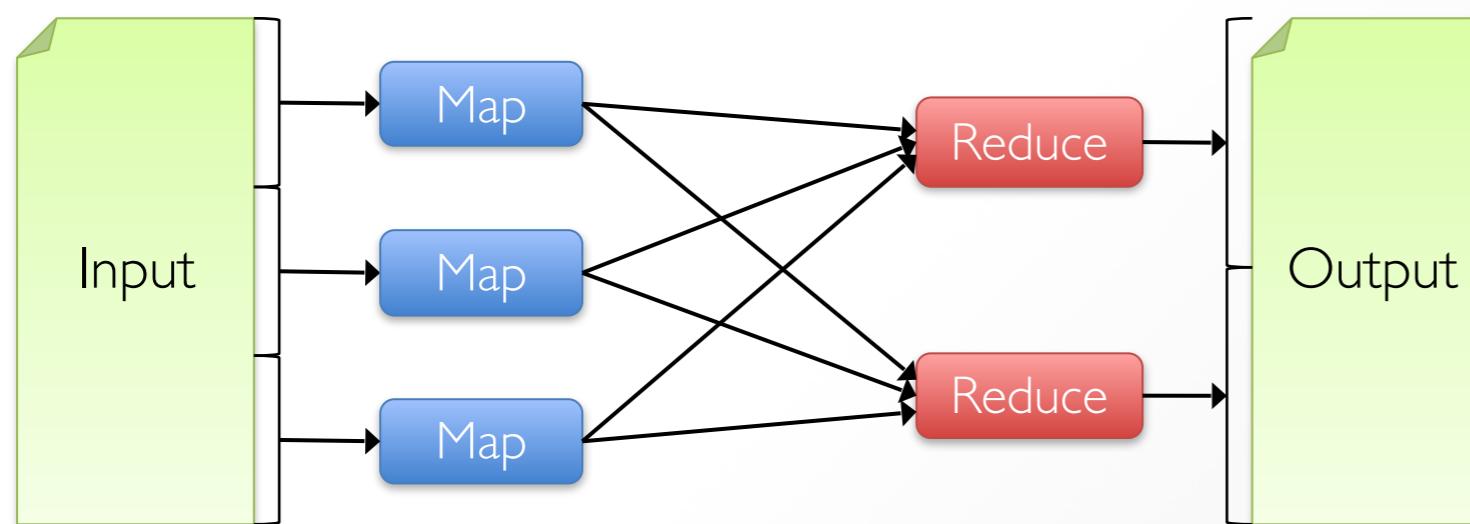
**SPARKLING  
WATER**

# **Who are we?**

- **Kuba**
  - Senior Software engineer at H2O.ai - Code Owner - Sparkling Water
  - Master's at Charles University (CZ)
  - Implemented high-performance cluster monitoring tool for JVM based languages (JNI, JVMTI, instrumentation)
- **Michal**
  - VP of Engineering at H2O.ai
  - Author of Sparkling Water
  - Ph.D at Charles University (CZ), PostDoc at Purdue (US)

# Data Flow in typical Map-Reduce

- Acyclic data flow, from storage A to storage B
- Benefit is that runtime can decide where to run tasks and can automatically recover in case of their failures



# Inefficiencies of acyclic Data Flow

- Inefficient for application that repeatedly reuse a *working set* of data
  - Iterative algorithms - machine learning
  - Interactive data mining tools ( R, Scala interpreter)
- With typical MapReduce, the data are reloaded for each query

# Solution: RDD

# **RDD: Resilient Distributed Dataset**

- Allow apps to keep working sets in memory for efficient reuse
- Whilst keep the benefits of MapReduce
  - Fault tolerance, data locality, scalability
- General approach

# Spark Programming Model

# **Spark Runs on**

- YARN, Mesos, Standalone cluster
- 1 driver node
- Several executor nodes
  - can be started dynamically when needed based on provided resources by resource manager

# RDDs

- Immutable, partitioned collections of objects
- Created through parallel transformations
  - **map, filter, groupBy, join...**
  - Can be cached for efficient use on various levels
- Transformations does not start any computations - lazy behaviour
- Actions start the computations
  - **count, reduce, collect, save**
- RDDs are split across multiple partitions on multiple nodes

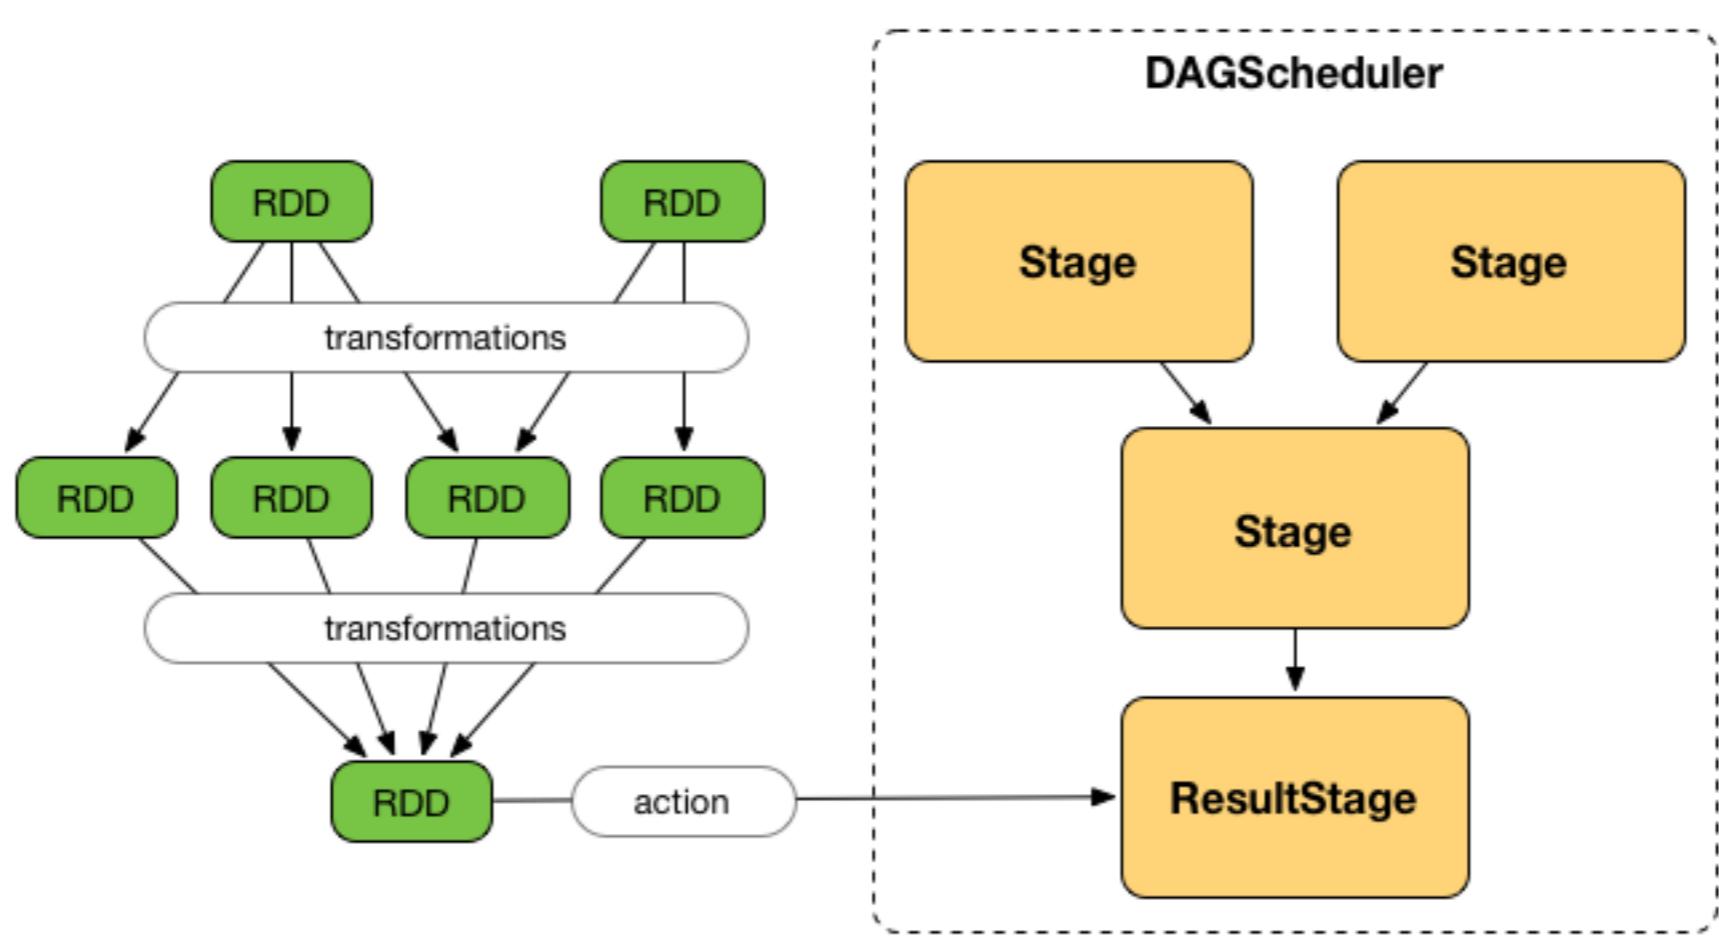
# RDD lineage

- Result of applying transformations on the RDD
- `val rdd = sc.textFile(...)`
- `val filtered = rdd.filter(...)`
- **toDebugString** method on RDD - see the lineage
- Plan describing the computation

# Spark Job

- Starts with an action called on RDD
- Job is split into multiple stages ( depends on type of transformations - **stage boundaries**)
- Stage is executed by DAG Scheduler
- Each stage is split into several tasks ( usually #partitions in the stage )

# DAG Scheduler



# Abstractions on top of RDD

# RDD

- Low-level transformations
- Data is unstructured
- Good when we want to manipulate with data using functional programming constructs
- Don't have schema of the data
- On heap, thus affected by GC and are using Java Serialisation

# RDD Example

```
rdd.filter(_.age > 21)          // transformation
  .map(_.last)                  // transformation
  .saveAsObjectFile("under21.bin") // action

rdd.filter(_.age > 21)
```

# Dataframes

- Build on top of RDDs => immutable distributed collection of data
- Data organised into columns like a table in relational database
- Have a schema
- DSL and SQL like specific language
- Better performance
  - Can use the schema and a lot of optimizations - Spark Query Optimiser
- Can be stored off-heap, Spark specific serialisation

# Datarame example

```
df.filter("age > 21")
```

```
df.filter(df.col("age").gt(21))
```

# Dataset

- First appeared in Preview of Spark 1.6
- Attempt to bring the best from RDDs and DataFrames
  - Compile type safety as in RDD
  - Performance as on DataFrames
- Off-heap storage ( outside the JVM memory )
- Usage of encoders to translate between JVM representation and Spark internal binary format, no need to de-serialise the object in order to read it

# Dataset example

```
val sc = new SparkContext(conf)
val sqlContext = new SQLContext(sc)
import sqlContext.implicits._
val sampleData: Seq[ScalaPerson] = ScalaData.sampleData()
val dataset = sqlContext.createDataset(sampleData)

dataset.filter(_.age < 21);
```

---

# Introduction to Scalable & Automatic Machine Learning with H<sub>2</sub>O

H<sub>2</sub>O.ai



# Agenda

- Talk 1: Introduction to H<sub>2</sub>O
  - Company and People
  - H<sub>2</sub>O Open Source ML Platform
  - Live Demos
- Talk 2: The Next Generation of Machine Learning on Apache Spark
  - Introduction to Sparkling Water
  - Live Demos
- Talk 3: Introduction to Driverless AI
  - Introduction to Driverless AI



# Company Overview

<b>Founded</b>	2012, Series C in Nov, 2017
<b>Products</b>	<ul style="list-style-type: none"><li>• Driverless AI – Automated Machine Learning</li><li>• H<sub>2</sub>O - Open Source Machine Learning Platform</li><li>• H2O4GPU - Lightning Fast Machine Learning on GPUs</li><li>• Sparkling Water - Integration of H<sub>2</sub>O and Apache Spark</li></ul>
<b>Mission</b>	Democratize AI. Do Good.
<b>Team</b>	<p>~100 employees</p> <ul style="list-style-type: none"><li>• Distributed Systems Engineers doing Machine Learning</li><li>• World-class visualization designers</li></ul>
<b>Offices</b>	Mountain View, London, Prague



H<sub>2</sub>O.ai

# Our Mission: Make Machine Learning Accessible to Everyone



Complexity is your enemy. Any fool can make something complicated. It is hard to keep things simple.

— *Richard Branson* —

AZ QUOTES

# Scientific Advisory Council



## Dr. Trevor Hastie

- John A. Overdeck Professor of Mathematics, Stanford University
- PhD in Statistics, Stanford University
- Co-author, *The Elements of Statistical Learning: Prediction, Inference and Data Mining*
- Co-author with John Chambers, *Statistical Models in S*
- Co-author, *Generalized Additive Models*



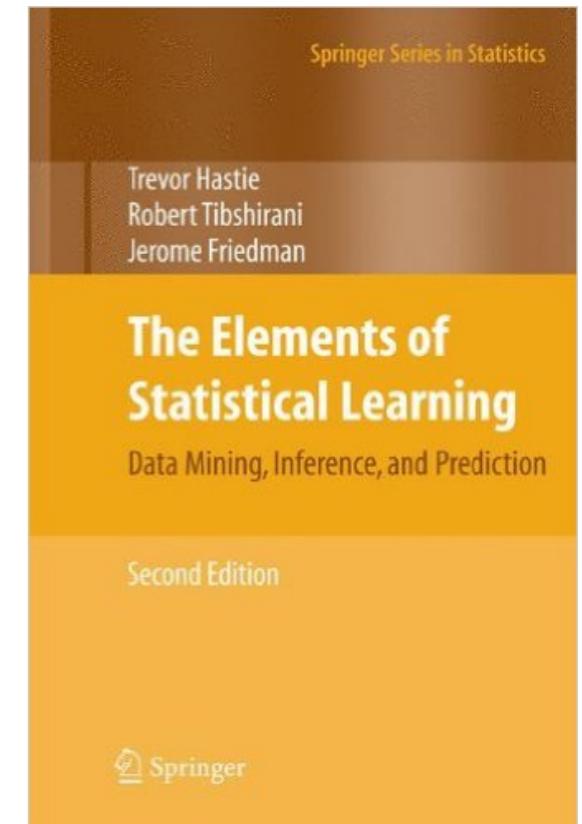
## Dr. Robert Tibshirani

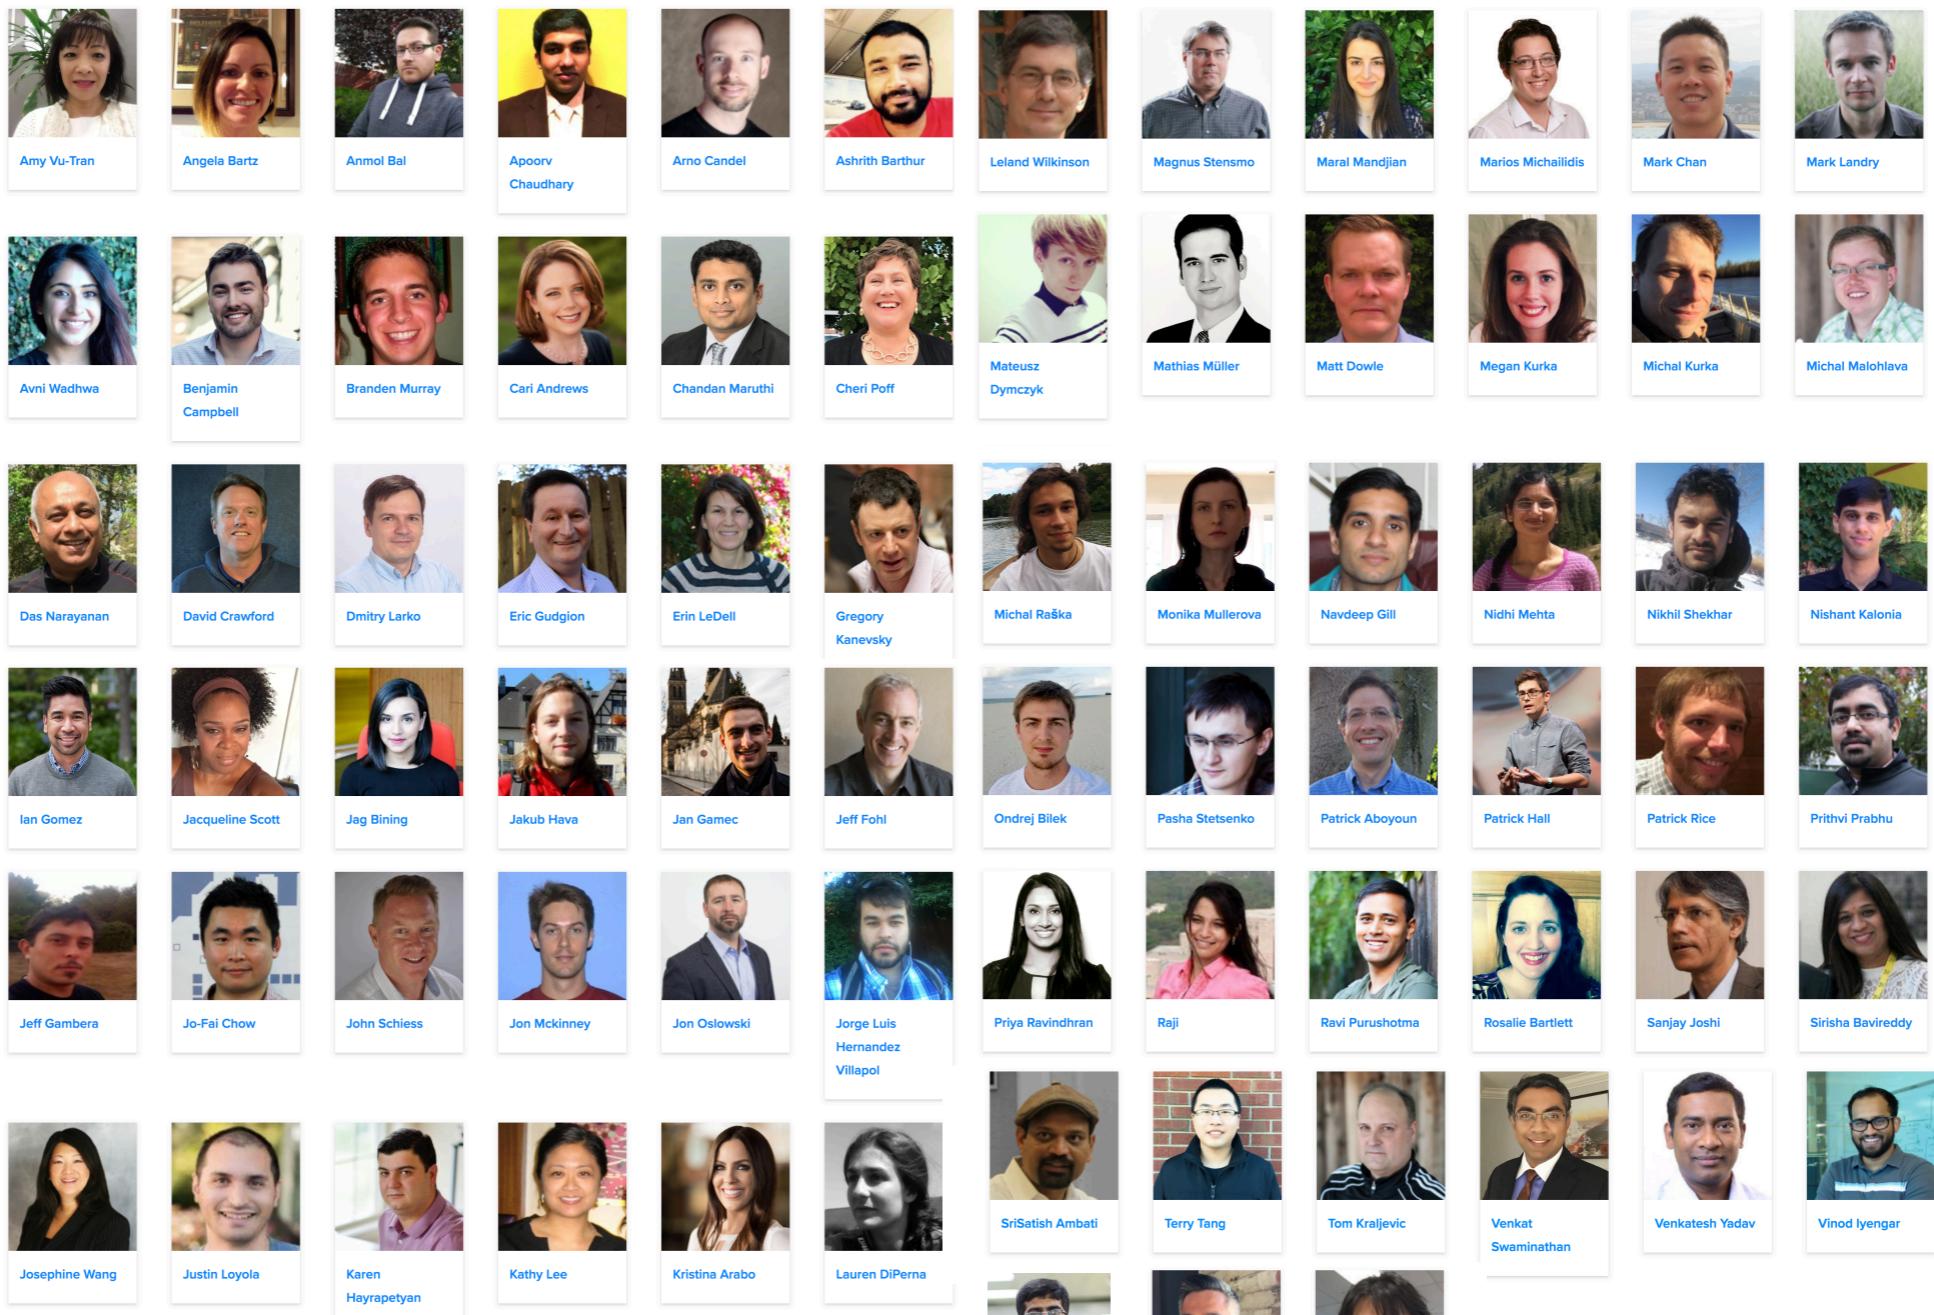
- Professor of Statistics and Health Research and Policy, Stanford University
- PhD in Statistics, Stanford University
- Co-author, *The Elements of Statistical Learning: Prediction, Inference and Data Mining*
- Author, *Regression Shrinkage and Selection via the Lasso*
- Co-author, *An Introduction to the Bootstrap*



## Dr. Steven Boyd

- Professor of Electrical Engineering and Computer Science, Stanford University
- PhD in Electrical Engineering and Computer Science, UC Berkeley
- Co-author, *Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers*
- Co-author, *Linear Matrix Inequalities in System and Control Theory*
- Co-author, *Convex Optimization*





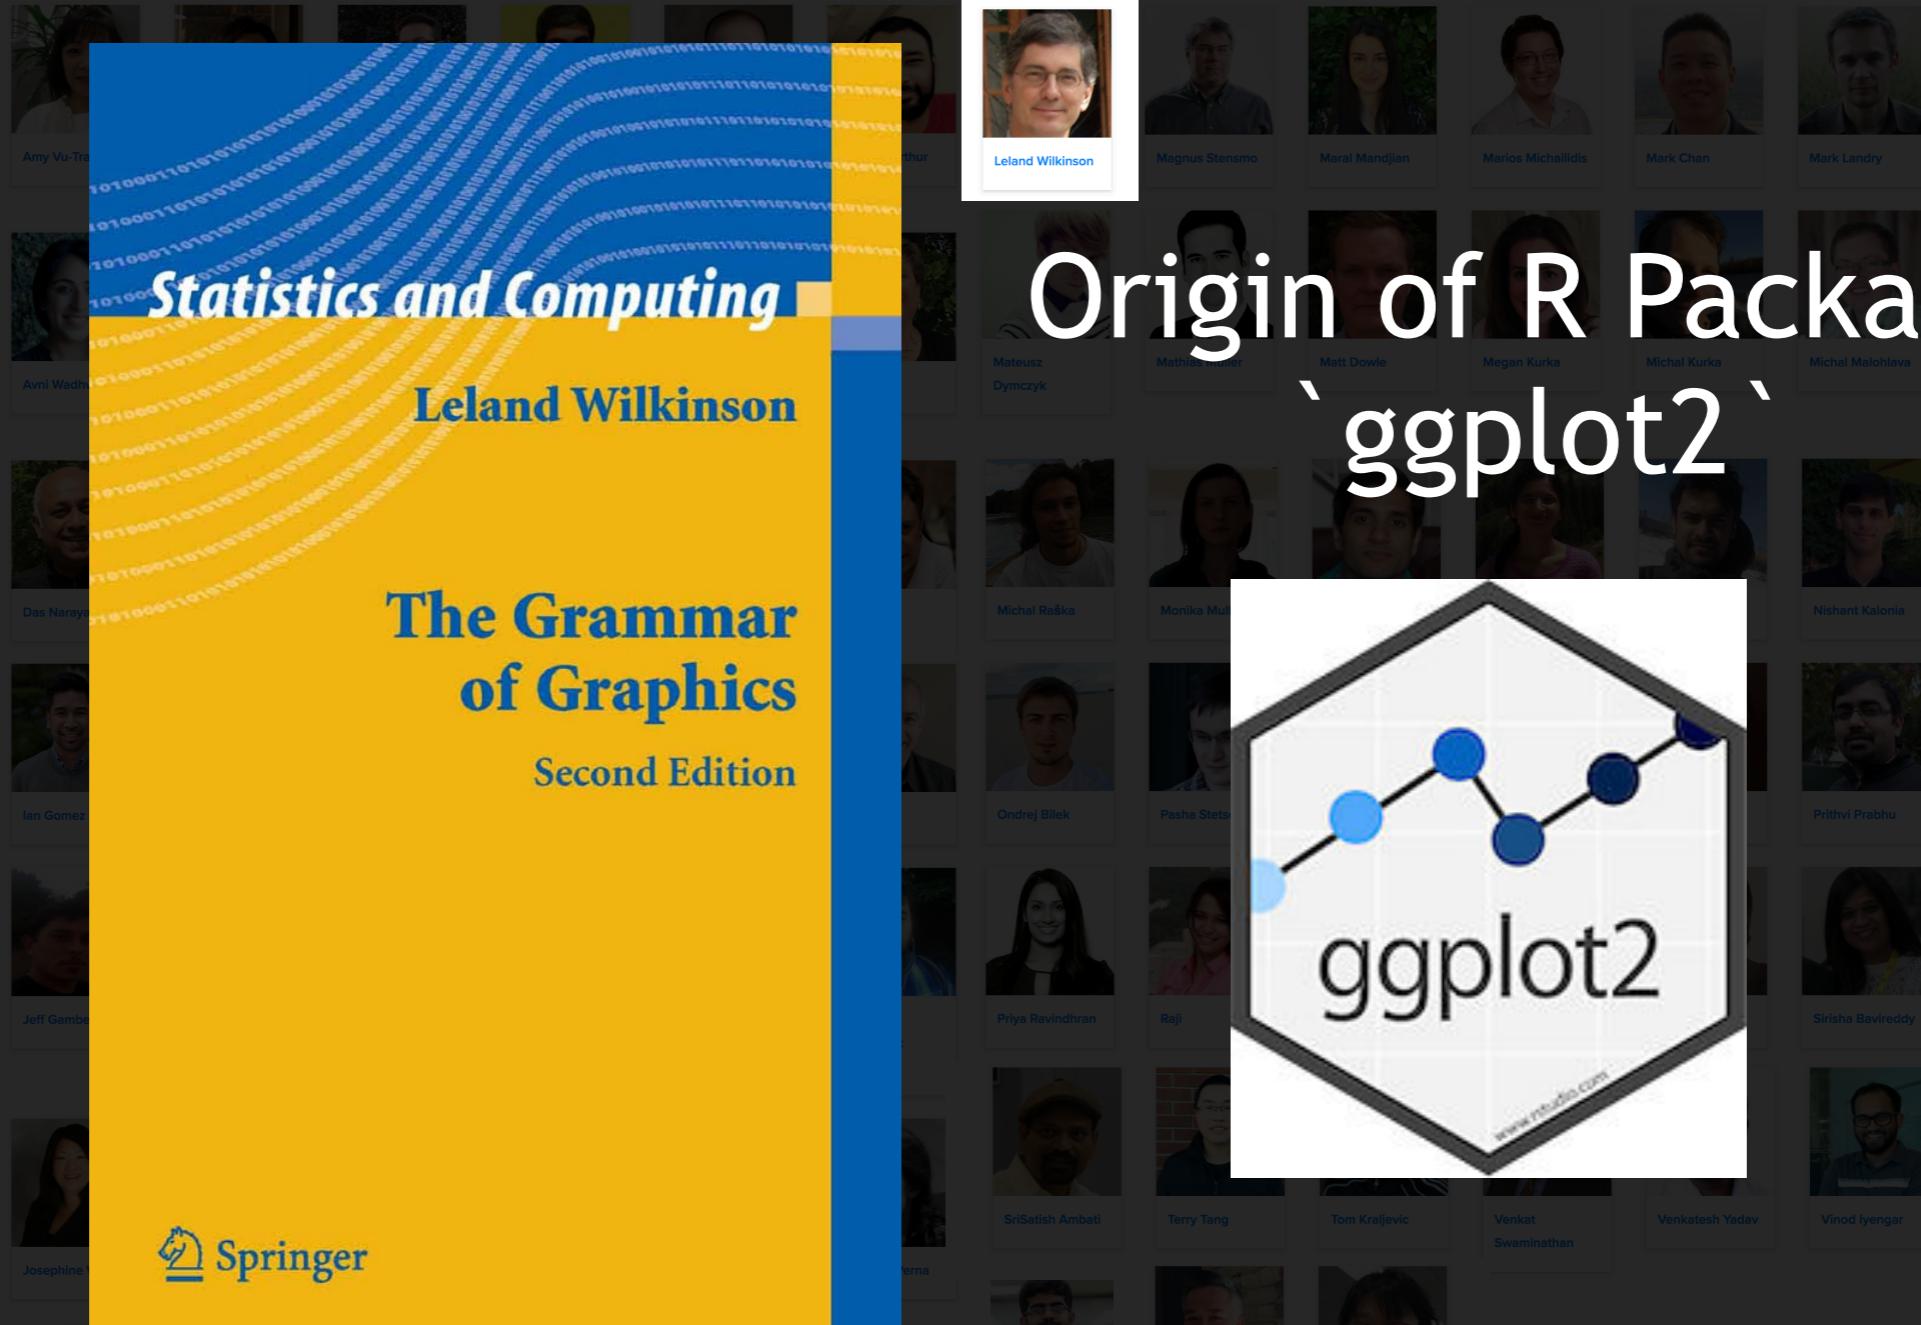
# H<sub>2</sub>O Team

# H<sub>2</sub>O Team

26

H<sub>2</sub>O.ai





H<sub>2</sub>O team



# Matt Dowle

# H<sub>2</sub>O Team

# H<sub>2</sub>O Products



In-Memory, Distributed Machine Learning Algorithms with H2O Flow GUI



H2O AI Open Source Engine Integration with Spark



Lightning Fast machine learning on GPUs

DRIVERLESSAI

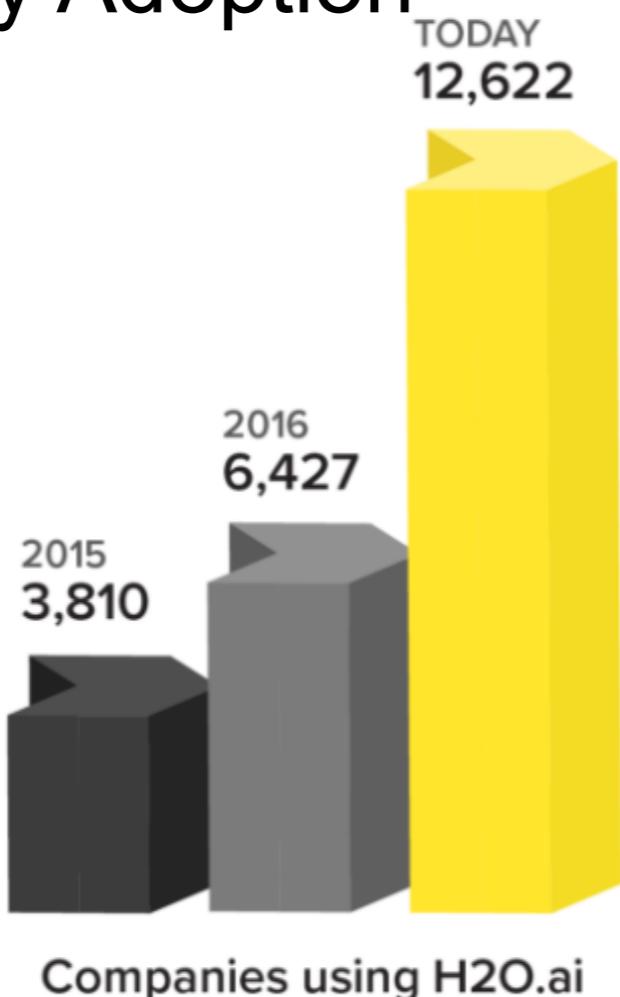
Automatic feature engineering, machine learning and interpretability

# Steam

Secure multi-tenant H2O clusters

**H<sub>2</sub>O.ai**

## Worldwide Community Adoption



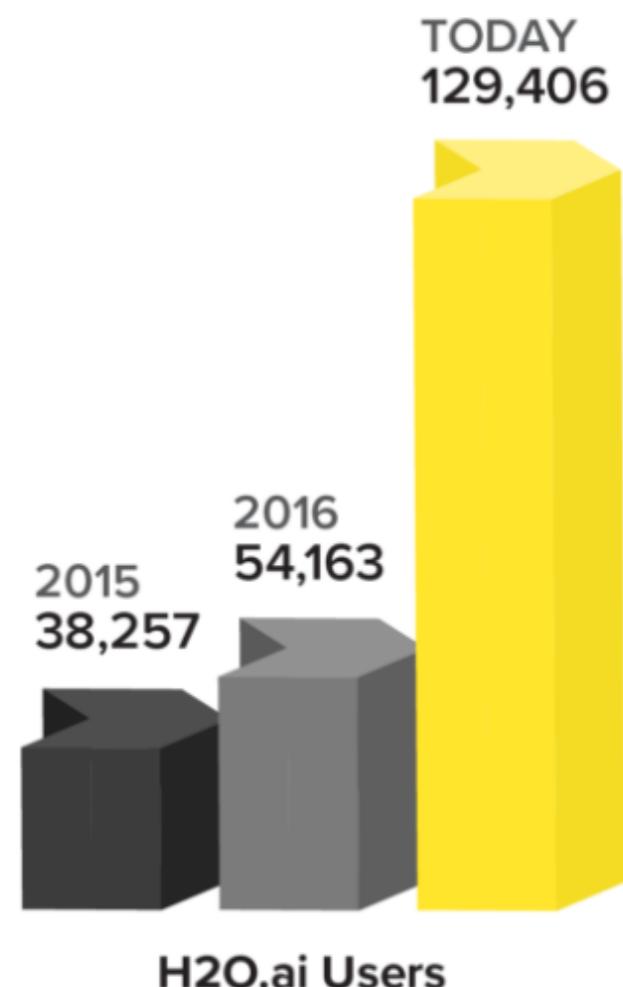
\* DATA FROM GOOGLE ANALYTICS EMBEDDED IN THE END USER PRODUCT

**222** OF **FORTUNE**  
**500**  
H<sub>2</sub>O

**8** OF TOP 10  
**BANKS**

**7** OF TOP 10  
**INSURANCE COMPANIES**

**4** OF TOP 10  
**HEALTHCARE COMPANIES**



# H2O.ai Solution Leadership Across Verticals



# Gartner names H2O as **Leader** with the **most completeness of vision**

- H2O.ai recognized as a **technology leader with most completeness of vision**
- H2O.ai was recognized for the mindshare, partner network and status as a **quasi-industry standard** for machine learning and AI.
- **H2O customers gave the highest overall score** among all the vendors for sales relationship and account management, customer support (onboarding, troubleshooting, etc.) and overall service and support.

Figure 1. Magic Quadrant for Data Science and Machine-Learning Platforms



# Platforms with H<sub>2</sub>O integration



srisatish  
@srisatish

Replies to @BobMuenchen @knime @h2oai

@KNIME gained the ability to run @H2O.ai algorithms, so these two may be viewed as complementary, not competitors  
#Ecosystem #OpenSource

3:32 PM - 2 Mar 2018

H<sub>2</sub>O + KNIME Talk  
at KNIME Summit  
Mar 2017



1:54 PM - 7 Mar 2018 from Hotel Berlin

Following

Figure 1. Magic Quadrant for Data Science and Machine-Learning Platforms



Source: Gartner (February 2018)

H<sub>2</sub>O.ai

# Community Expansion



88,286 members    43 interested    52 Meetups    48 cities    18 countries

Find out more: [www.h2o.ai/community/](http://www.h2o.ai/community/)

---

# H<sub>2</sub>O Machine Learning Platform

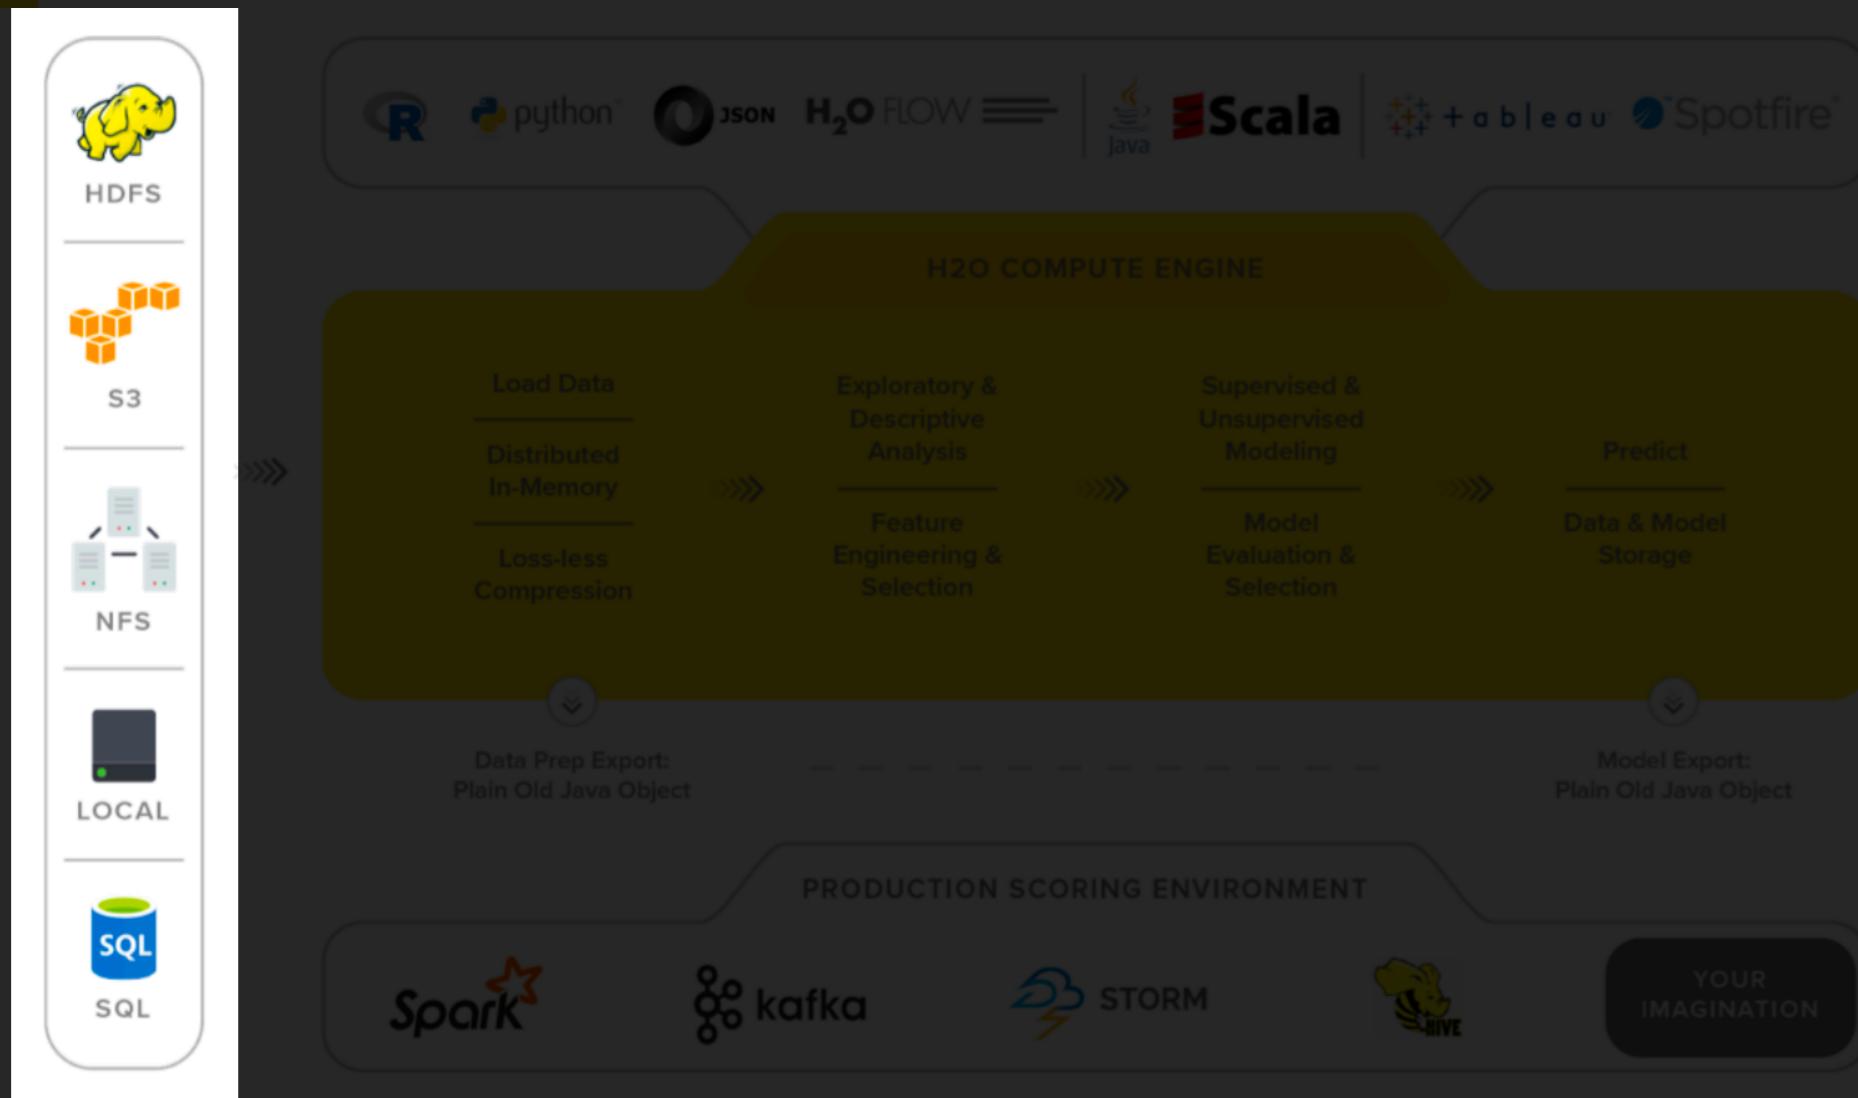


# High Level Architecture

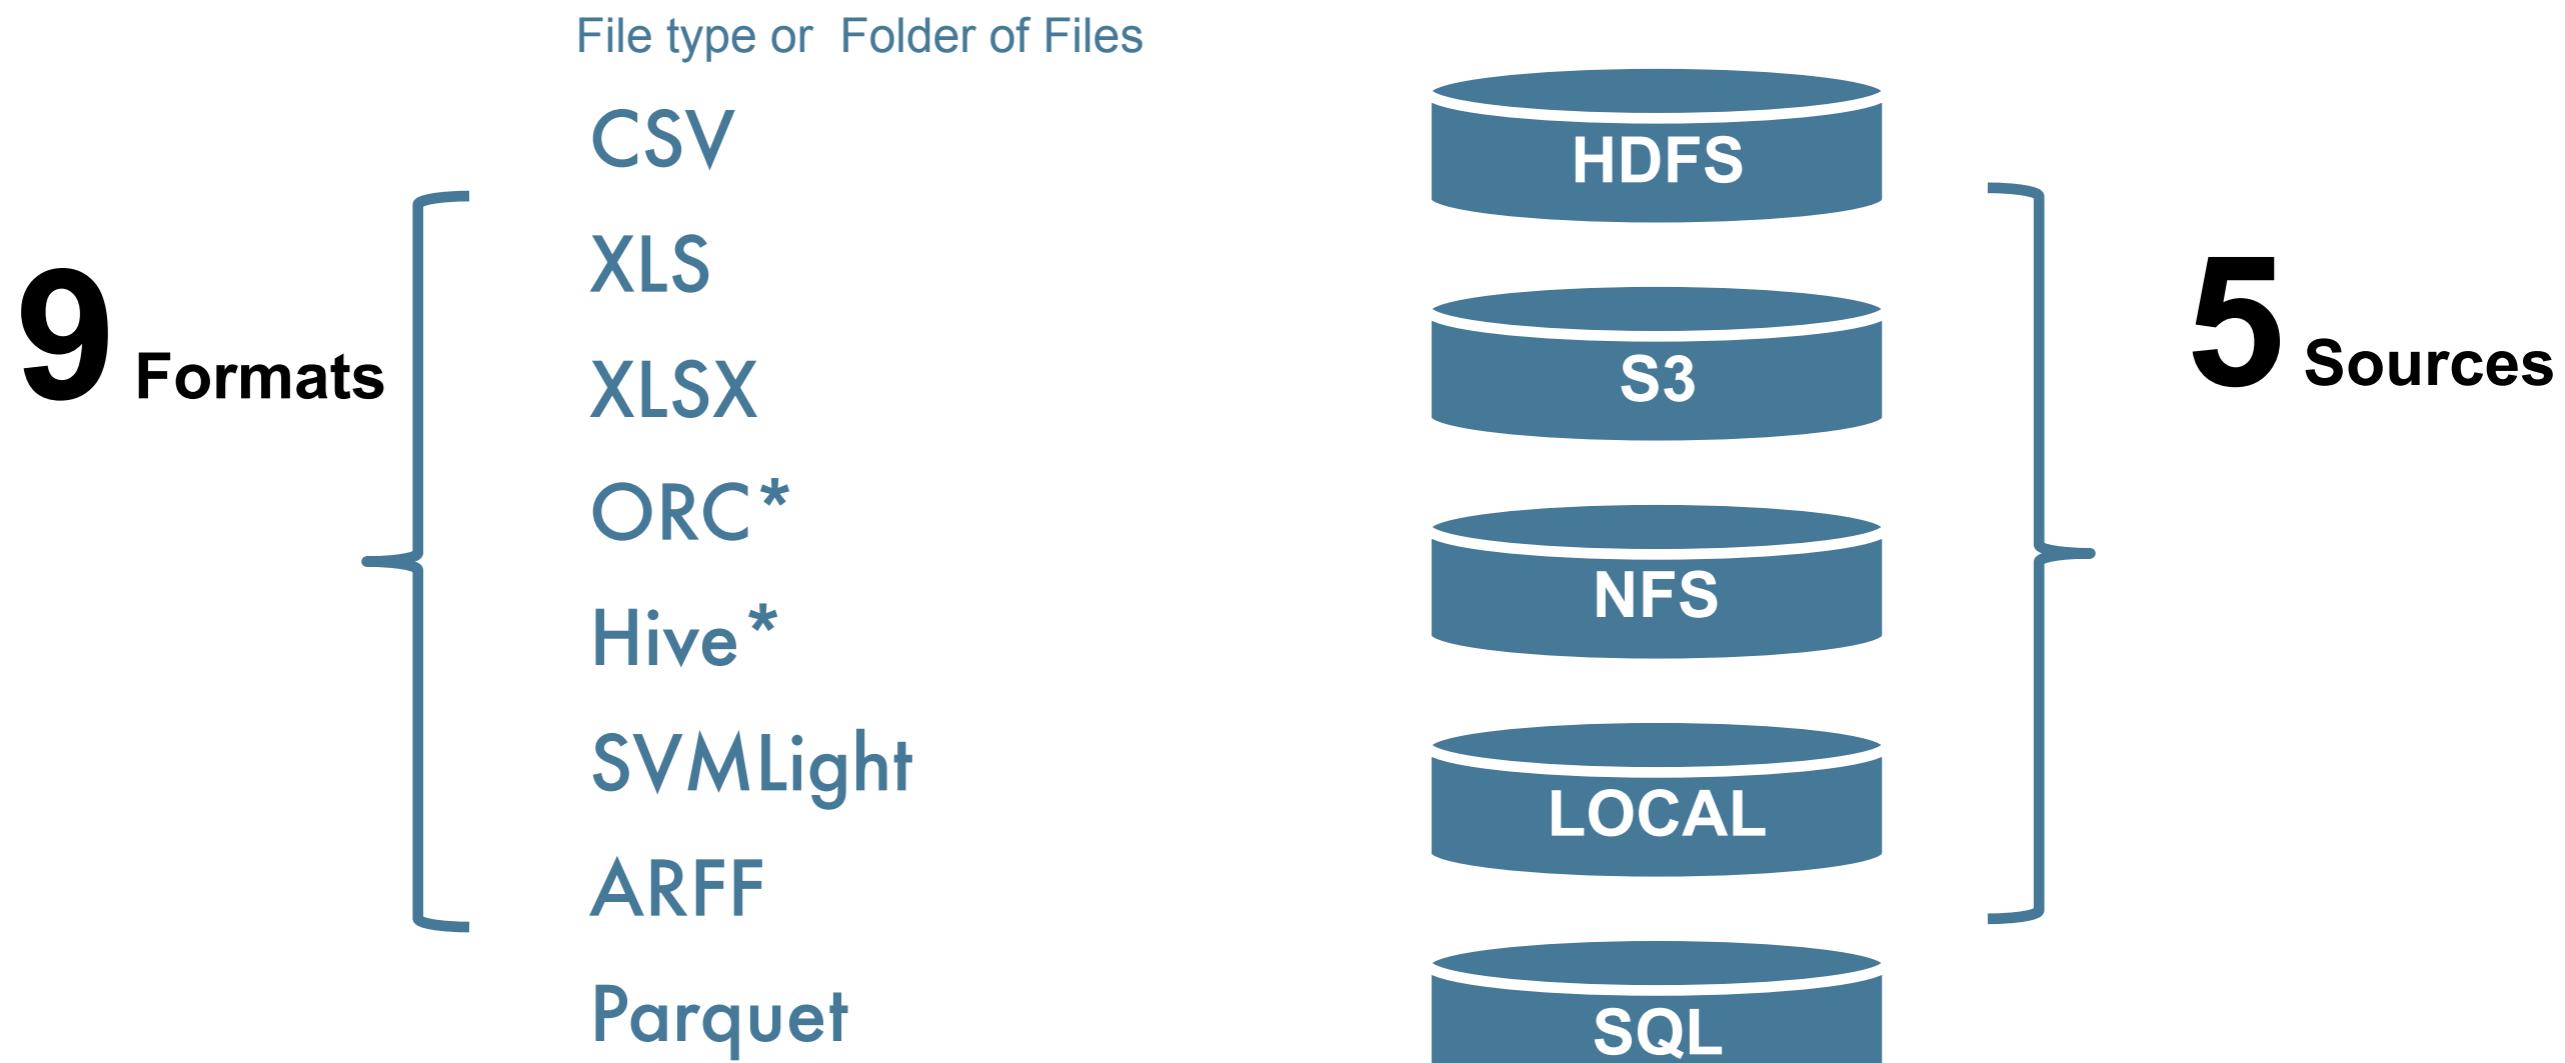


# High Level Architecture

Import Data from Multiple Sources



# Supported Formats & Data Sources



\* 1. only if H2O is running as a [Apache 1.8.0\\*](#)

\* 2. Hive files that are saved in ORC format

\* 3. without multi-file parsing or column type modification

# High Level Architecture

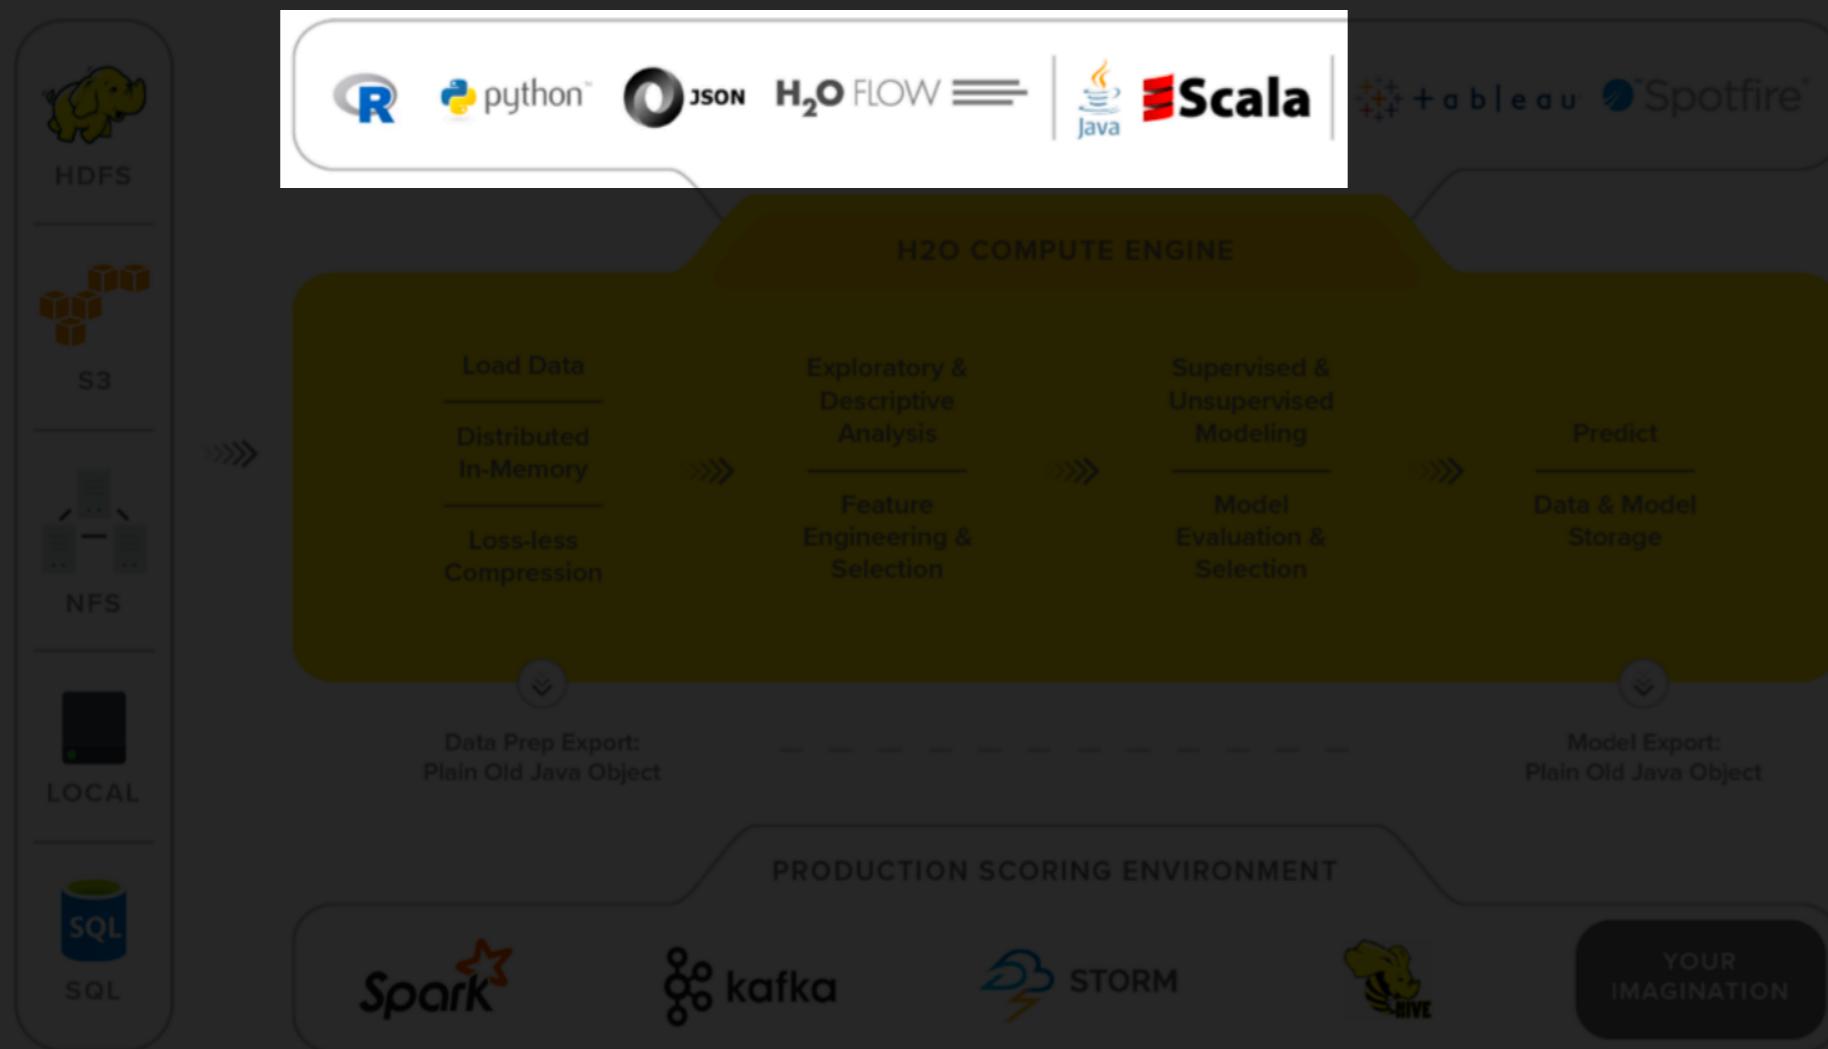
Fast, Scalable &  
Distributed  
Compute Engine  
Written in Java



# Multiple Interfaces

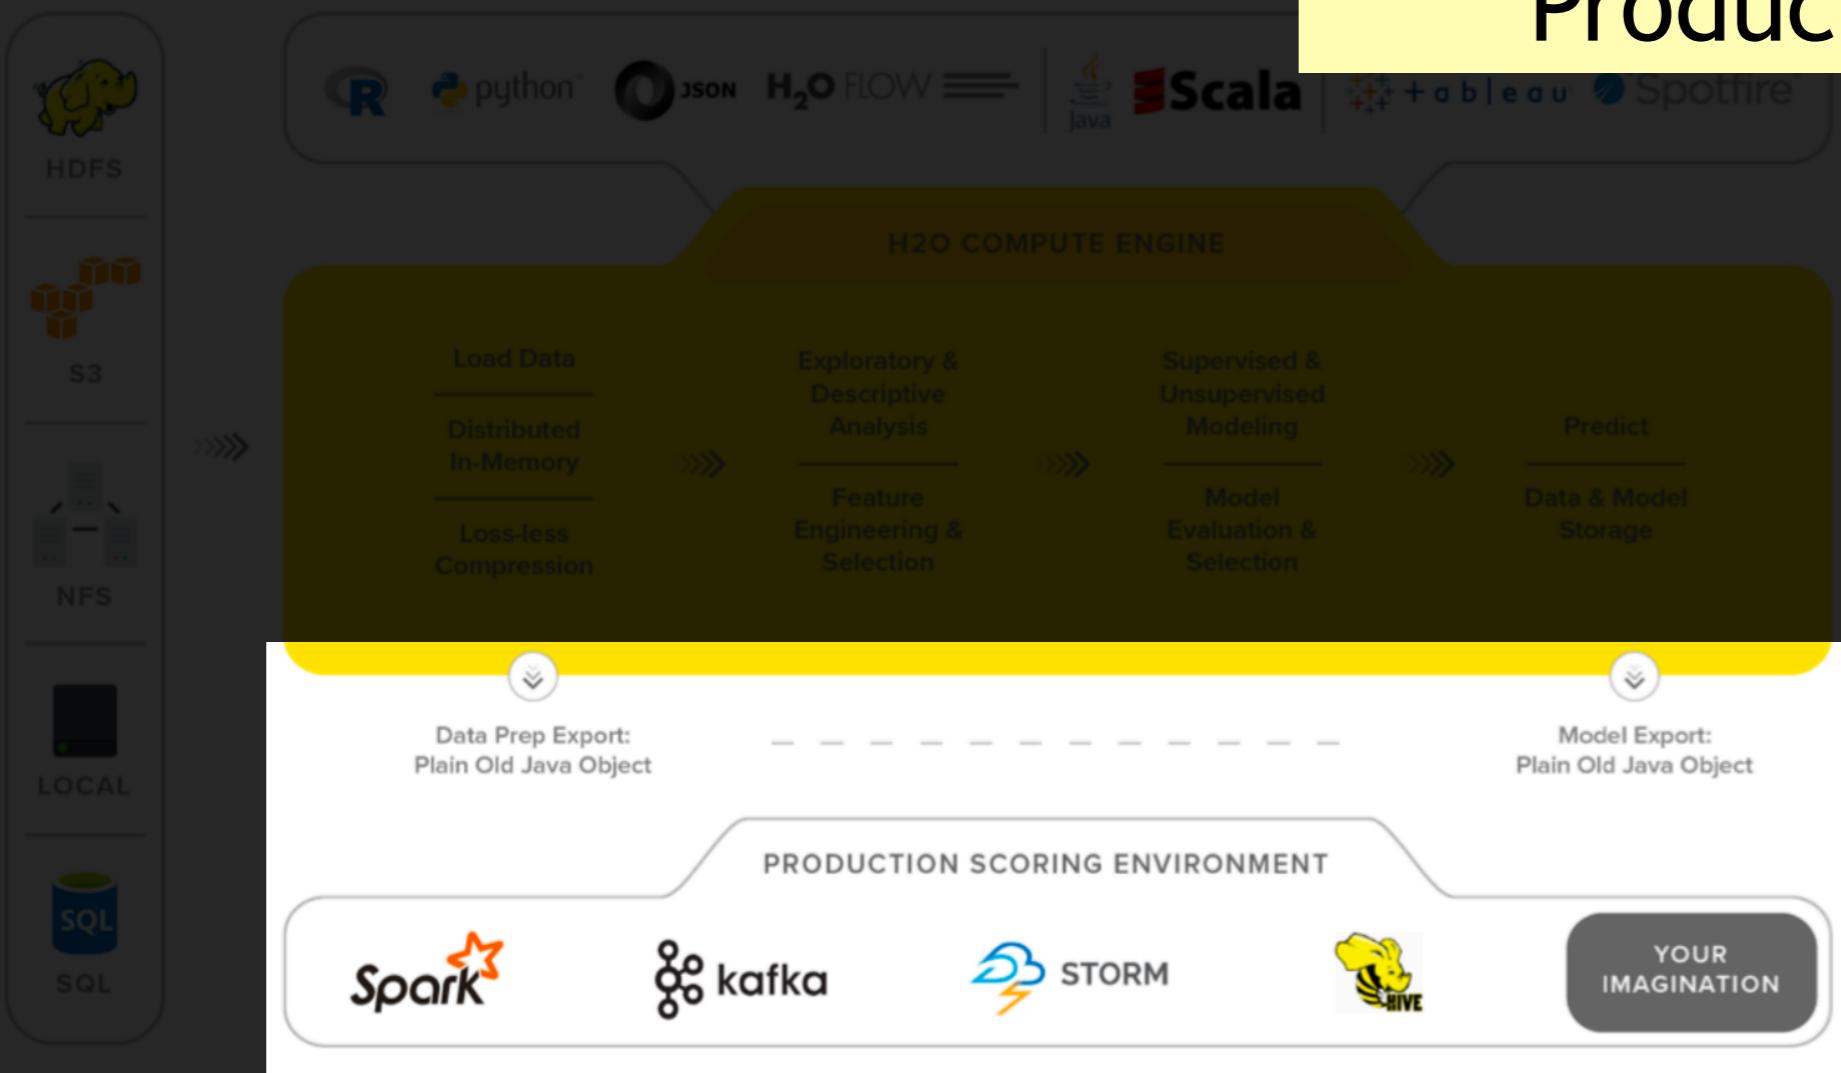
H<sub>2</sub>O.ai

# High Level Architecture



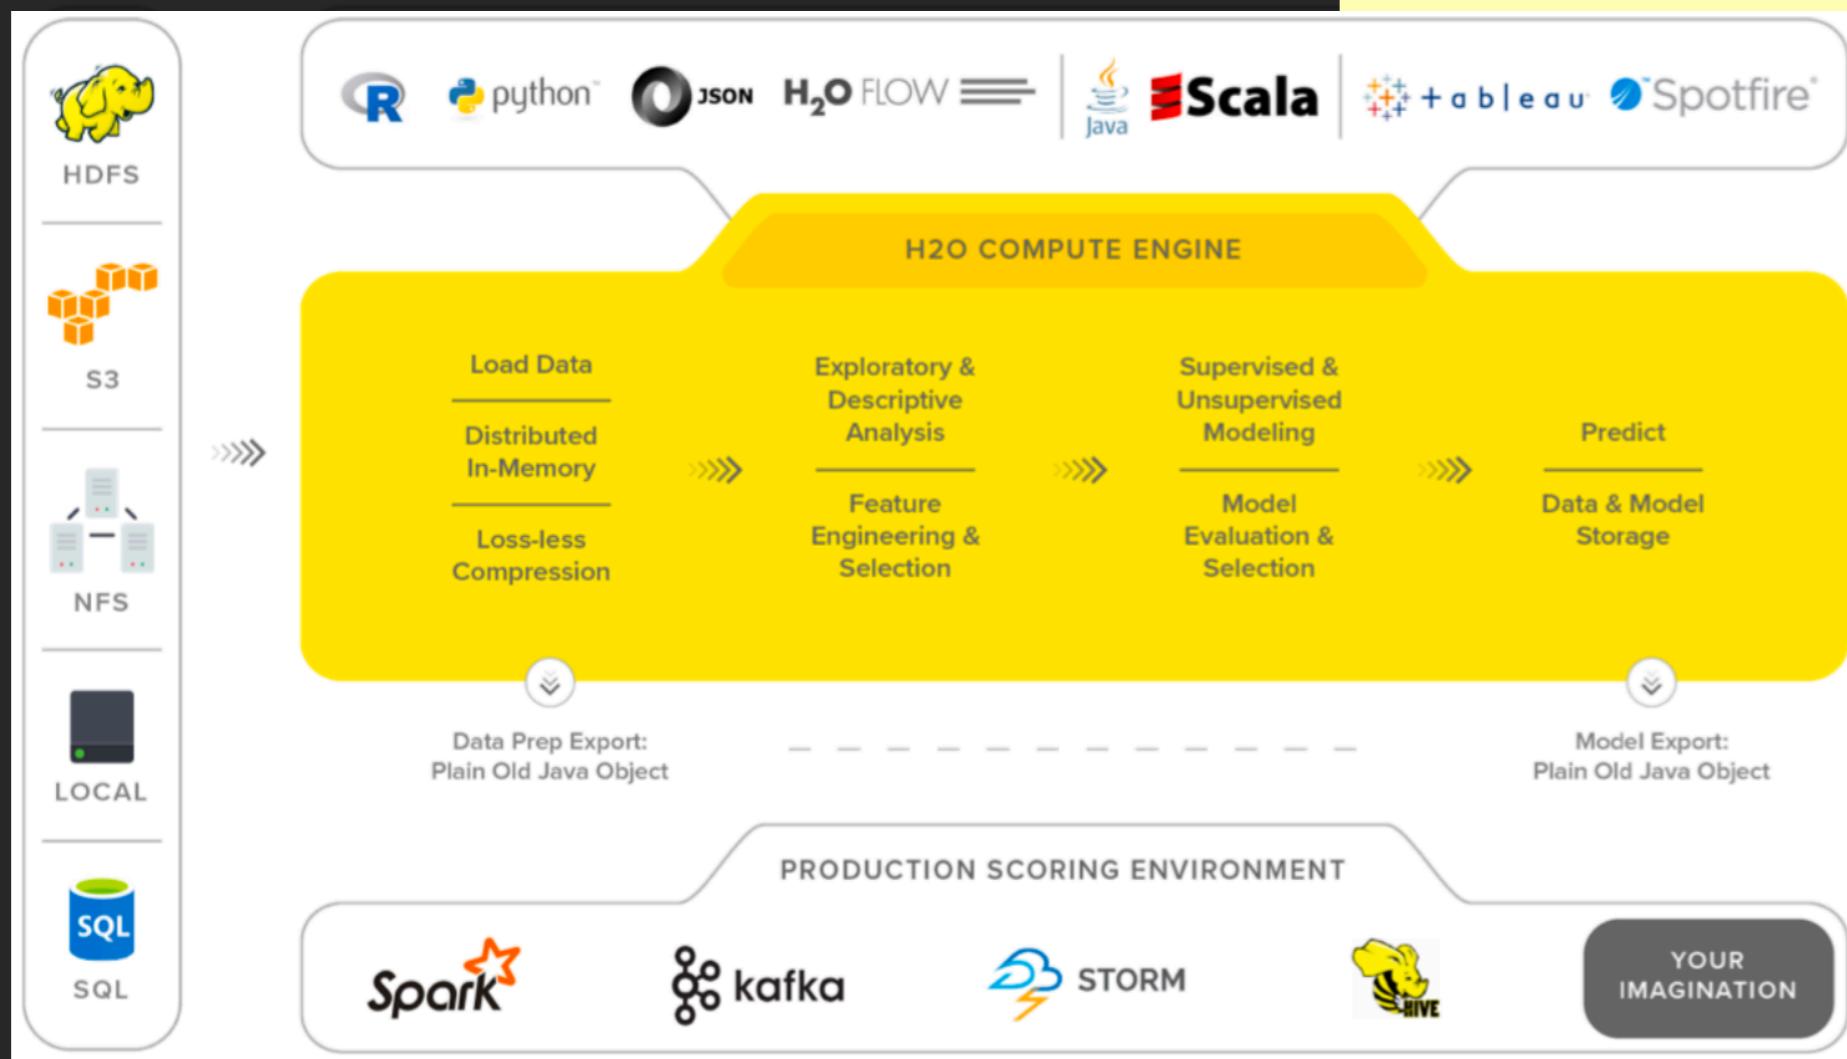
# High Level Architecture

## Export Standalone Models for Production



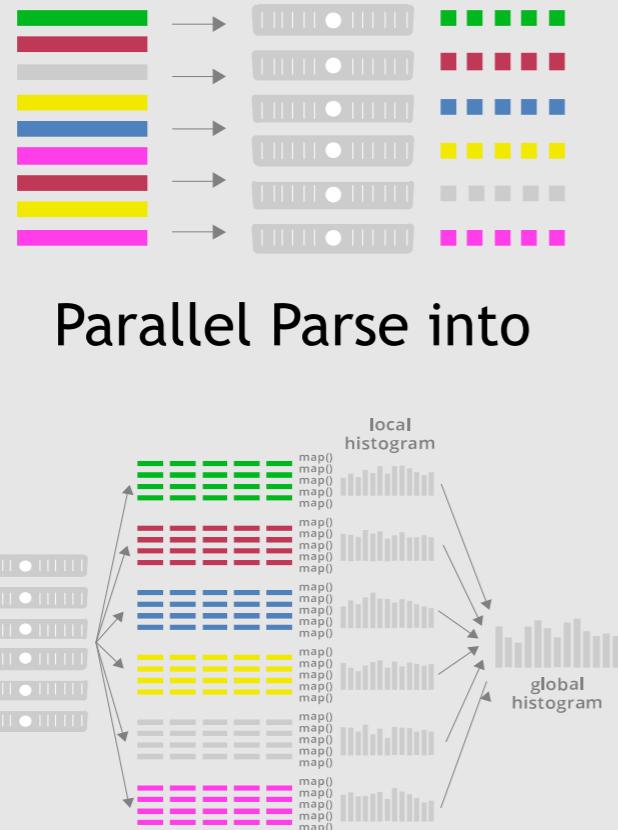
# High Level Architecture

## H<sub>2</sub>O - Open Source Machine Learning Platform



# Distributed Algorithms

## Foundation for Distributed Algorithms



**Parallel Parse into**  
**Fine Grain Map Reduce Illustration: Scalable**

## Advantageous Foundation

- Foundation for In-Memory Distributed Algorithm Calculation - **Distributed Data Frames** and **columnar compression**
- All algorithms are distributed in H<sub>2</sub>O: GBM, GLM, DRF, Deep Learning and more. Fine-grained map-reduce iterations.
- **Only enterprise-grade, open-source distributed algorithms in the market**

## User Benefits

- “Out-of-box” functionalities for all algorithms (**NO MORE SCRIPTING**) and uniform interface across all languages: R, Python, Java
- **Designed for all sizes of data sets, especially large data**
- **Highly optimized Java code for model exports**
- **In-house expertise for all algorithms**

# H<sub>2</sub>O Core



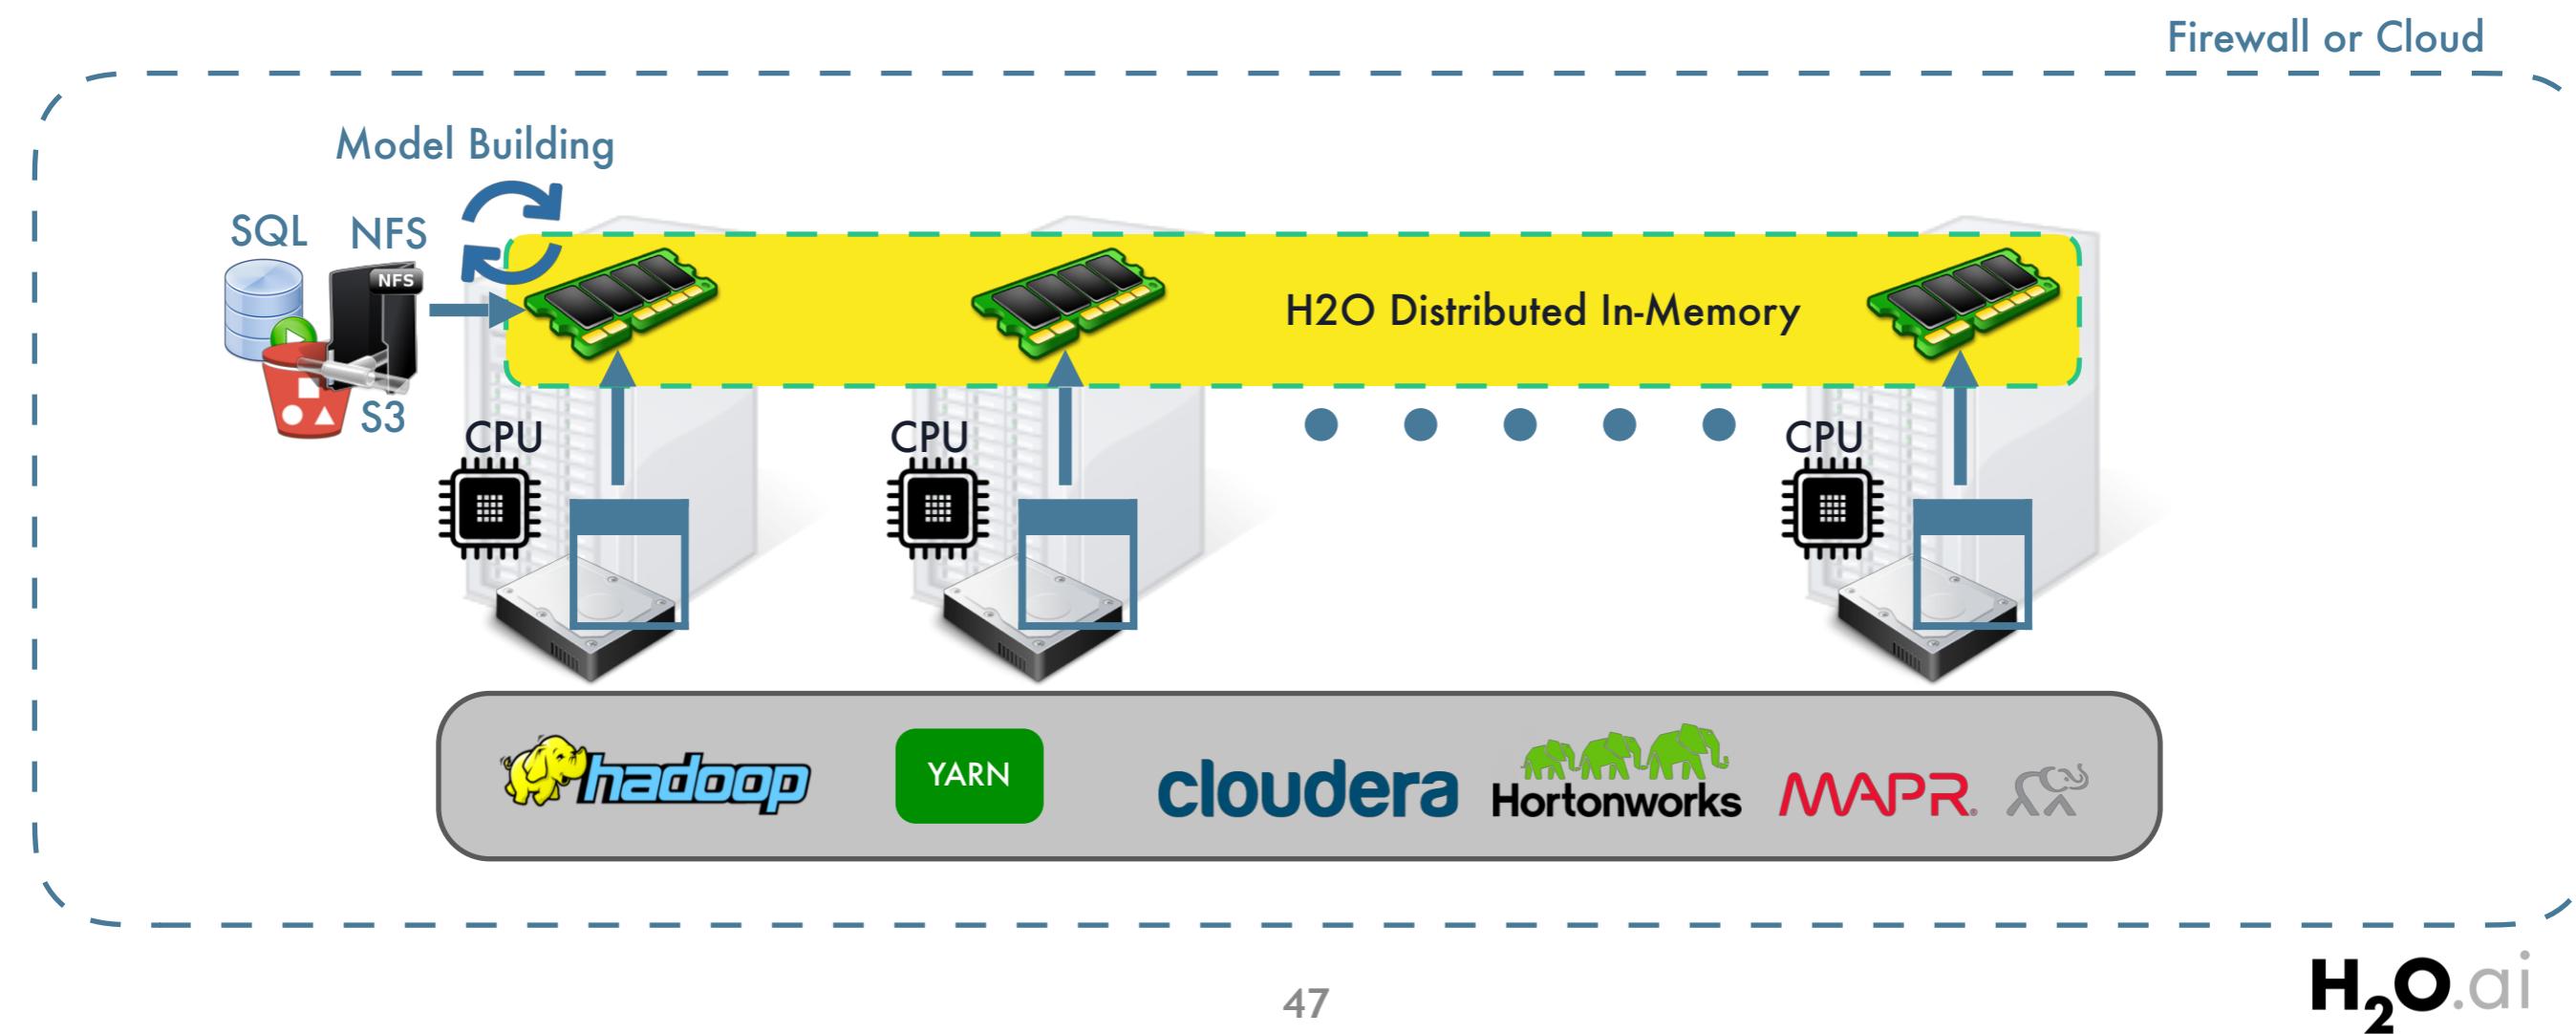
# H<sub>2</sub>O Core



# H<sub>2</sub>O Core



# H<sub>2</sub>O Core



# H<sub>2</sub>O-3 Algorithms Overview

## Supervised Learning

### Statistical Analysis

- **Generalized Linear Models:** Binomial, Gaussian, Gamma, Poisson and Tweedie
- **Naïve Bayes**

### Ensembles

- **Distributed Random Forest:** Classification or regression models
- **Gradient Boosting Machine:** Produces an ensemble of decision trees with increasing refined approximations

### Deep Neural Networks

- **Deep learning:** Create multi-layer feed forward neural networks starting with an input layer followed by multiple layers of nonlinear transformations

## Unsupervised Learning

### Clustering

- **K-means:** Partitions observations into k clusters/groups of the same spatial size. Automatically detect optimal k

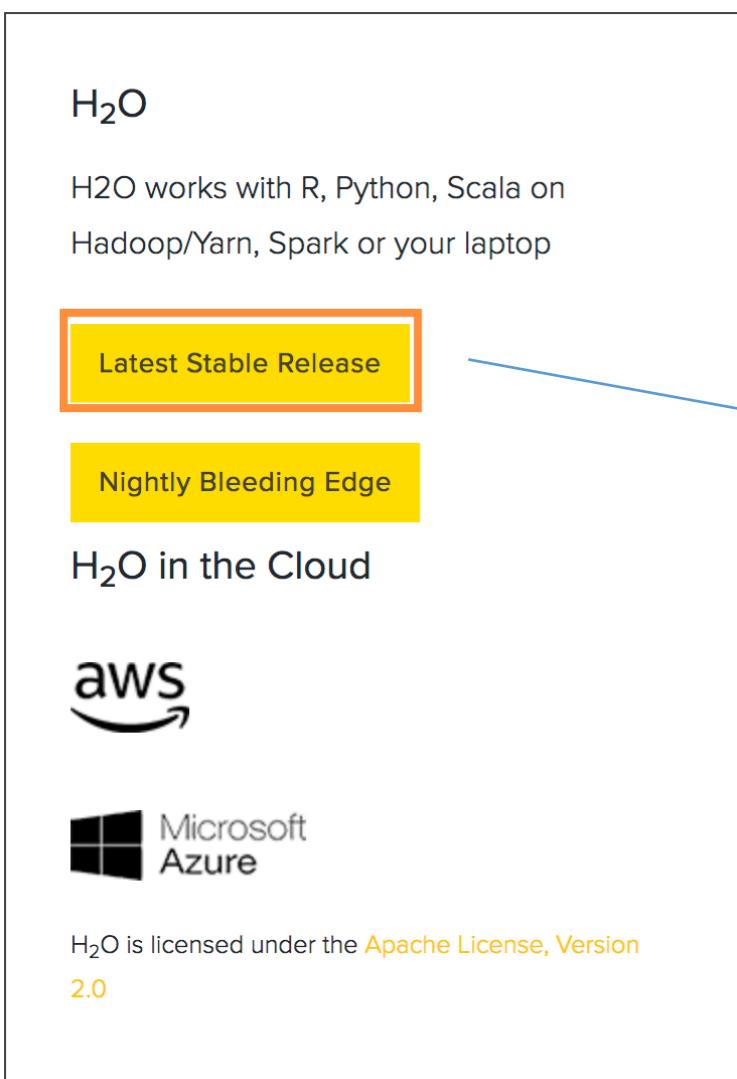
### Dimensionality Reduction

- **Principal Component Analysis:** Linearly transforms correlated variables to independent components
- **Generalized Low Rank Models:** extend the idea of PCA to handle arbitrary data consisting of numerical, Boolean, categorical, and missing data

### Anomaly Detection

- **Autoencoders:** Find outliers using a nonlinear dimensionality reduction using deep learning

# Downloading H<sub>2</sub>O



**H<sub>2</sub>O**  
Version 3.18.0.10

Fast Scalable Machine Learning API  
For Smarter Applications

DOWNLOAD AND RUN    INSTALL IN R    INSTALL IN PYTHON    INSTALL ON HADOOP    USE FROM MAVEN

**DOWNLOAD H<sub>2</sub>O**

Get started with H<sub>2</sub>O in 3 easy steps

1. Download H<sub>2</sub>O. This is a zip file that contains everything you need to get started.
2. From your terminal, run:  

```
cd ~/Downloads
unzip h2o-3.18.0.10.zip
cd h2o-3.18.0.10
java -jar h2o.jar
```
3. Point your browser to <http://localhost:54321>

**copy**

**H<sub>2</sub>O.ai**

**h2o.ai/download/**

# H<sub>2</sub>O Flow

The screenshot shows the H2O Flow web application running in a browser window. The title bar reads "H2O Flow". The address bar shows "localhost:54321/flow/index.html". The top navigation bar includes "Flow", "Cell", "Data", "Model" (which is currently selected), "Score", "Admin", and "Help". A toolbar with various icons is visible above the main content area.

The main content area has a sidebar on the left titled "assit" which contains a "Assistance" section listing various H2O routines with descriptions:

Routine	Description
importFiles	Import file(s) into H2O
getFrames	Get a list of frames in H2O
splitFrame	Split a frame into two or more
mergeFrames	Merge two frames into one
getModels	Get a list of models in H2O
getGrids	Get a list of grid search results
getPredictions	Get a list of predictions in H2O
getJobs	Get a list of jobs running in H2O
buildModel	Build a model
runAutoML	Automatically train and tune a model
importModel	Import a saved model
predict	Make a prediction

A dropdown menu is open under the "Model" tab, listing various machine learning models:

- Aggregator...
- Deep Learning...
- Distributed Random Forest...
- Gradient Boosting Machine...
- Generalized Linear Modeling...
- Generalized Low Rank Modeling...
- K-means...
- Naive Bayes...
- Principal Components Analysis...
- Stacked Ensemble...
- Word2Vec...
- XGBoost...

The right side of the screen features a "HELP" panel with sections for "Using Flow for the first time?", "Quickstart Videos", "STAR H2O ON GITHUB!", "GENERAL" (with links to Flow Web UI, Importing Data, Building Models, Making Predictions, Using Flows, and Troubleshooting), and "EXAMPLES" (with a note about Flow packs and a link to Browse installed packs...).

At the bottom, the URL "localhost:54321/flow/index.html#" is shown in the footer, along with "Connections: 0" and the H2O logo.

H<sub>2</sub>O.ai

## H<sub>2</sub>O Documentation

Getting Started & User Guides | Q & A | Algorithms | Languages | Tutorials, Examples, & Presentations | API & Developer Docs | For the Enterprise

### Getting Started & User Guides

Open Source

Commercial

**H<sub>2</sub>O**

What is H<sub>2</sub>O?  
**H<sub>2</sub>O User Guide** (Main docs)  
H<sub>2</sub>O Book (O'Reilly)  
Recent Changes  
Open Source License (Apache V2)

Quick Start Video - Flow Web UI  
Quick Start Video - R  
Quick Start Video - Python

**Download H<sub>2</sub>O**

**Sparkling Water**

What is Sparkling Water?  
**Sparkling Water User Guide** 2.3 2.2 2.1  
Sparkling Water Booklet  
RSparkling Readme  
PySparkling User Guide 2.3 2.2 2.1  
Recent Changes 2.3 2.2 2.1  
Open Source License (Apache V2)

Quick Start Video - Scala

**Download Sparkling Water**

**Driverless AI**

What is Driverless AI?  
Driverless AI User Guide   
Recent Changes  
Driverless AI Booklet  
MLI with Driverless AI Booklet

Quick Start Video - Downloading Driverless AI  
Quick Start Video - Launching an Experiment  
Driverless AI Webinars

**Download Driverless AI**

**H<sub>2</sub>O4GPU (alpha)**

H<sub>2</sub>O4GPU Readme  
Open Source License (Apache V2)

**Download H<sub>2</sub>O4GPU**

Jakub Háva  
[jakub.hava@h2o.ai](mailto:jakub.hava@h2o.ai)

# The Next Generation of Machine Learning on Apache Spark

Firemní semináře @ MFF UK, Praha  
October 24, 2018

Spark<sup>®</sup> + H<sub>2</sub>O

---

SPARKLING  
WATER

# H<sub>2</sub>O Products

**H<sub>2</sub>O.ai**

In-Memory, Distributed  
Machine Learning Algorithms  
with H2O Flow GUI



H2O AI Open Source Engine  
Integration with Spark

**H<sub>2</sub>O4GPU**

Lightning Fast machine  
learning on GPUs

DRIVERLESSAI

Automatic feature  
engineering, machine learning  
and interpretability

**Steam**

Secure multi-tenant H2O clusters

H<sub>2</sub>O+Spark =  
Sparkling  
Water

# Sparkling Water

- Transparent integration of H2O with Spark ecosystem - MLlib and H2O side-by-side
- Transparent use of H2O data structures and algorithms with Spark API
- Excels in existing Spark workflows requiring advanced Machine Learning algorithms
- Deployment tool for Driverless AI MOJOs

Functionality missing in H2O can be replaced by Spark and vice versa

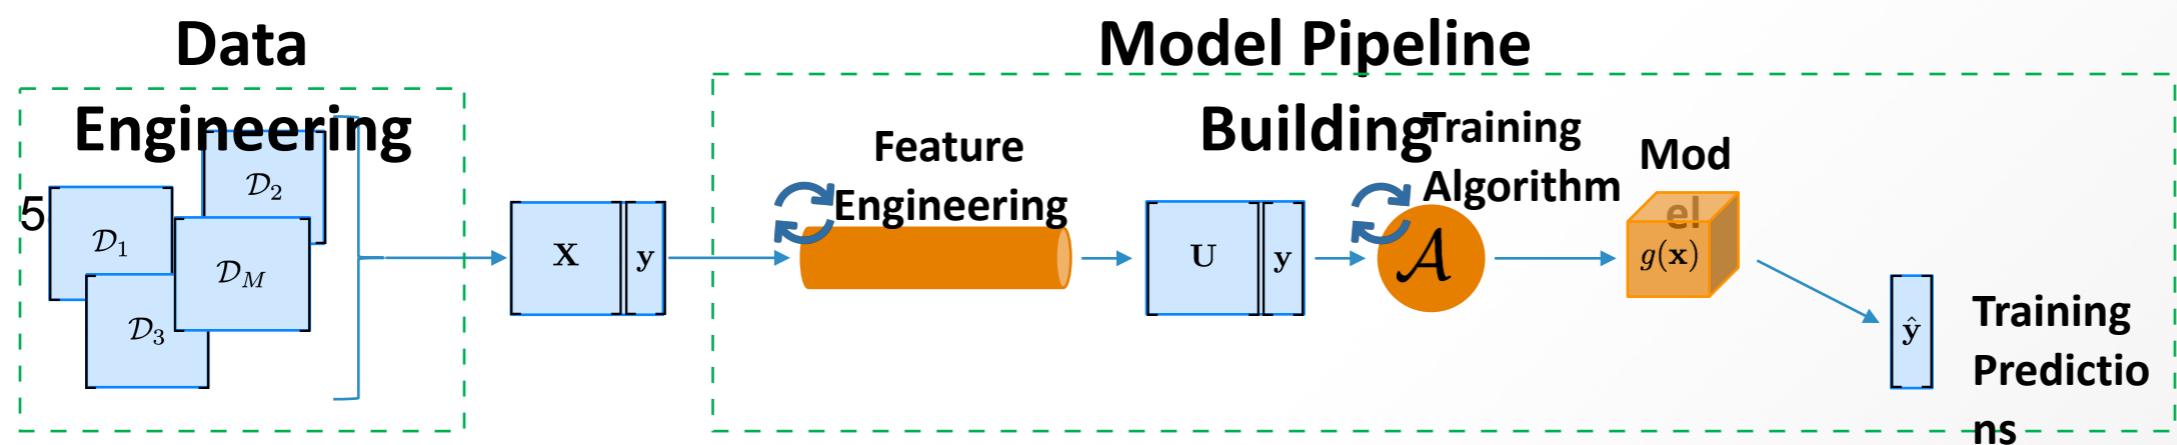
# Benefits



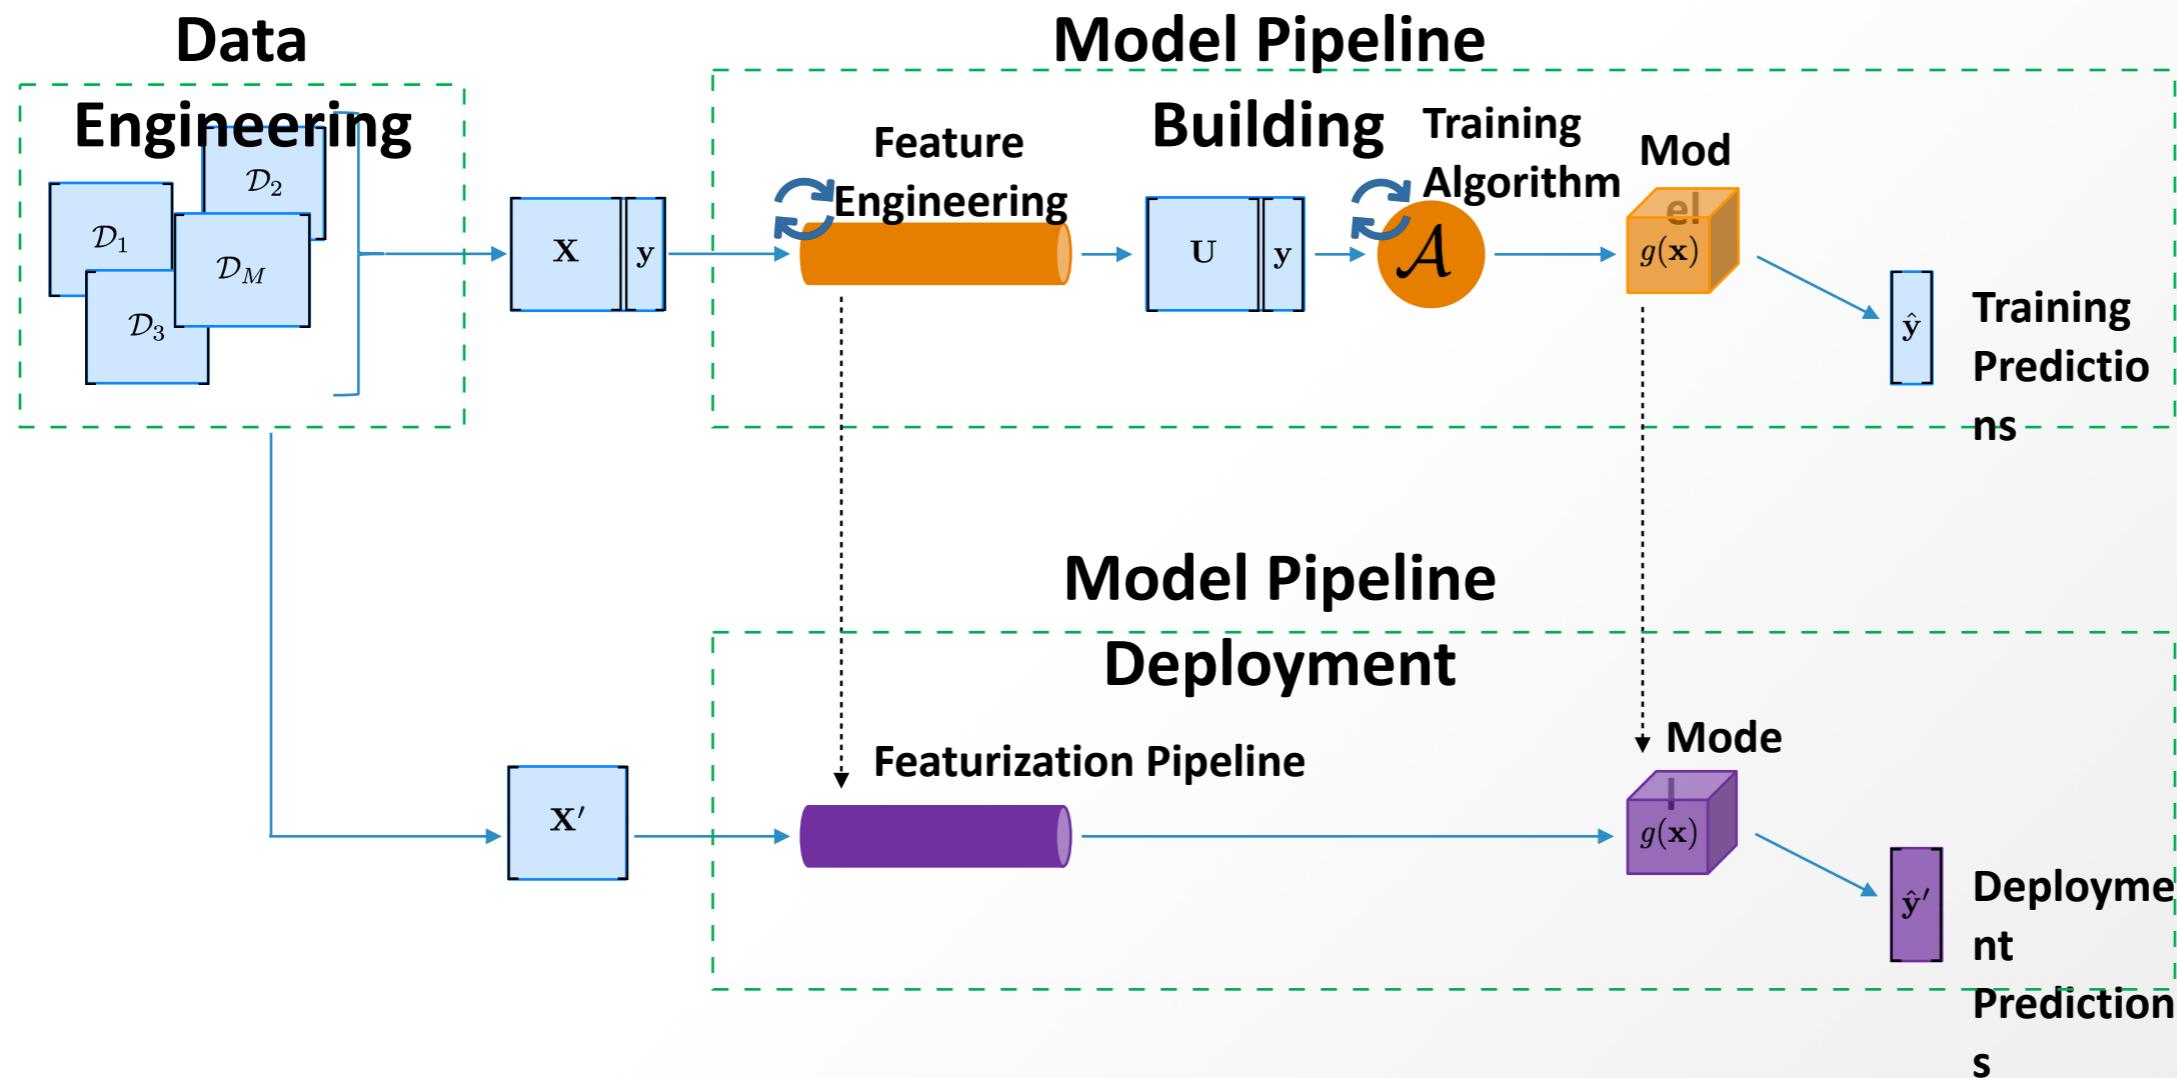
- Additional algorithms
  - NLP
- Powerful data munging
- ML Pipelines
- Advanced algorithms
  - speed v. accuracy
  - advanced parameters
- Fully distributed and parallelised
- Graphical environment
- R/Python interface

How it fits to  
ML Life-  
Cycle ?

# Basic ML Lifecycle



# Basic ML Lifecycle



#ML4SAIS

# Example Implementation

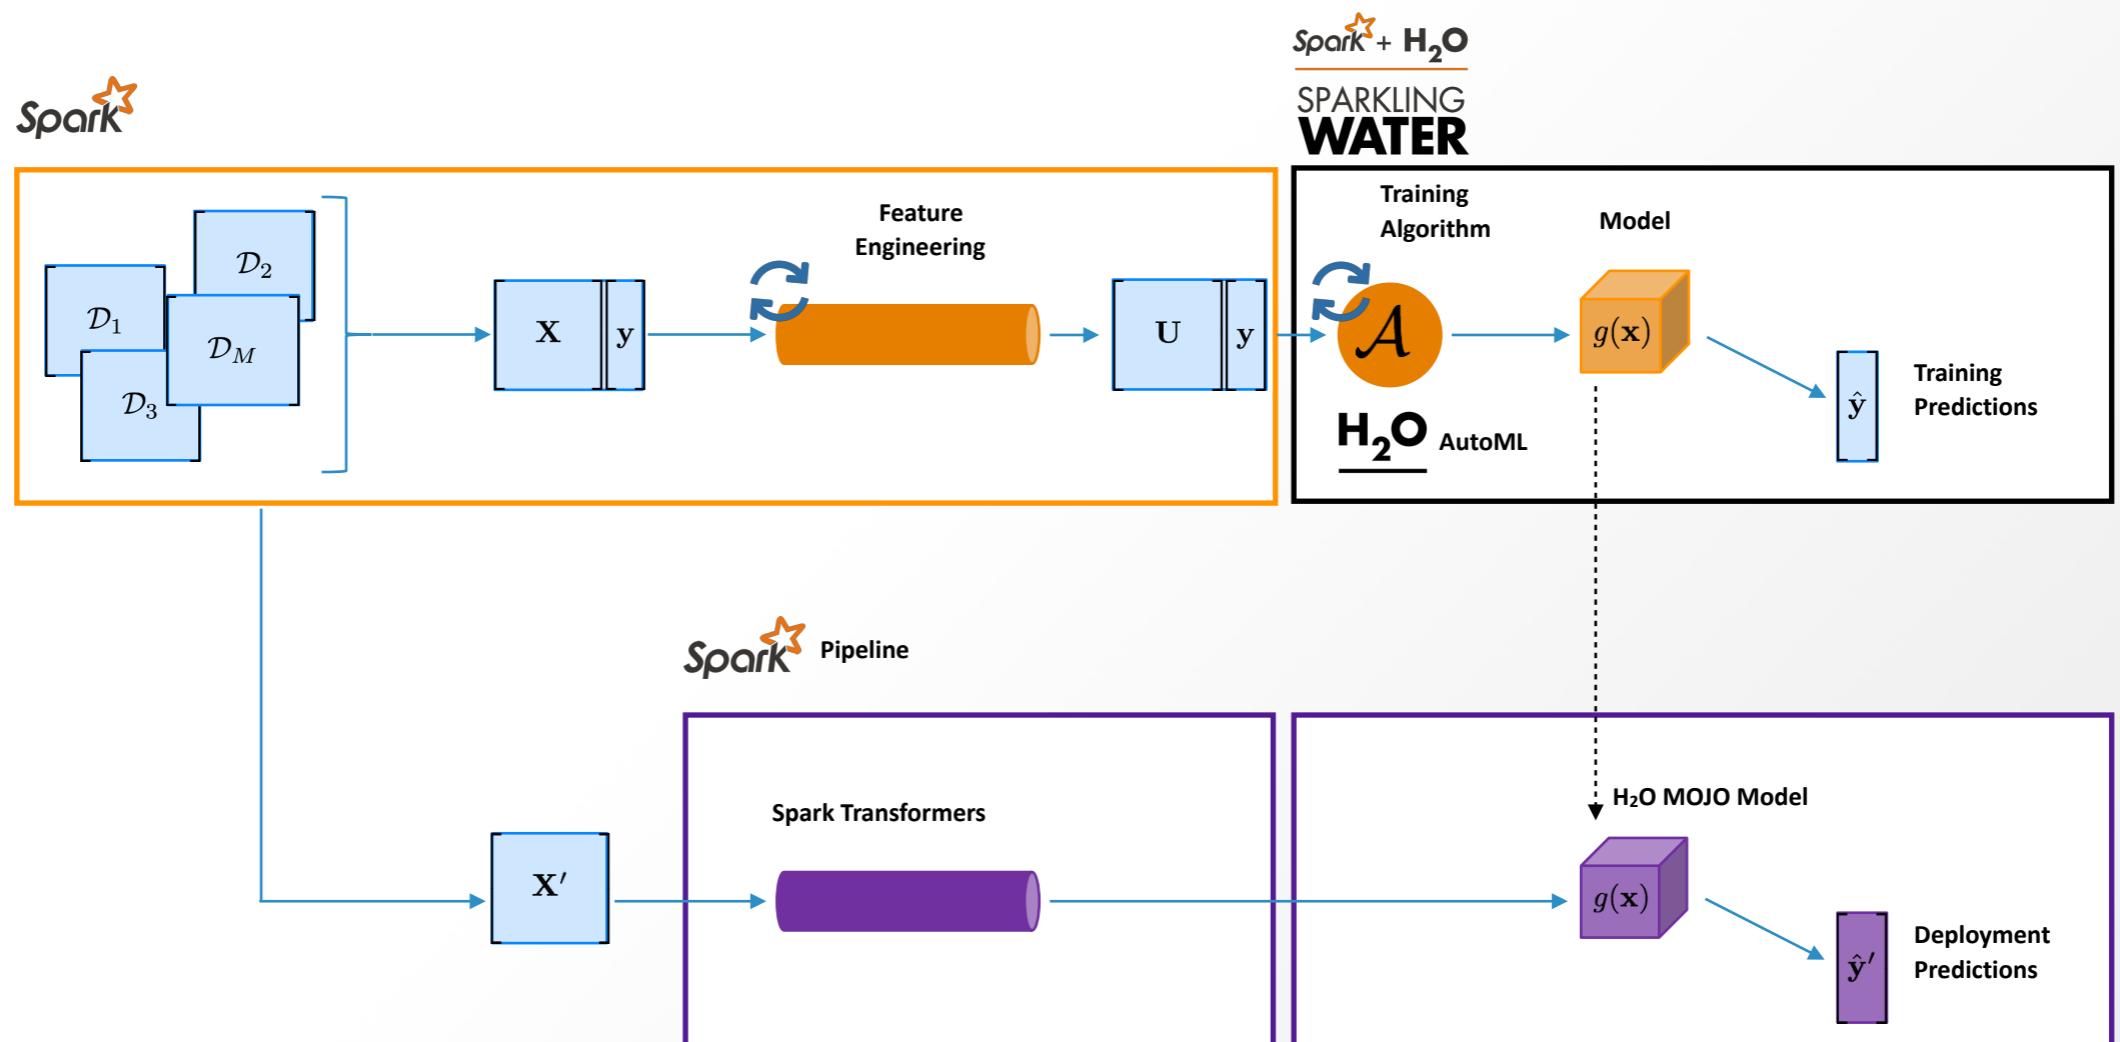
## Model Building

## Model Deployment

Data Engineering	Feature Engineering	Training Algorithm	Deployment Pipeline	Model
Spark	H2O	Spark	H2O MOJO	
Spark	H2O Driverless AI		Spark	H2O Driverless AI MOJO

#ML4SAIS

# Basic ML Lifecycle



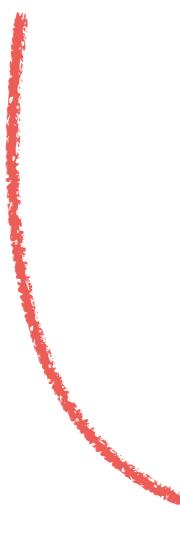
# How to use Sparkling Water?

# Start spark with Sparkling Water

`start.sh`

```
1 $SPARK_HOME/bin/spark-submit \
2   --class water.SparklingWaterDriver \
3   --packages ai.h2o:sparkling-water-examples_2.10:1.6.3 \
4   --executor-memory=6g \
5   --driver-class-path scalastyle.jar /dev/null
```

Raw

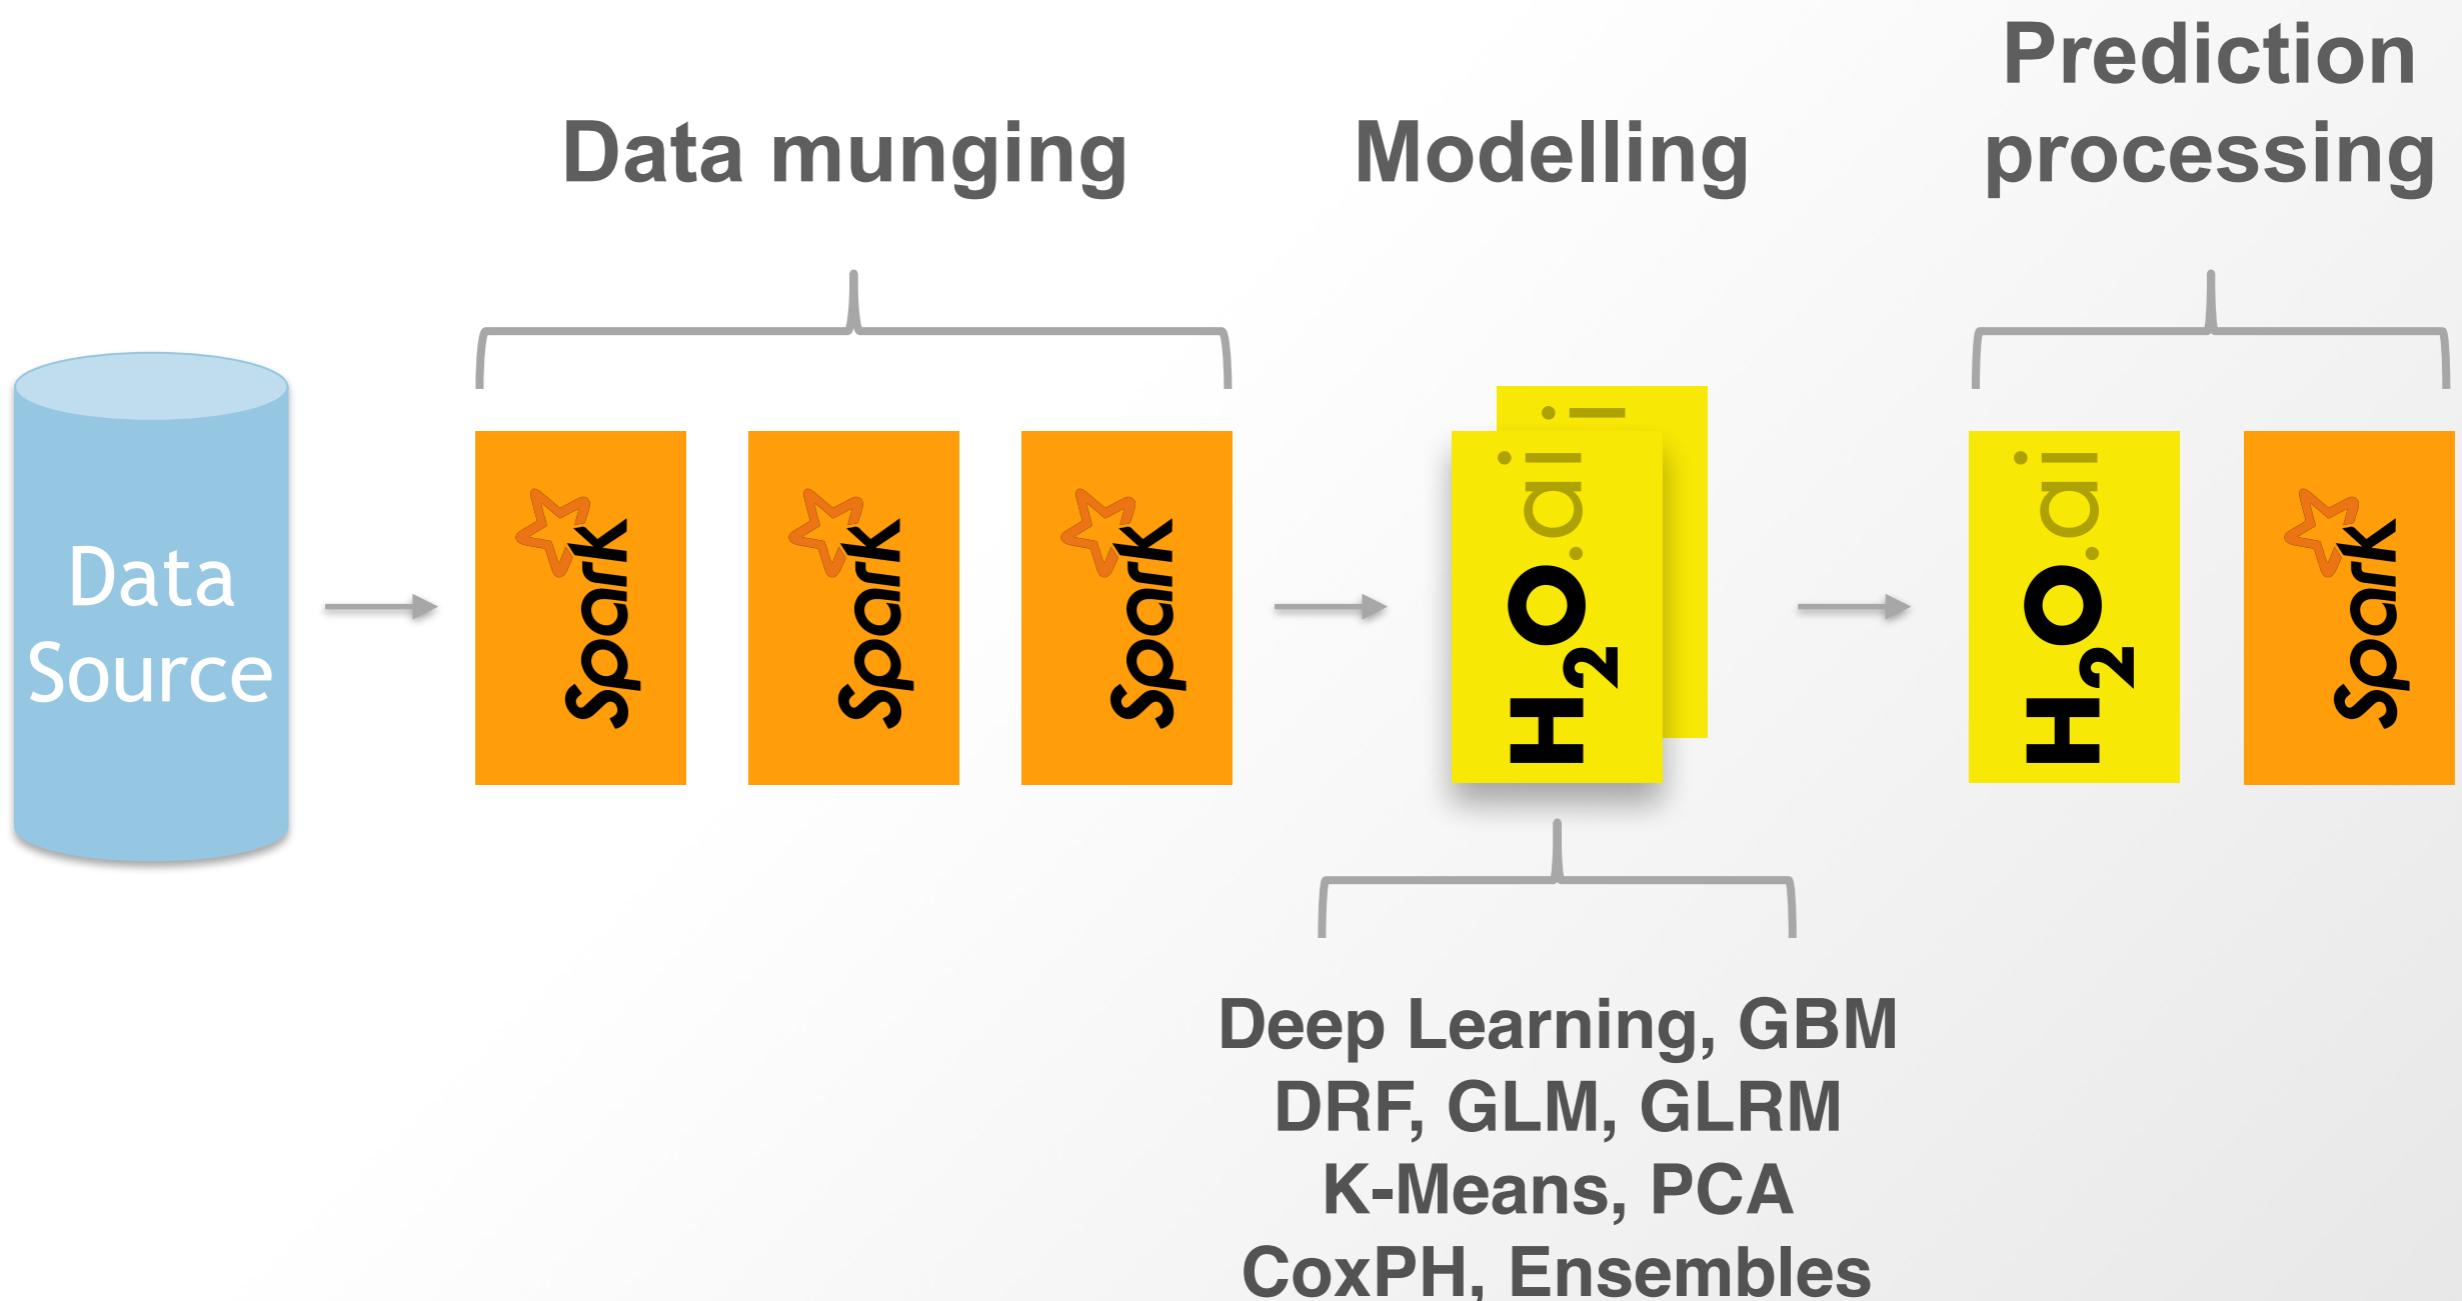


The screenshot shows the H2O Flow web application running at 127.0.0.1. The main title bar says "Start Spark with Sparkling Water". The interface has a toolbar with various icons for file operations like Open, Save, and Print. Below the toolbar is a navigation bar with tabs: OUTLINE, FLOWS, CLIPS, and HELP (which is currently selected). The main content area is titled "Untitled Flow". It contains a sidebar with the text "assist" and a list of "Assistance" routines:

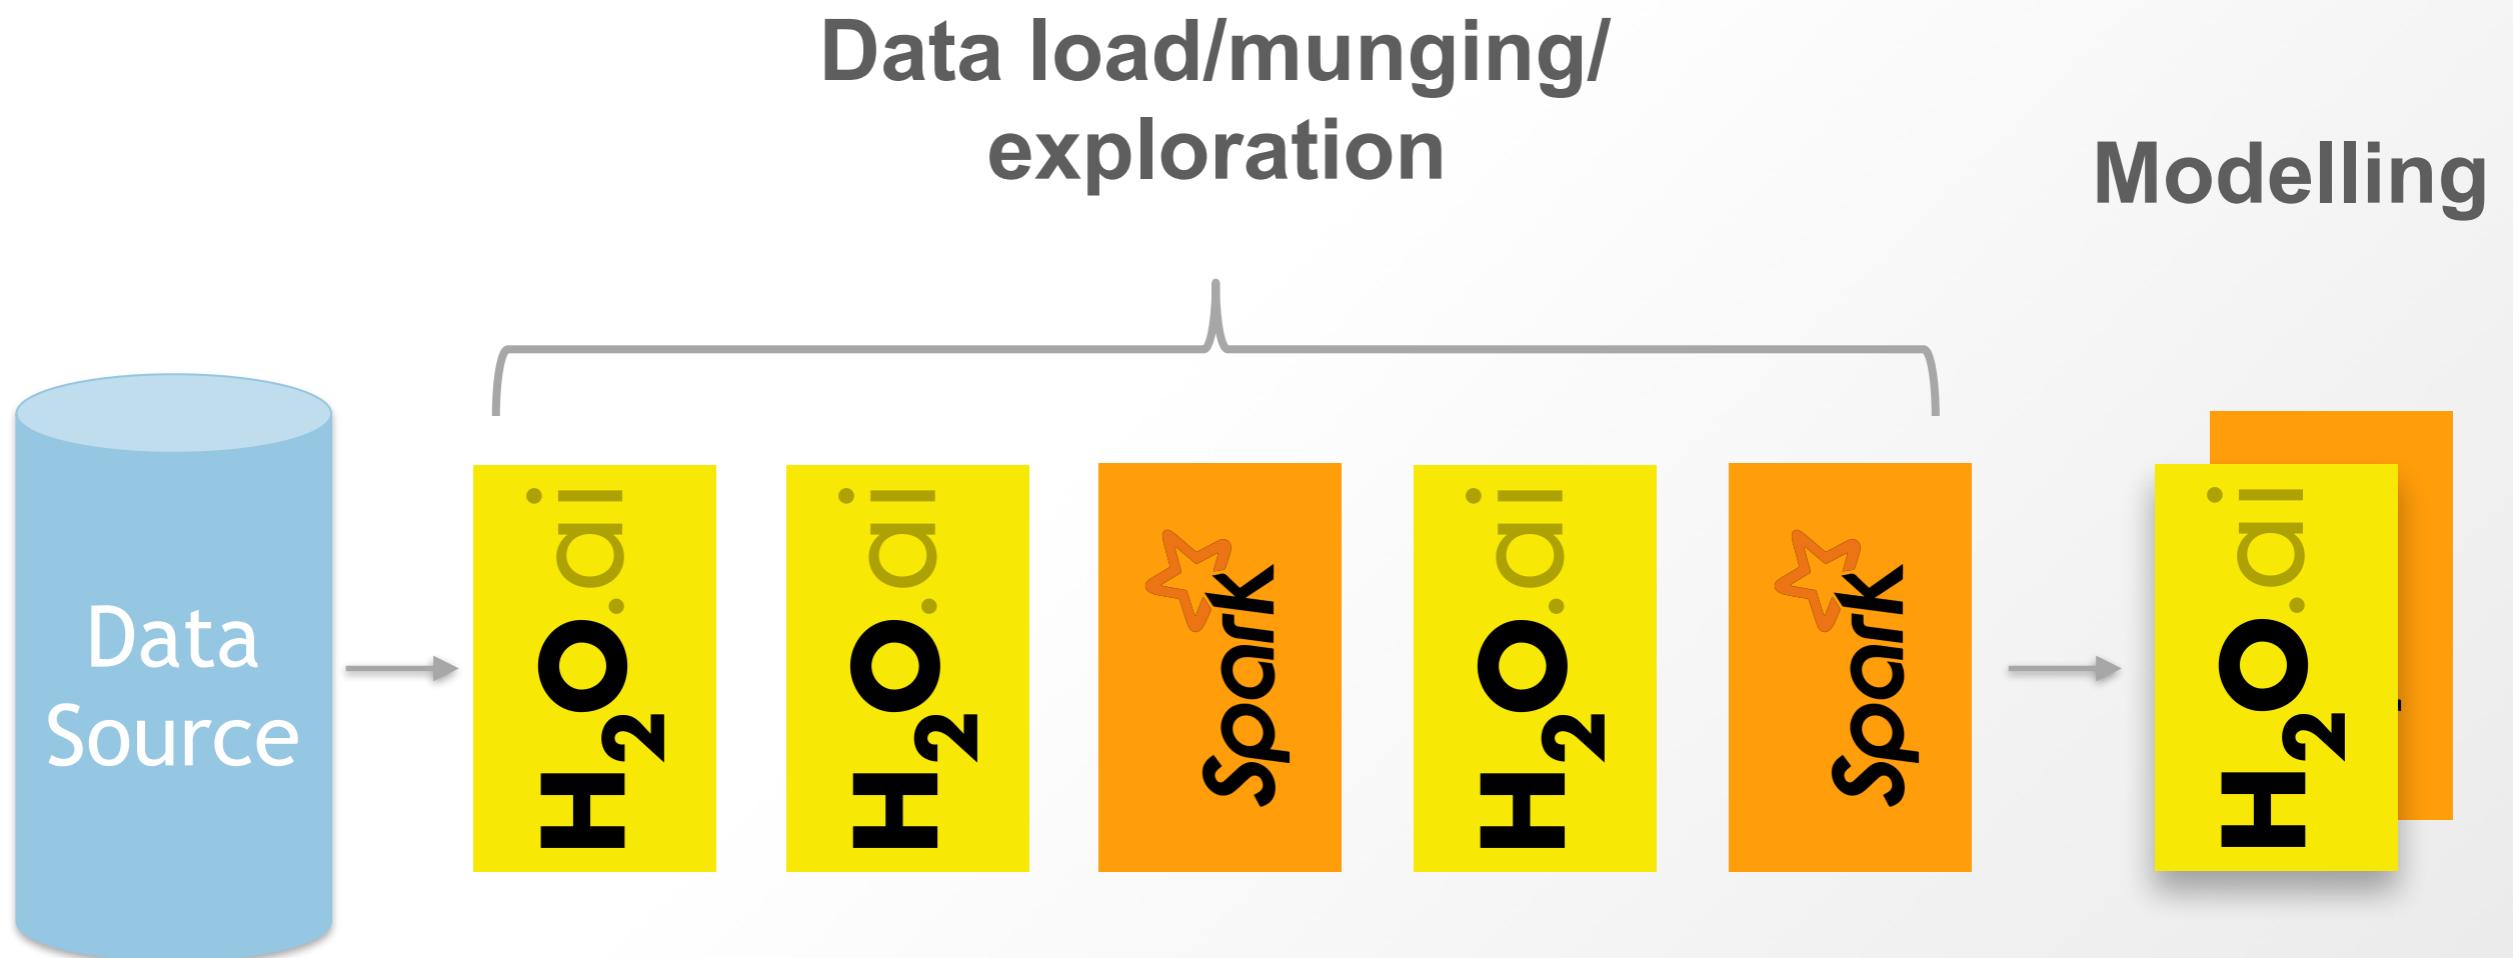
Routine	Description
importFiles	Import file(s) into H <sub>2</sub> O
getFrames	Get a list of frames in H <sub>2</sub> O
splitFrame	Split a frame into two or more frames
getModels	Get a list of models in H <sub>2</sub> O
getGrids	Get a list of grid search results in H <sub>2</sub> O
getPredictions	Get a list of predictions in H <sub>2</sub> O
getJobs	Get a list of jobs running in H <sub>2</sub> O
buildModel	Build a model
importModel	Import a saved model
predict	Make a prediction
getRDDs	Get a list of Spark's RDDs
getDataFrames	Get a list of Spark's data frames

To the right of the sidebar, there is a "Help" section with links to "Quickstart Videos" and "example Flows". It also features a GitHub star button ("Star 1,018") and sections for "GENERAL" and "EXAMPLES". The bottom status bar indicates "Ready" and "Connections: 0".

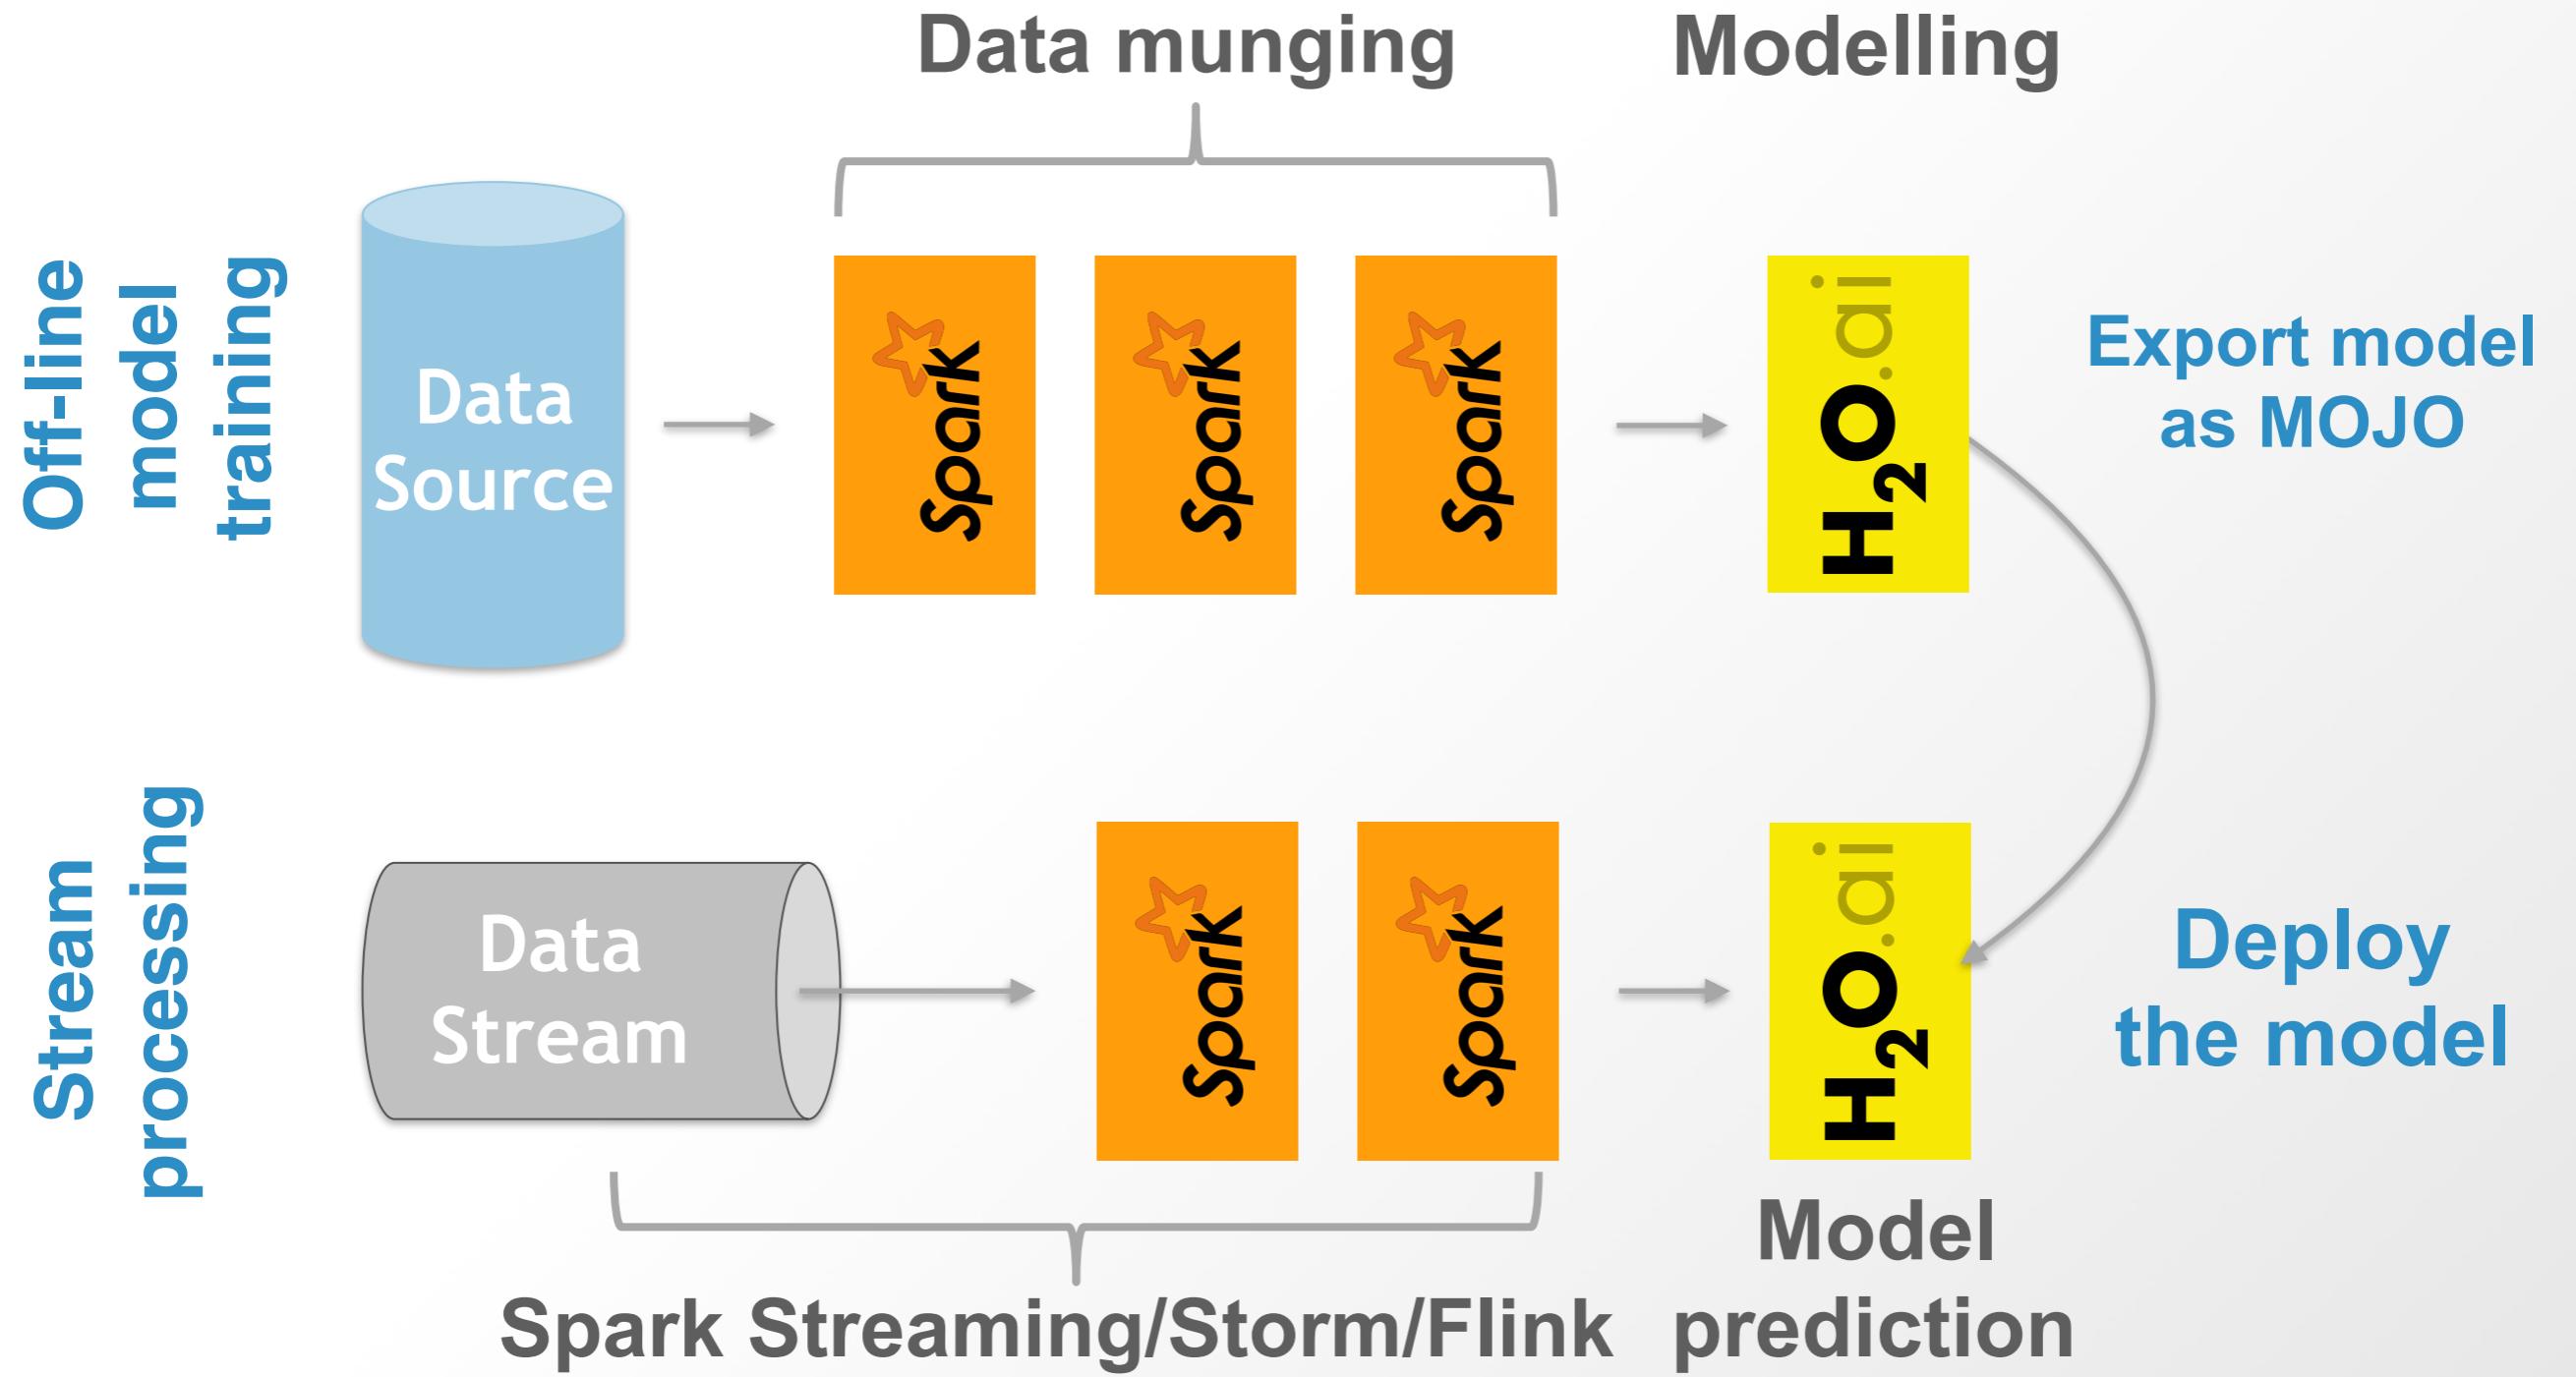
# Model Building



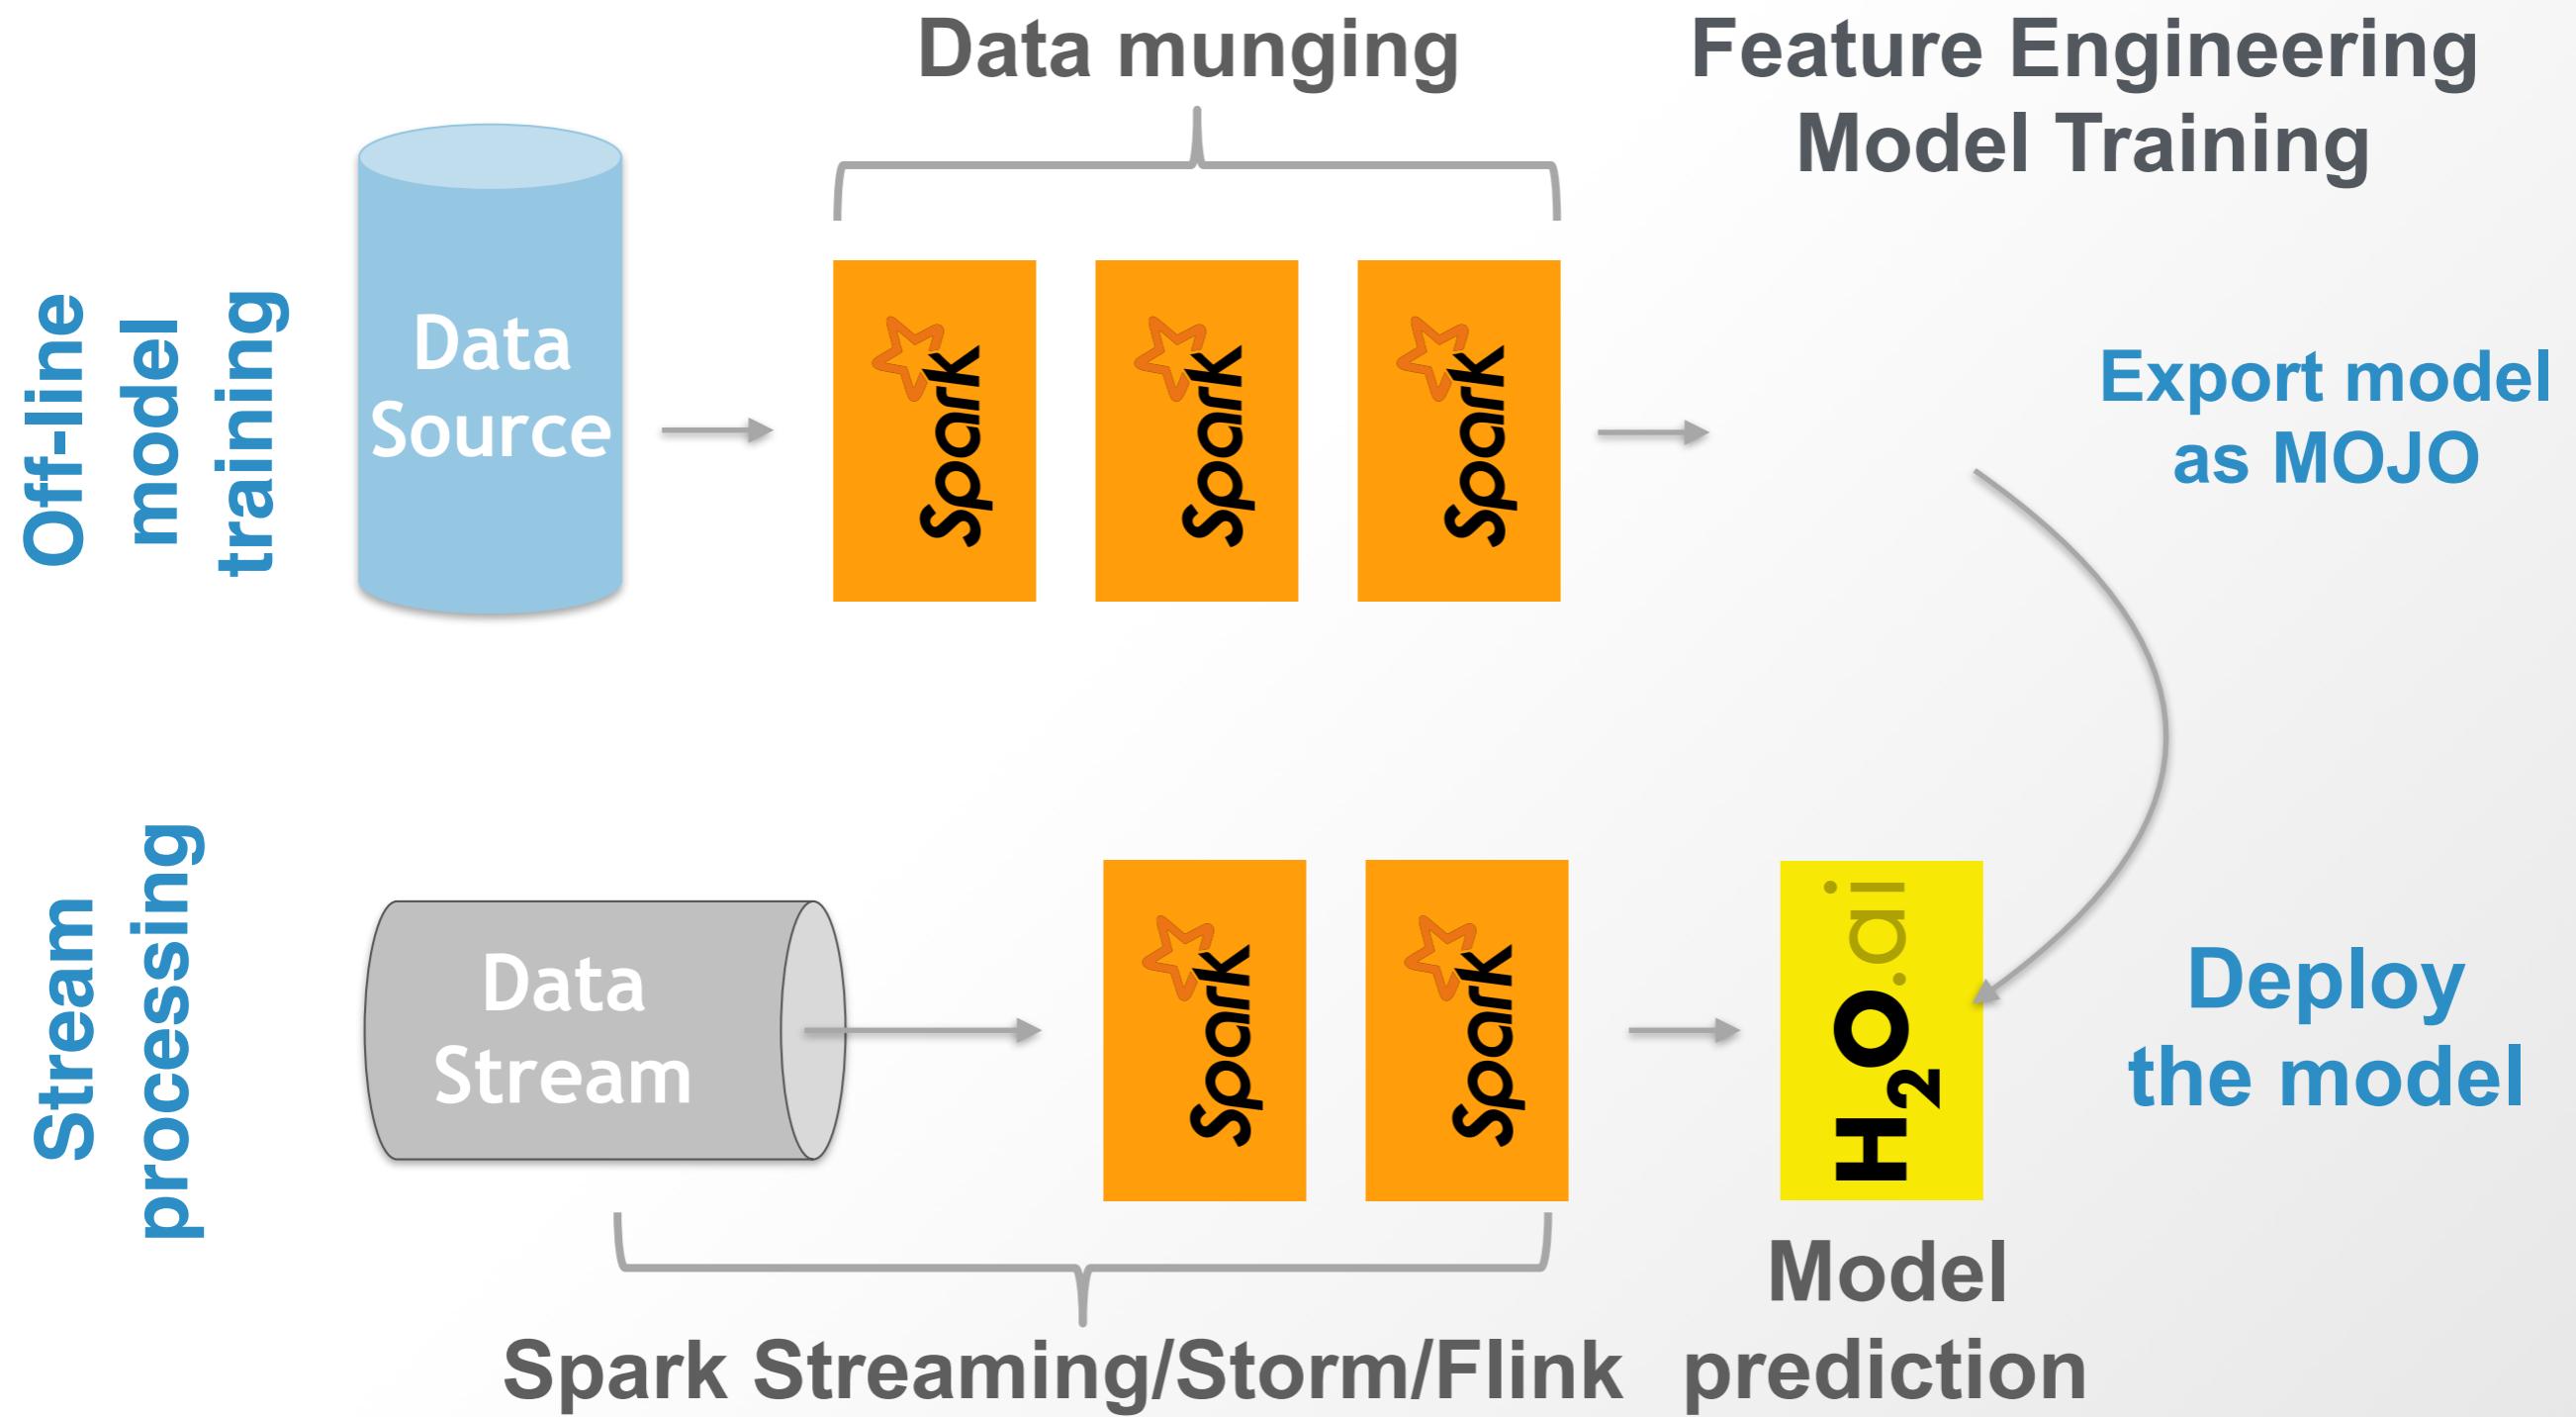
# Data Munging



# Stream Processing



# Stream Processing 2



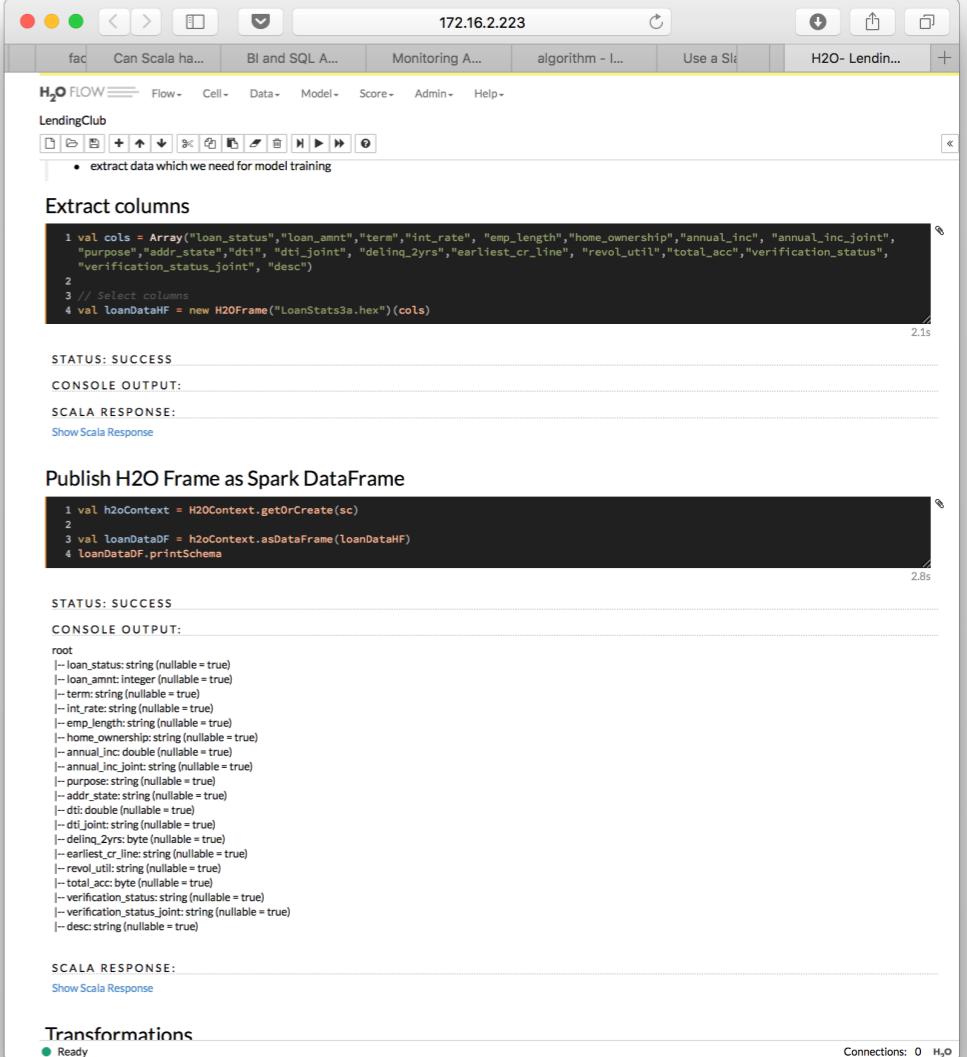
# Scoring

- POJO
  - Plain Old Java Object
- MOJO
  - Model Object Optimised
- MOJO Pipeline
  - MOJO from DAI with additional future transformations
- No runtime dependency on H2O/Driverless AI frameworks

# Features Overview

# Scala code in H2O Flow

- New type of cell
- Access Spark from Flow UI
- Experimenting made easy



The screenshot shows the H2O Flow interface running on a Mac OS X system. The window title is "172.16.2.223". The main area displays two code snippets in a Scala code editor.

**Extract columns:**

```
1 val cols = Array("loan_status","loan_amnt","term","int_rate", "emp_length","home_ownership","annual_inc", "annual_inc_joint",
2 "purpose","addr_state","dti","dti_joint", "delinq_2yrs","earliest_cr_line", "revol_util","total_acc","verification_status",
3 // Select columns
4 val loanDataHF = new H2OFrame("LoanStats3a.hex")(cols)
```

**PUBLISH H2O FRAME AS SPARK DATAFRAME:**

```
1 val h2oContext = H2OContext.getOrCreate(sc)
2
3 val loanDataDF = h2oContext.asDataFrame(loanDataHF)
4 loanDataDF.printSchema
```

Both code blocks result in a **STATUS: SUCCESS** message. The second code block also shows the **SCALA RESPONSE** which lists all the columns and their types for the DataFrame.

# **H2O Frame as Spark's Datasource**

- Use native Spark API to load and save data
- Spark can optimise the queries when loading data from H2O Frame
- Use of Spark query optimiser

# Machine learning pipelines

- Wrap our algorithms as Transformers and Estimators
- Support for embedding them into Spark ML Pipelines
- Can serialise fitted/unfitted pipelines
- Unified API => Arguments are set in the same way for Spark and H2O Models
- Integration with Mojo pipelines

# **PySparkling**

- PySparkling is on PyPi
- Contains all H2O and Sparkling Water dependencies, no need to worry about them
- Just add in on your Python path and that's it
- Correctly specifies dependency on PySpark

# RSparkling

- Sparkling Water for R
- Based on SparklyR package
- Independent on Spark and Sparkling Water version

# **And others!**

- Support for Datasets
- Zeppelin notebook support
- XGBoost Support in Distributed Mode
- Support for high-cardinality fast joins
- Secure Communication - SSL
- Support for Sparse Data conversions
- External Cluster Stabilisation
- Enterprise security support - LDAP, Kerberos
- Driverless AI Pipeline deployment availability

# **Coming features**

- Integration with Steam
  - Deploy on remises
  - Deploy to Cloud

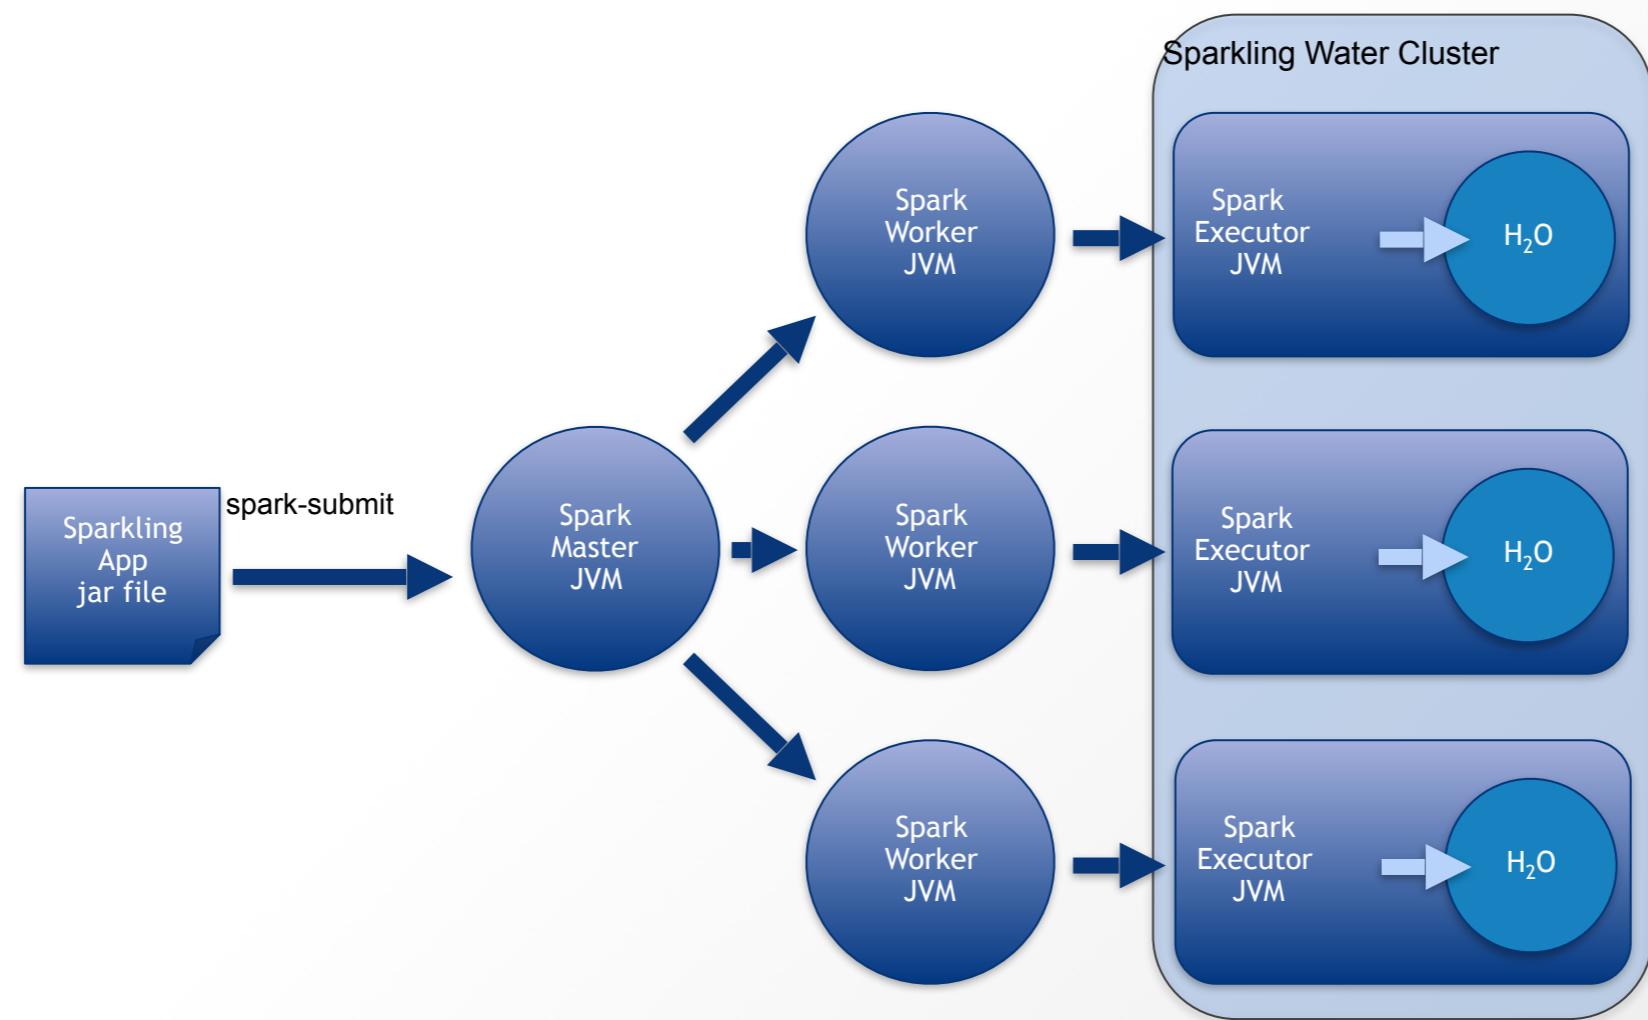
**What is  
inside?**

# Two Backends

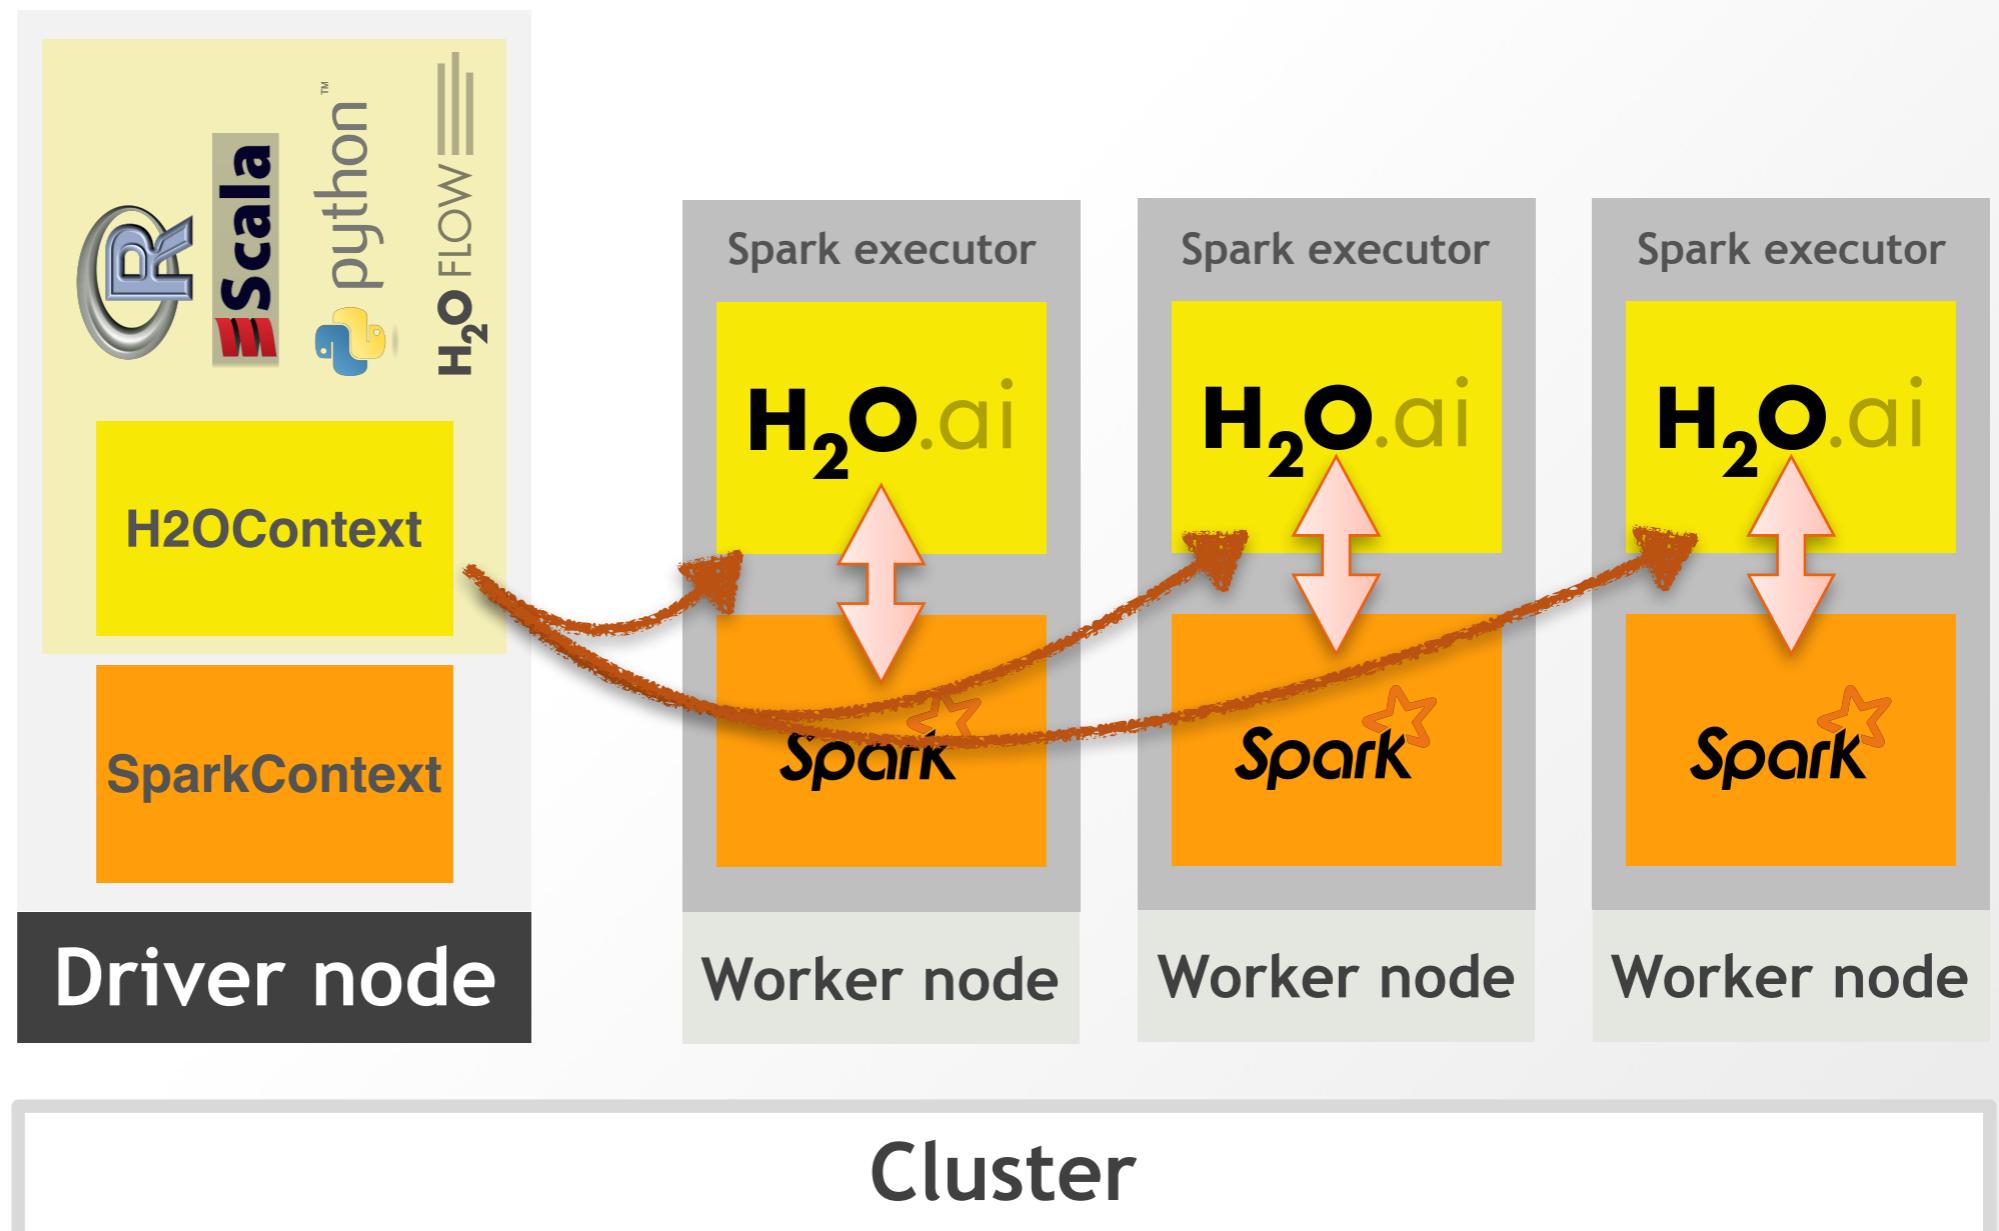
- Sparkling Water has two backends
  - External
  - Internal
- The backend determines where H2O cluster is located
- Each backend is good for different use-cases

# Internal Backend

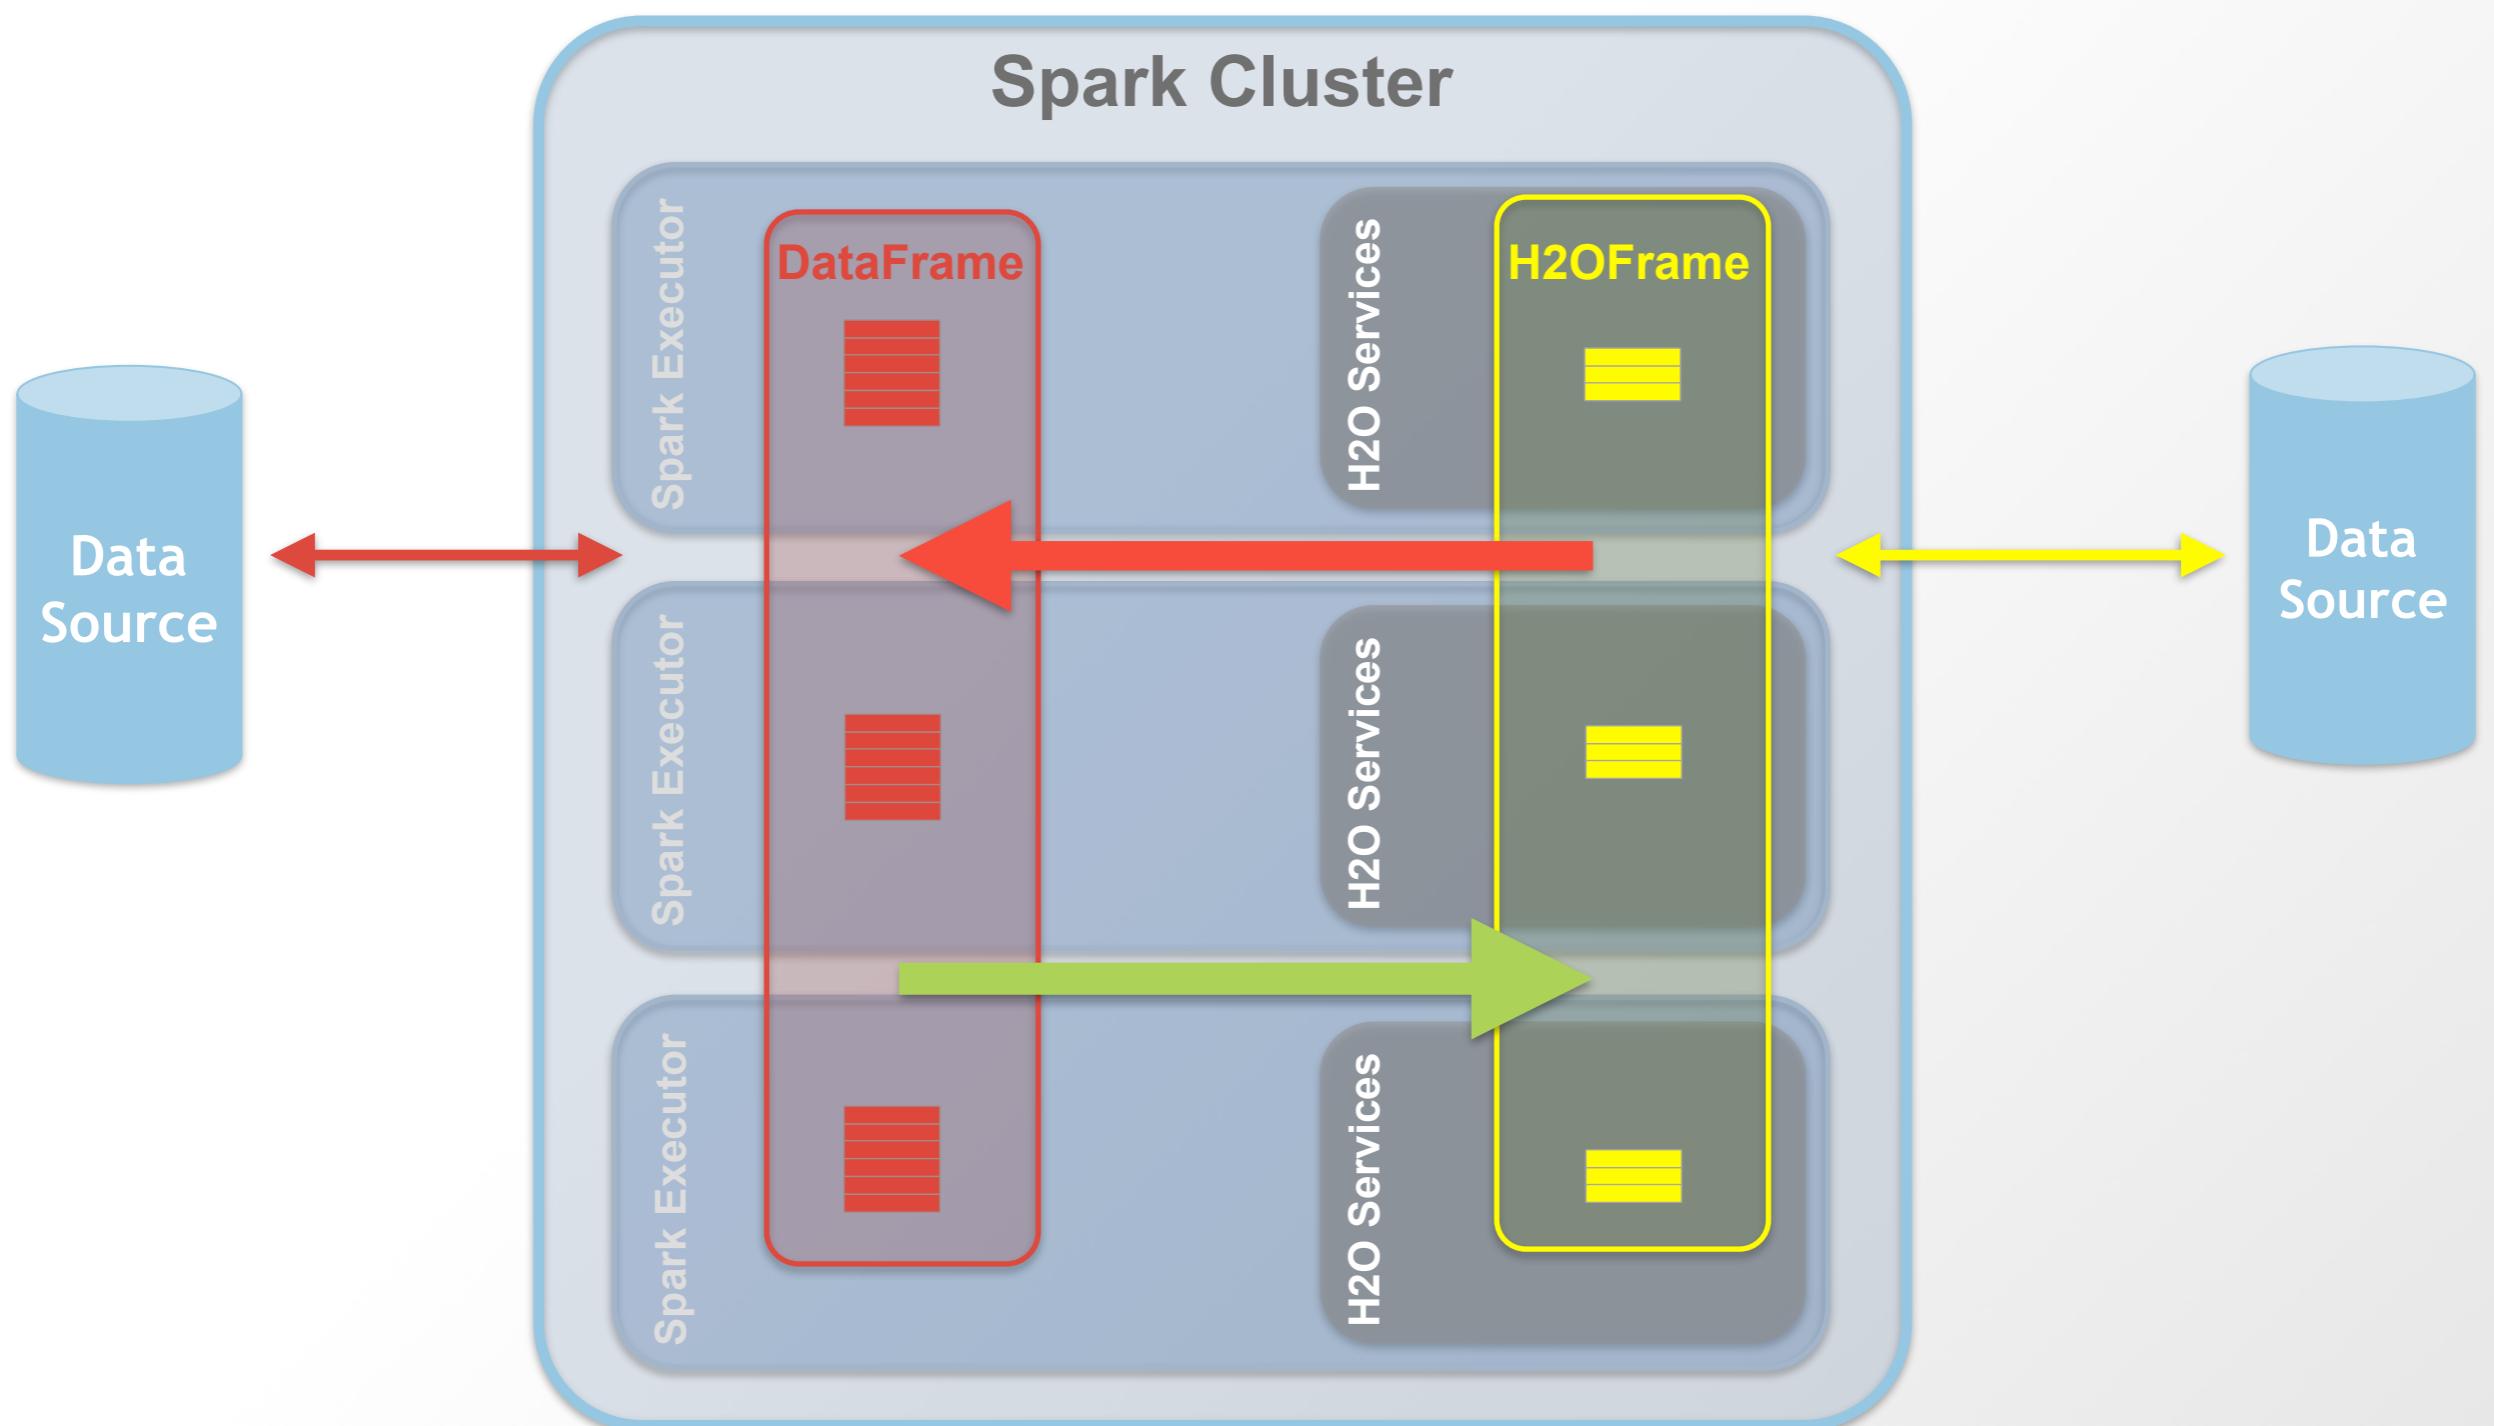
# Sparkling Water Internal Backend



# Internal Backend



# Data Transfers



## **Pros & Cons**

- Advantages
  - Easy to configure
  - Faster ( no need to send data to different cluster)
- Disadvantages
  - Spark kills or joins new executor => H2O goes down
  - No way how to discover all executors

# External Backend

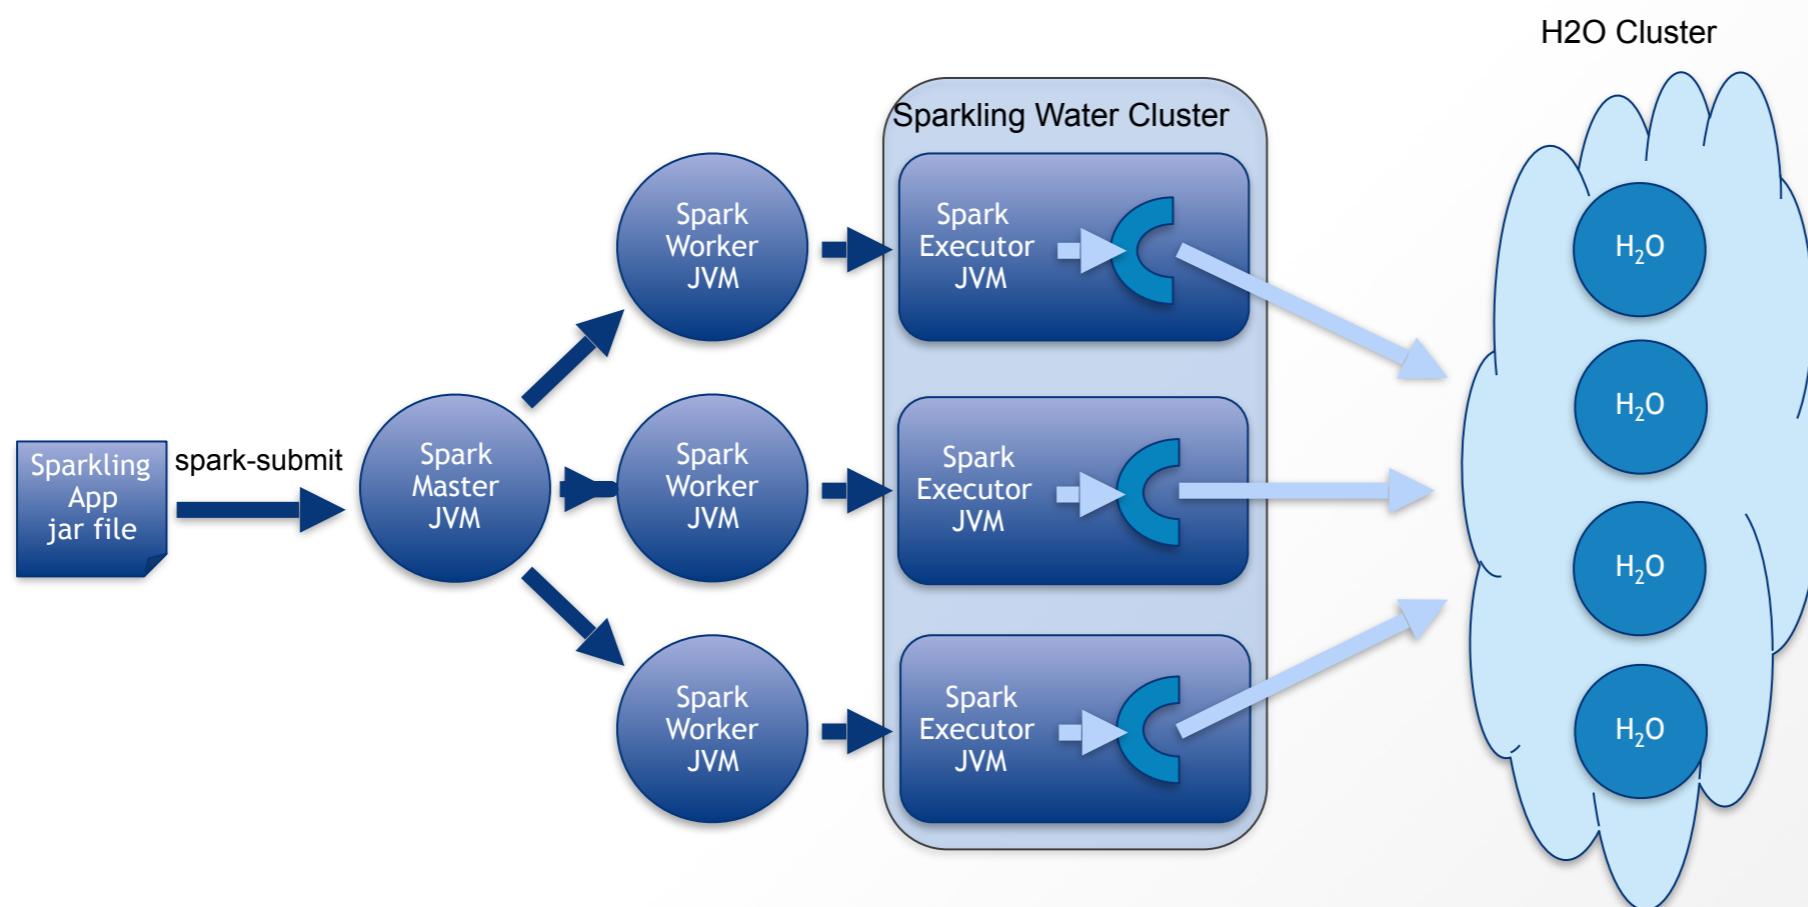
# Overview

- Sparkling Water is using external H2O cluster instead of starting H2O in each executor
- Spark executors can come and go and H2O won't be affected
- Start H2O cluster on YARN automatically

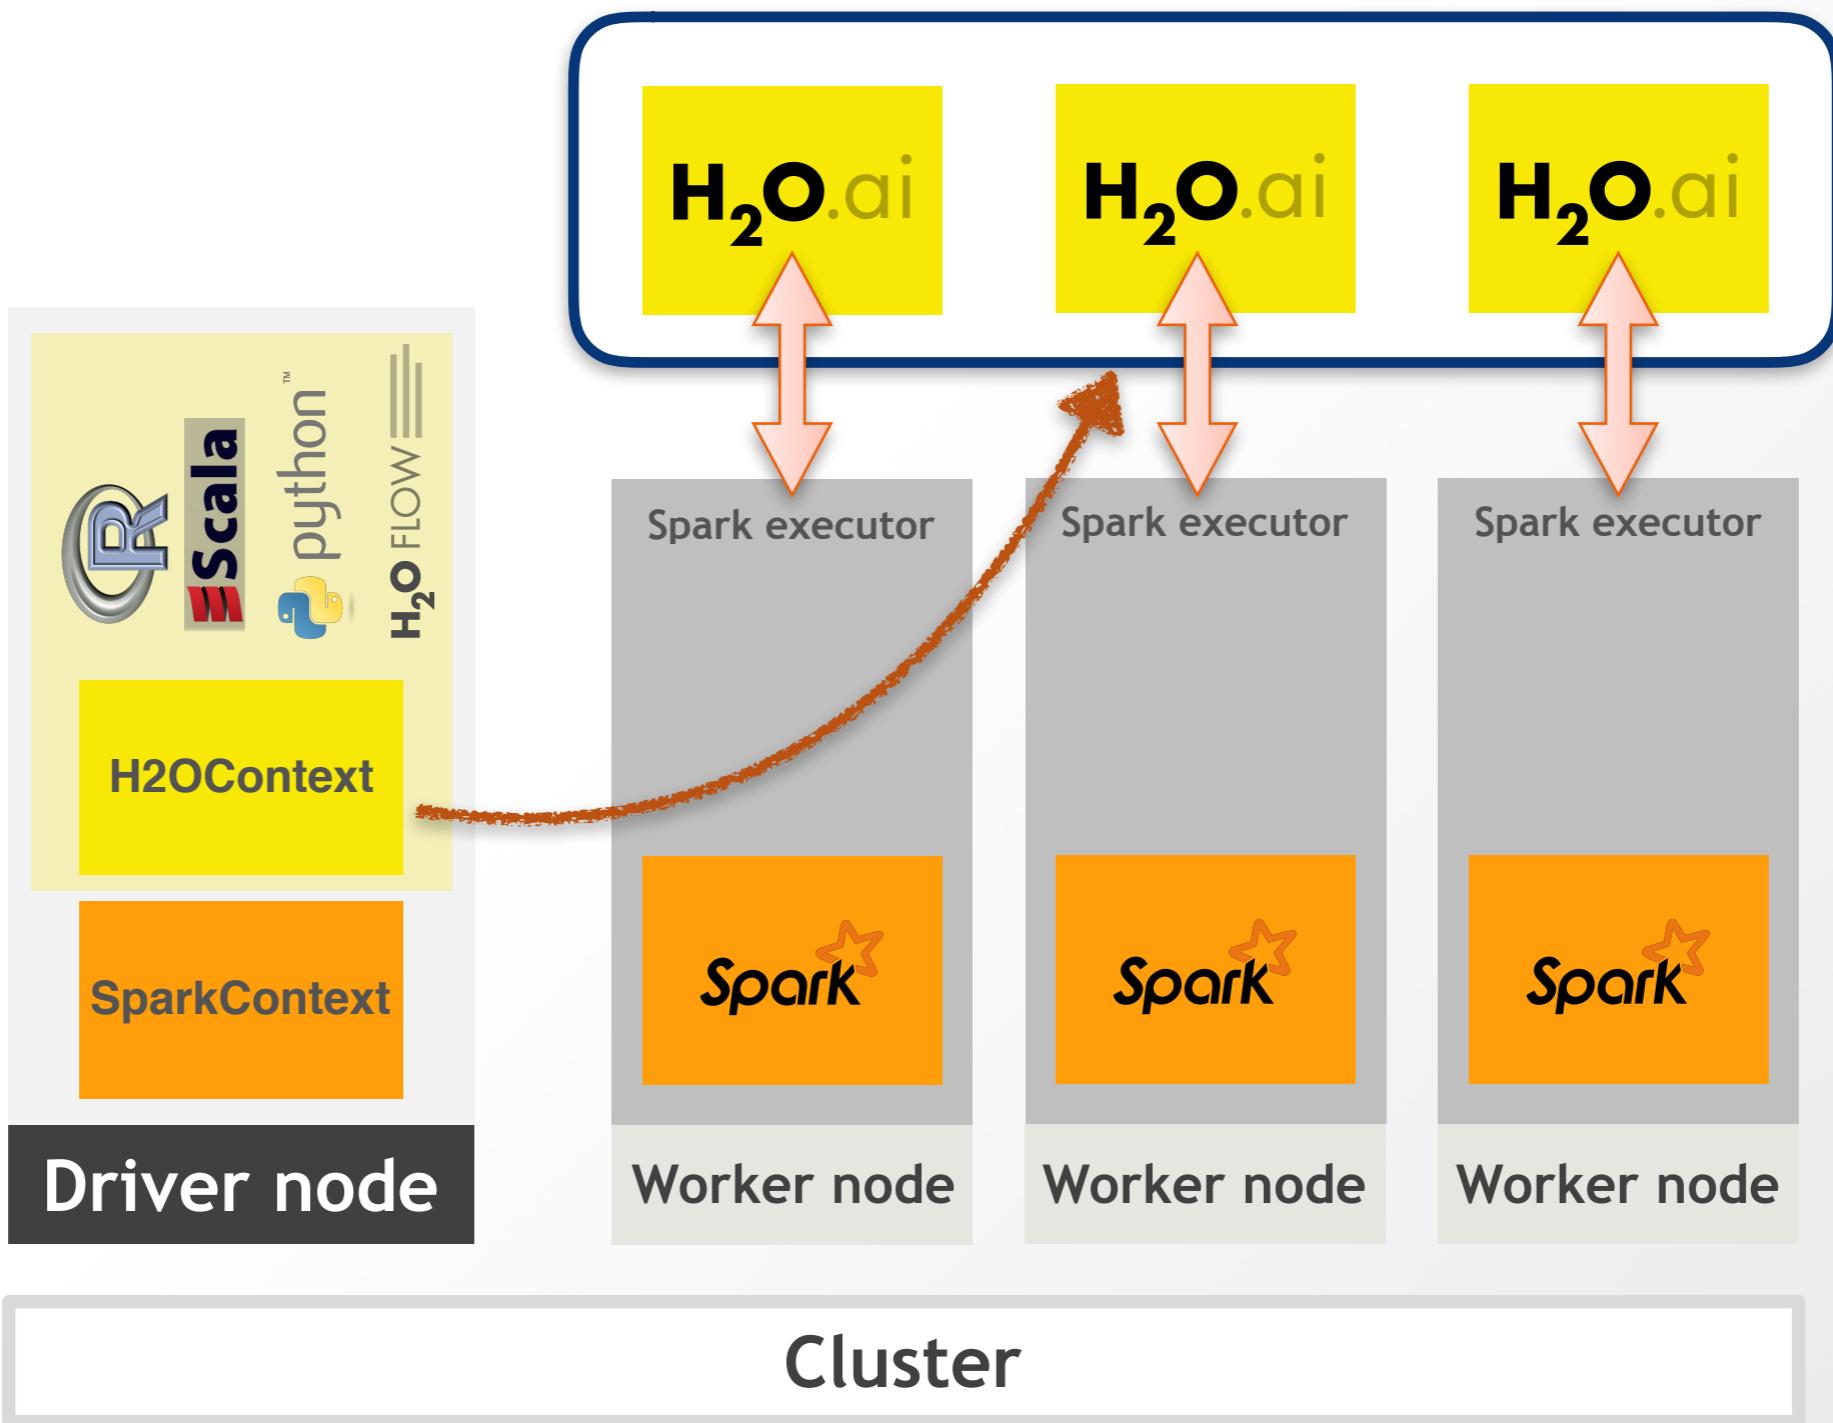
# Separation Approach

- Separating Spark and H2O
  - But preserving same API:  
`val h2oContext = H2OContext.getOrCreate("http://h2o:54321/")`
- Spark and H2O can be submitted as Yarn job and controlled in separation
  - But H2O still needs non-elastic environment (H2O itself does not implement HA yet)

# Sparkling Water External Backend



# External Backend



# **Pros & Cons**

- **Advantages**
  - H2O does not crash when Spark executor goes down
  - Better resource management since resources can be planned per tool
- **Disadvantages**
  - Transfer overhead between Spark and H2O processes
    - under measurement with cooperation of a customer

# Modes

- **Auto Start mode**
  - Start h2o cluster automatically on YARN
- **Manual Start Mode**
  - User is responsible for starting the cluster manually

# More Info

- Documentation: <http://docs.h2o.ai>
- Tutorials: <https://github.com/h2oai/h2o-tutorials>
- Slidedecks: <https://github.com/h2oai/h2o-meetups>
- Videos: <https://www.youtube.com/user/0xdata>
- Events & Meetups: <http://h2o.ai/events>
- Stack Overflow: <https://stackoverflow.com/tags/sparkling-water>
- Google Group: <https://tinyurl.com/h2ostream>
- Gitter: <http://gitter.im/h2oai/sparkling-water>

# Thank you!

Sparkling Water is  
open-source  
ML application platform  
combining  
power of Spark and H2O

Learn more at [h2o.ai](https://h2o.ai)

Follow us at [@h2oai](https://twitter.com/h2oai)

PS: We are hiring!



[www:says-it.com/unclesam/](http://www:says-it.com/unclesam/)



H<sub>2</sub>O.ai





H2O AI World London  
29 & 30 October

- H2O.ai is a **Leader** in 2018 Gartner Magic Quadrant for Data Science & Machine Learning Platform.
- H2O AI World is our **Community Event**.
  - 29th October: Hands-on H2O Training
  - 30th October: Tech Talks, Real-World Use Cases, Kaggle Grandmasters Panel & More

Register at: [world.h2o.ai](http://world.h2o.ai)  
Promotional Code: LASTCHANCE